

Toolbox for perception-based music analysis

Concepts, demos, and reference manual

Marc Leman, Micheline Lesaffre, Koen Tanghe
Institute for Psychoacoustics and Electronic Music (IPEM)
Ghent University
Blandijnberg 2
9000 Ghent, Belgium
toolbox@ipem.ugent.be
www.ipem.ugent.be

Version: 1.02 (beta)
Date: 20140106

Copyright ©2014 Ghent University

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled "[GNU Free Documentation License](#)".

Contents

I	Introduction	7
II	Concepts	12
1	Introduction	13
2	Modules	16
2.1	Introduction	16
2.2	Auditory Peripheral Module	19
2.2.1	Introductory description	19
2.2.2	Functional-logical description	20
2.2.3	Signal processing description	21
2.2.4	Implementation	22
2.2.5	Examples	22
2.3	Roughness Module	26
2.3.1	Introductory description	26
2.3.2	Functional-logical description	27
2.3.3	Signal processing description	28
2.3.4	Implementation	29
2.3.5	Examples	29
2.4	Onset Module	31
2.4.1	Introductory description	31
2.4.2	Functional-logical description	31
2.4.3	Signal processing description	32
2.4.4	Implementation	34
2.4.5	Examples	34
2.5	Pitch Completion Module	38
2.5.1	Introductory description	38
2.5.2	Functional-logical description	39
2.5.3	Signal processing description	40
2.5.4	Implementation	41
2.5.5	Examples	41
2.6	Rhythm Module	43
2.6.1	Introductory description	43
2.6.2	Functional-logical description	44

2.6.3	Signal processing description	44
2.6.4	Implementation	45
2.6.5	Examples	45
2.7	Echoic Memory Module	47
2.7.1	Introductory description	47
2.7.2	Functional-logical description	48
2.7.3	Signal processing description	48
2.7.4	Implementation	48
2.7.5	Examples	48
2.8	Contextuality Module	50
2.8.1	Introductory description	50
2.8.2	Functional-logical description	50
2.8.3	Signal processing description	50
2.8.4	Implementation	51
2.8.5	Examples	51
3	Experiments	56
3.1	Introduction	56
3.2	Roughness experiments	57
3.2.1	Introduction	57
3.2.2	Method	57
3.2.3	Application	57
3.2.4	Results and discussion	62
3.3	Tonality induction experiments	65
3.3.1	Introduction	65
3.3.2	Method	65
3.3.3	Application	66
3.3.4	Results and discussion	74
4	Applications	78
4.1	Introduction	78
4.2	Roughness applications	79
4.2.1	Introduction	79
4.2.2	Method	79
4.2.3	Application	79
4.2.4	Results and discussion	79
4.3	Rhythmic pattern extraction	80
4.3.1	Introduction	80
4.3.2	Method	80
4.3.3	Application	80
4.3.4	Results and discussion	86
Acronyms		92
References		93

III Reference manual	95
5 Introduction	96
6 General information	97
6.1 Conditions/Disclaimer	97
6.2 GNU Free Documentation License	97
6.3 System requirements	107
6.4 Installation	107
6.5 Reporting problems	108
6.6 Updates, ideas and remarks	109
6.7 About the IPEM Toolbox concepts	109
7 About the documentation	110
7.1 Documentation structure	110
7.2 Used terminology	111
7.2.1 Prefixes	111
7.2.2 Signal and SampleFreq	111
7.2.3 ANI, ANIFreq and ANIFilterFREQs	111
7.2.4 FileName and FilePath	111
7.2.5 Empty and not specified input arguments	111
7.2.6 PlotFlag	112
8 Function reference	113
General functions	114
IPEMAMTone	115
IPEMAadaptLevel	116
IPEMAnimateSlices	117
IPEMBatchExtractSoundFragment	119
IPEMBellShape	121
IPEMBlockDC	122
IPEMCalcANI	123
IPEMCalcANIFromFile	125
IPEMCalcCentroid	127
IPEMCalcCentroidWidth	128
IPEMCalcFFT	129
IPEMCalcFlux	131
IPEMCalcMeanAndVariance	132
IPEMCalcNoteFrequency	133
IPEMCalcOnsets	135
IPEMCalcOnsetsFromANI	136
IPEMCalcPeakLevel	137
IPEMCalcRMS	138
IPEMCalcRoughnessOfToneComplex	139
IPEMCalcRoughnessOverSubparts	140

IPEMCalcSpectrogram	141
IPEMCalcZeroCrossingRate	143
IPEMCheckVersion	144
IPEMClip	145
IPEMCombFilter	146
IPEMContextualityIndex	147
IPEMConvertToAMNoise	149
IPEMConvertToClickSound	150
IPEMConvertToMIDINoteNr	151
IPEMConvertToNoteString	152
IPEMCountZeroCrossings	153
IPEMCreateMask	154
IPEMDoOnsets	155
IPEMEnsureDirectory	156
IPEMEnvelopeFollower	157
IPEMExportFigures	158
IPEMExtractSegments	159
IPEMFadeLinear	160
IPEMFindAllPeaks	161
IPEMFindNearestMinima	162
IPEMFindNoteFromFrequency	163
IPEMGenerateBandPassedNoise	165
IPEMGenerateFrameBasedSegments	166
IPEMGeneratePitchShiftScript	168
IPEM GetAllCombinations	170
IPEMGetCrestFactor	171
IPEMGetFileList	172
IPEMGetKurtosis	173
IPEMGetLevel	174
IPEMGetRolloff	175
IPEMGetRoughnessFFTReference	176
IPEMGetSkew	177
IPEMHandleInputArguments	178
IPEMHarmonicTone	180
IPEMHarmonicToneComplex	181
IPEMLeakyIntegration	183
IPEMLoadANI	184
IPEMMECAnalysis	185
IPEMMECExtractPatterns	186
IPEMMECFindBestPeriods	188
IPEMMECReSynthUI	190
IPEMMECSaveResults	192
IPEMMECSynthesis	194
IPEMNormalize	196
IPEMOnsetPattern	197

IPEMOnsetPatternFilter	198
IPEMOnsetPeakDetection	199
IPEMOnsetPeakDetection1Channel	200
IPEMPeriodicityPitch	201
IPEMPlaySoundWithCursor	203
IPEMPlotMultiChannel	204
IPEMReadSoundFile	206
IPEMReplaceSubStringInFileNames	208
IPEMRescale	209
IPEMReshape	210
IPEMRippleFilter	212
IPEMRootDir	213
IPEMRotateMatrix	215
IPEMRoughnessFFT	216
IPEMRoughnessOfSoundPairs	218
IPEMSaveANI	220
IPEMSaveVar	221
IPEMSetFigureLayout	222
IPEMSetup	223
IPEMSetupPreferences	224
IPEMShepardTone	225
IPEMShepardToneComplex	226
IPEMShowSoundFile	227
IPEMSineComplex	228
IPEMSnipSoundFile	229
IPEMSnipSoundFileAtOnsets	231
IPEMStripFileSpecification	232
IPEMToneScaleComparison	233
IPEMToolboxVersion	234
IPEMWriteSoundFile	235
Demo functions - MEC pattern extraction application	236
IPEMDemoMECRhythmExtraction	237
IPEMDemoStartMEC	238

Part I

Introduction

Toolbox for perception-based music analysis

Our motivation

The motivation for developing a toolbox for perception-based music analysis starts from the observation that much of the world music production (seen over different cultures and time periods) has no score representation, and that if a score representation is available, it still represents but a small percentage of what musical communication really is about. A sonological analysis of musical sounds may provide steps towards a better understanding of the components that contribute to musical information processing but is also insufficient in view of how humans deal with musical signals.

Up to the 19th century, the score was indeed the main carrier for musical representation. And as a consequence musicology had a focus on pitch and duration analysis, the main musical parameters to encode in scores. But in the 20th century, thanks to the breakthrough of electronic means for music production and analysis, more powerful devices allowed for the control of physical aspects of music. New methods for musical sound analysis have been developed and the focus on pitch has been extended towards timbre and sound effects. By the end of the 20th century, digital technology laid the foundation for computer modeling of perception and cognition, thereby offering a global human information processing approach to music analysis. In this approach aspects of musical communication that are based on pitch, timbre, texture, expression, spatialization etc... can be studied in close connection with empirical based scientific disciplines.

Technology has revolutionized our conception of what kind of music research will be possible in the 21st century. But it is the task of musicologists, in collaboration with other scientists, such as engineers, psychologists, and brain scientists, to make these dreams more concrete. Our basic motivation for developing this toolbox is a *pragmatic* one: if the musicology aims at understanding music as a social and cultural phenomenon embedded in the physical world and mediated through human faculties of perception, cognition, and processing of expressive communication, then new tools must be developed that allow a fully integrated approach.

Our aims

Our aim is to provide a foundation of music analysis in terms of human perception. We start from sound and take human perception as the basis for musical feature extraction and higher-level conceptualization and representation. To achieve our goals we use computer modeling of the

human auditory system. These models are quite complex, and there is indeed a thorough need for higher-level working tools and companion manuals that focus on music. Well, this is exactly what we offer: a toolbox for perception-based music analysis within the modern laboratory environment of MATLAB.

Our approach

The approach advocated here is both ecological and naturalistic. *Ecological* means that the human information processing system *and* the musical environment are considered as a global unity in which musical content is an emerging outcome of an interactive process. *Naturalistic* means that methods from the natural sciences are used to study phenomena of the musical mind. In particular, we refer to computer modeling and experimenting (see [Leman and Schneider \(1997\)](#), and [Leman \(1999\)](#) for more details about the historical, epistemological and methodological foundations).

The content of this document

The toolbox provides a set of functions that allow researchers to deal with different aspects of feature extraction in perception (i.e. chord and tonality, pitch, roughness, onset detection, beat and meter extraction, ...) but we also offer a global concept, a *philosophy*, say, of how to deal with perception-based music analysis. In order to make clear what we mean by this we present the main concepts and offer some elaborate demonstrations of how the toolbox has been used in our research. The reference manual gives an overview of all available functions, which are in fact the tools you can work with.

Our audience

The IPEM Toolbox is useful for researchers interested in content analysis and automatic extraction and description of musical features. Students may use the toolbox for acquiring practical knowledge in perception-based auditory modeling. We encourage cooperation and solicit feedback.

Our policy

Users are kindly requested to send their comments to [toolbox\[at\]ipem\[dot\]ugent\[dot\]be](mailto:toolbox@ipem.ugent.be). Users can also send us a request to add their own modules, scripts or demos. We will only consider fully documented and working code, however (see the [Function Reference](#) for more details on what we see as "good coding behaviour").

Acknowledgement

The toolbox for perception-based music analysis (Version: 1.02 (beta)) has been developed at IPEM (Institute for Psychoacoustics and Electronic Music) research center of the Department of Musicology at the Ghent University, with support from the Research Council of the University (BOF) and the Fund for Scientific Research of Flanders (FWO). The promoter of this ongoing project is Marc Leman.

Thanks

The authors would like to thank Dirk Moelants and Francesco Carreras for testing this toolbox and giving valuable comments during the course of its development.

How to read this text?

Part II introduces the reader to the concepts and worked out applications of the toolbox. Part III contains the reference manual. We strongly recommend the lecture of the concepts part, in particular, the introductory descriptions of the modules, as well as the associated examples. They will give you a flavor of what the toolbox can offer in terms of music analysis. The signal processing sections can be left out for a second reading.

The concepts part also contains an examples section that shows how to use the very basic toolbox commands. Please take into account that we don't present an exhaustive tutorial of all functions described in the reference manual!

If you want to have an idea of how to apply the toolbox for somewhat larger projects in music analysis, you can have a look at the applications. For the list of all functions, see the reference manual.

It's true, you can read the first part (concepts and applications) without using the toolbox. But we strongly recommend to use the toolbox from the very beginning. Please, go to [Reference Manual](#), [General Information](#), and install the toolbox... and don't forget to read the conditions/disclaimer section!

Part II

Concepts

Chapter 1

Introduction

The toolbox is a collection of MATLAB *functions* for perception-based music analysis. These "building blocks" are combined to form *modules*, related to musical perceptual phenomena, which will be described in the following chapter. The next chapter will evaluate some of these modules with respect to data from psychoacoustical experiments. In a subsequent chapter these modules are then used in some applications to demonstrate their effectiveness.

The modules are described from different points of view so that you can access them at different levels. We offer you an introductory, functional-logical, signal processing, implementation, and example-based description.

1. The *introductory description* provides just a simple verbal description of what the module does, how it is situated in our global conception, what the inputs are and what outputs are generated.
2. The *functional-logical description* allows a concise description which resembles the way in which the MATLAB functions are to be used while programming in MATLAB. The approach is vector oriented (see next paragraph).
3. The *signal processing description* gives a mathematical description of the modules in terms of a functional equivalence model of physiological processes. A trade-off had to be found between psychoacoustical and physiological accurateness and processing efficiency. Our approach is based on filter theory and vector processing rather than physical modeling.¹
4. The *implementation description* concerns the way in which the MATLAB functions are implemented. Signal processing functions of the MATLAB signal processing toolbox have been used where possible. See the [Reference Manual](#) for more details on the implemented MATLAB functions.
5. Finally for each module we also provide a number of *examples*.

¹An essential aspect, for example, concerns the auditory peripheral part. Rather than having a physical model of the resonance of the basilar membrane and subsequent hair cells in response to sound, the effects of the cochlear hydromechanics and hair cells are simulated by means of an array of overlapping (asymmetric) band-pass filters and subsequent non-linear transformation into neural firing rate-code (not firing spike-code).

There is another important thing that you should know. The modules are related to auditory information processing. That means: we start from sound and transform sounds into either auditory images or inferences. *Auditory images*, or images in short, reflect features of the sound as internal representations, in other words as brain activations so to speak. *Inferences*, on the other hand, provide derived information that can be compared to human behavior.

Hence our dual validation model: images and associated processes are compared with human physiology, while inferences are compared with human behavioral responses (fig. 1.1). This is why we introduce the following important notions in our approach, which define our global internal representational framework:

- An *image* has two aspects: (i) its content represents features related to the musical signal, and (ii) it is assumed to be carried by an array of neurons. From the point of view of computer modeling, an image is an ordered array of numbers (= a vector) whose values represent neural activation. We conceived of neural activation in terms of firing *rate-code*, that is, the probability of neuronal spiking during a certain time interval. Hence images are to be conceived of as snapshots of the brain activation (or better: what we believe to be snapshots of brain activation). A distinction will be made between different types of images (such as *primary* images, *pitch* images, *spectral* images, etc...). A main characteristic and limitation of our current state-of-the-art is that images are basically *frame based*. In other words, we do not yet consider an *object-based* representation of musical content.
- Processes dealing with image transformations are called *image transformation processes*. They transform sounds into images, and images into other images. This is a good place to say that different memory systems carry different kinds of images, hence our inclination to speak about short-term images and long-term images in association with short-term memories and long-term memories.
- *Inferences* are outputs that can be directly compared with listener's behavioral responses to musical information processing.
- *Inference processes* compare images, inspect images or extract features from images.

Figure 1.1 gives an overview of the internal representational framework.

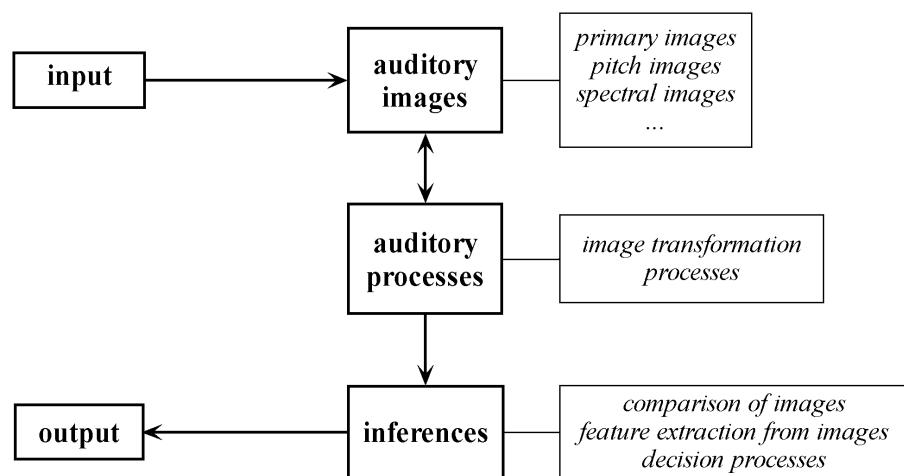


Figure 1.1: Overview of the internal representational framework.

Chapter 2

Modules

2.1 Introduction

This chapter describes the modules in the toolbox. A *module* represents a well-defined process leading to images and (possibly) an inference. *Well-defined* here means that it corresponds to a relevant musical perceptual phenomenon. Figure 2.1 gives an overview of what we have in mind as a global picture.

The modules on the left are sensory based, the ones on the right are cognition based and the ones in the middle are perception based.

- Sensory modules involve immediate memory and are based on stimulus-driven processing. We assume that they are located in the periphery of the auditory system.
- Perception modules involve short-term memory and specific image mapping from the temporal domain into the spatial domain. We assume that they are located in the brain stem.
- Cognitive modules involve long-term memory, learning and categorization. We assume that they are located in the cortex.

The modules are furthermore organized from texture (top) to structure (bottom). This is obviously a simplification but one that according to our feeling makes sense, at least to some extent. You should be aware of the fact that this is just one possible way to visualize the modules and that other visualizations of the global picture are possible. Notice one more feature of this chart, in particular the expressions between quotes, such as "roughness". These labels refer to inferences associated with the modules. In figure 2.2 the modules in the toolbox are highlighted.

Recall that a module is a well-defined process but that it may consist of one or more toolbox functions. We learn you how to use these functions while exploring the modules. If it is the first time that you read this introduction you can skip the functional-logical and signal processing descriptions and limit yourself to the introductory description and the examples. This should provide enough information to acquire an idea of what the module is doing.

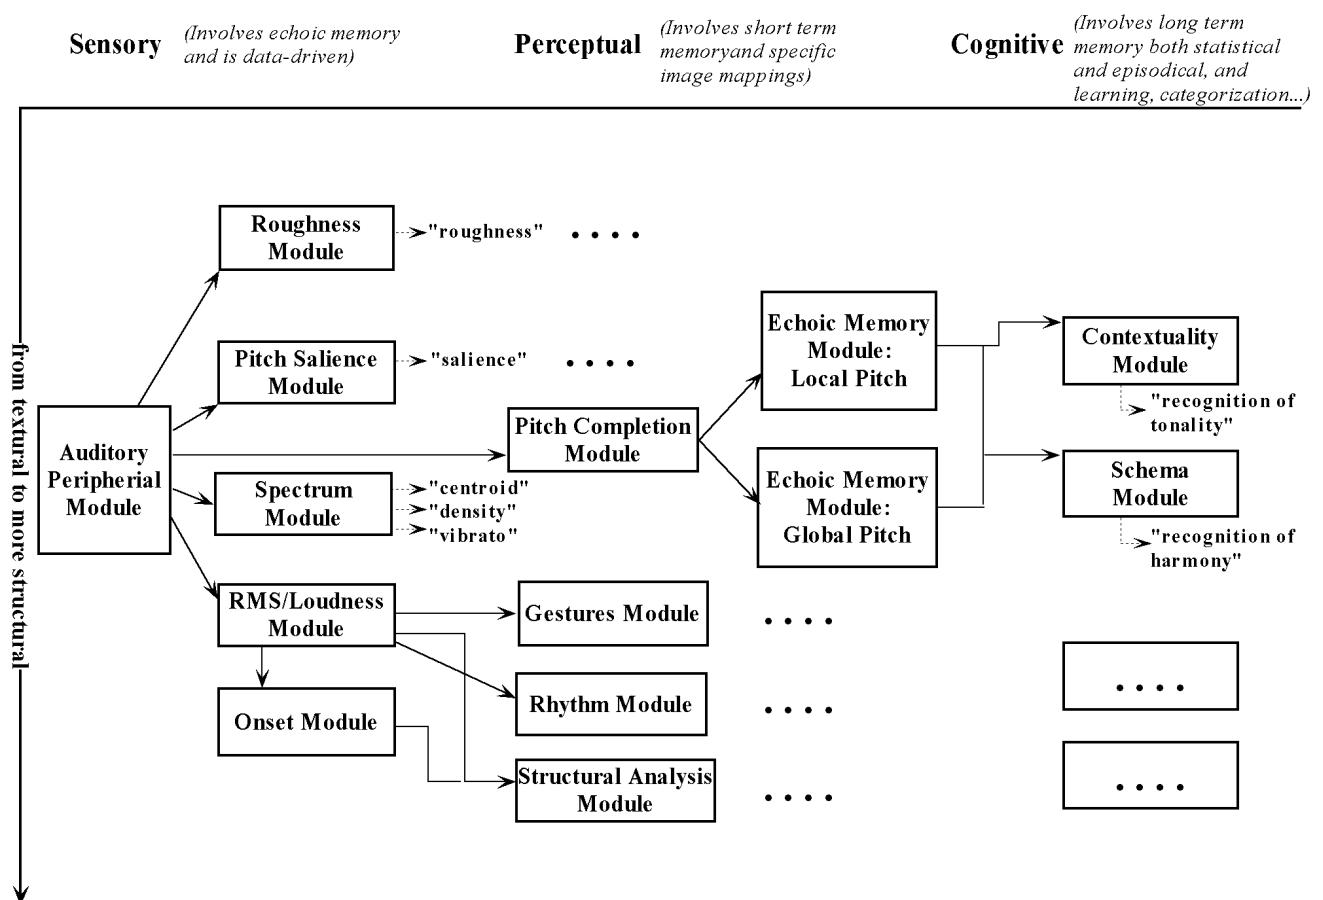


Figure 2.1: Chart of image transformation modules, organized according to the distinction between sensory, perceptual, and cognitive information processing (horizontal axis), and from textural to structural (vertical axis). The chart is not exhaustive.

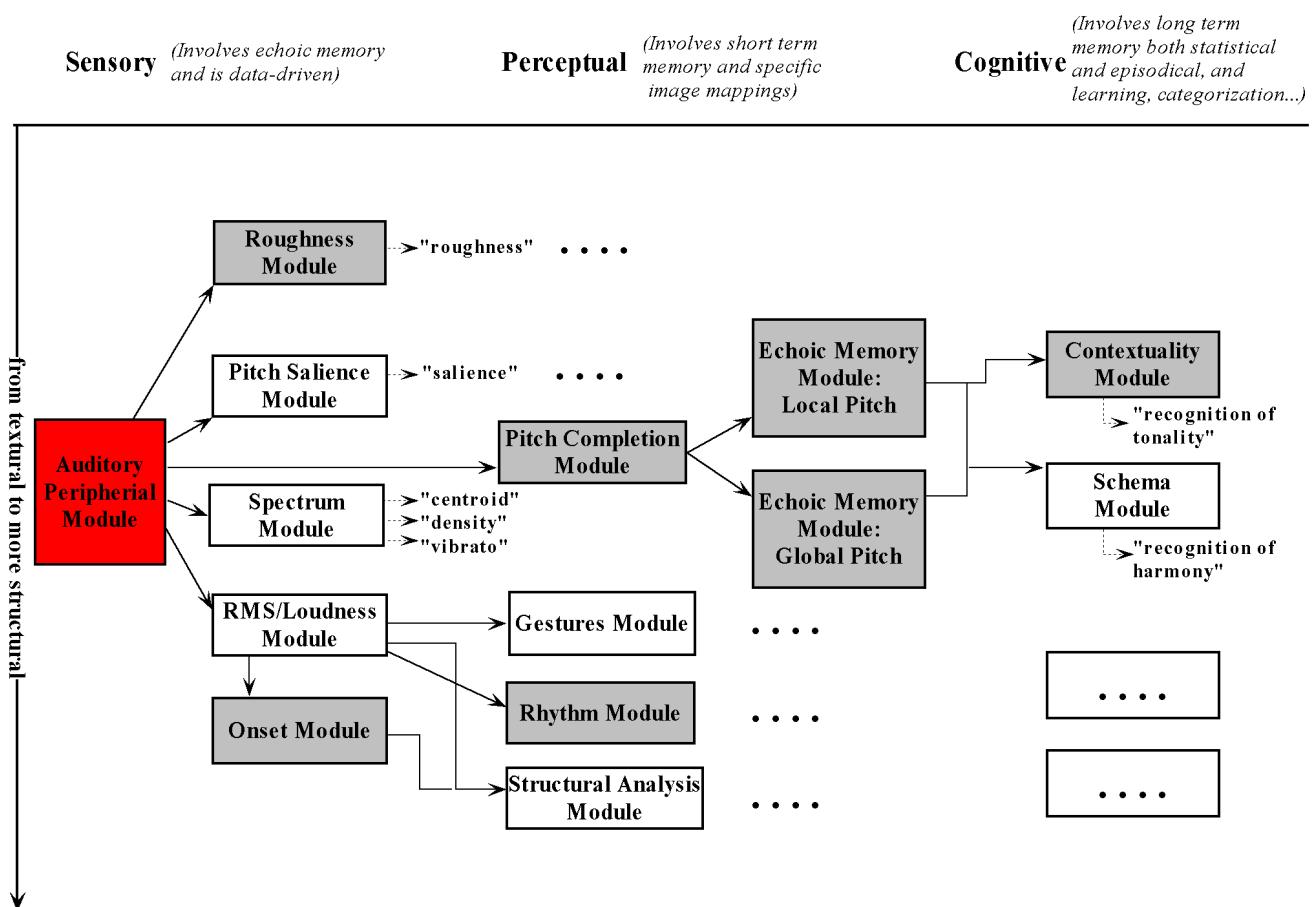


Figure 2.2: Situation of the toolbox modules in the global conception

2.2 Auditory Peripheral Module

2.2.1 Introductory description

The Auditory Peripheral Module (APM) (fig. 2.3) takes as input a sound and gives as output the *auditory primary image* which is a kind of physiological justified representation of the auditory information stream along the VIIth cranial nerve. The musical signal is decomposed in different subbands and represented as neural patterns. The patterns are rate-codes, which means that they provide the probability of neuronal firing during a short interval of time.

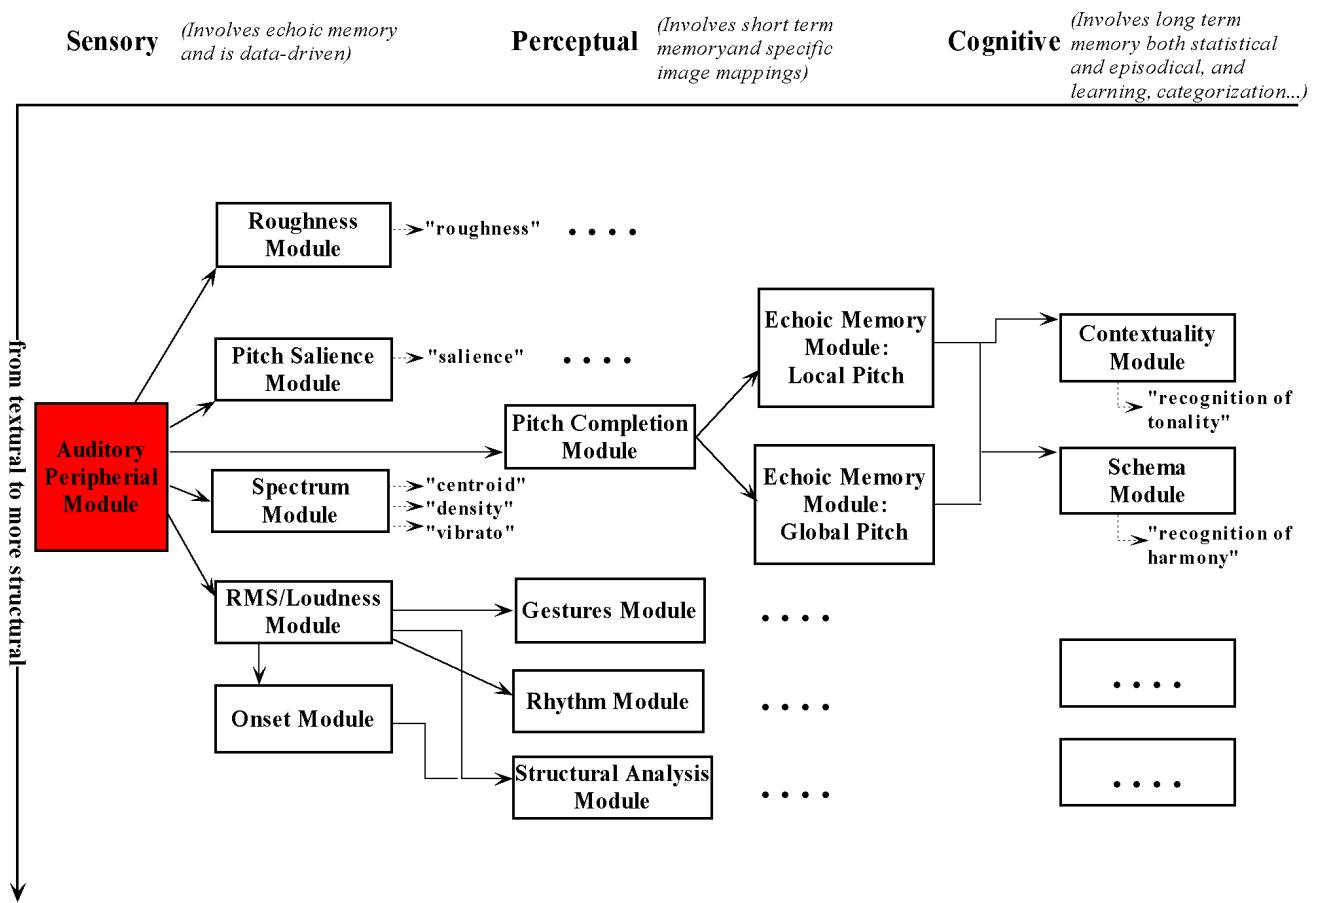


Figure 2.3: Chart of image transformation modules, with APM highlighted

The Auditory Peripheral Module that we use is an adapted version of [Van Immerseel and Martens \(1992\)](#) model of the auditory periphery. The processing stages involve:

- Simulation of the filtering of the outer and middle ear.
- Simulation of basilar membrane resonance in the inner ear. This is implemented by an array of band-pass filters whose center frequencies are spaced on a critical band scale (Bark scale). The bandwidth of each filter equals a critical band.

- Simulation of a hair cell model. This converts the band-pass filtered signals into neural rate-code patterns. This operation deals with half-wave rectification and dynamic range compression.

Figure 2.4 gives a view of the transformation of sound into primary images.

Image Transformation Process

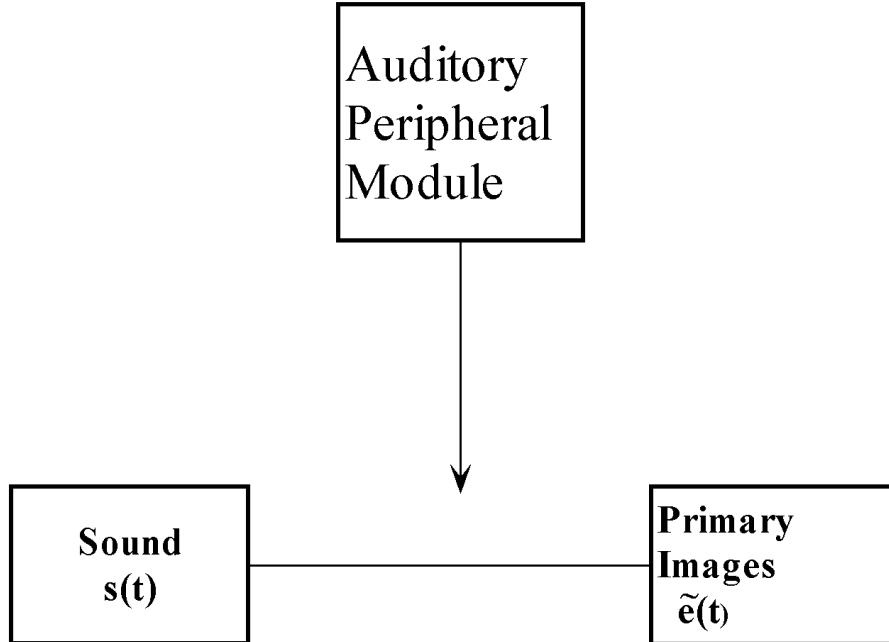


Figure 2.4: Image Transformation Process: Auditory Peripheral Module

2.2.2 Functional-logical description

The APM can be described as a function which transforms a musical sound signal $s(t)$ into a set of patterns $e_n(t)$ (with $n = 1, 2, \dots, C$) that encode the responses of auditory neuronal fibers spread along the basilar membrane (see Equation 2.1). The auditory filters, sometimes also called *channels*, simulate the neuronal band-pass characteristics and the firing rate encoding. The frequency range covered by the auditory filters depends on the center frequencies of the filters, the number of chosen channels and the chosen distance between the channels. In most of our simulations we chose forty channels ($C = 40$), half a critical band apart from each other. This covers a range from 140 Hz to 8877 Hz. As a short-hand we use the notation:

$$APM : s(t) \rightarrow \begin{bmatrix} e_1(t) \\ e_2(t) \\ \dots \\ e_C(t) \end{bmatrix} = e(t, c) = \tilde{e}(t)$$

or, alternatively,

$$APM : s(t) \rightarrow \tilde{e}(t) \quad (2.1)$$

The pattern denoted $\tilde{e}(t)$ is called the *primary image* or *auditory nerve image (ANI)* of $s(t)$. It should be considered the auditory counterpart of D. Marr's *primal sketch* in the domain of visual perception (Marr, 1982). The tilde character here denotes a vector in which each vector-component corresponds to one auditory filter. The arrow represents a causal transformation, in this case from a sound into the auditory nerve image.

In our terminology, a *running vector* means that the values of the vector-components change over time, where t denotes time. The pattern $\tilde{e}(t)$ thus represents the neural activation in all subbands over subsequent time steps. The notation $e(t, c)$ means the same, whereas $e(t, c)$ for fixed c is the neural activation in one particular channel c .

2.2.3 Signal processing description

For a full description of this module we refer to Van Immerseel and Martens (1992). We limit ourselves to a more technical verbal summary.

The auditory peripheral module simulates the cochlear mechanical filtering using an array of overlapping band-pass filters. The basic steps of this model, shown in figure 2.5 can be summarized as follows.

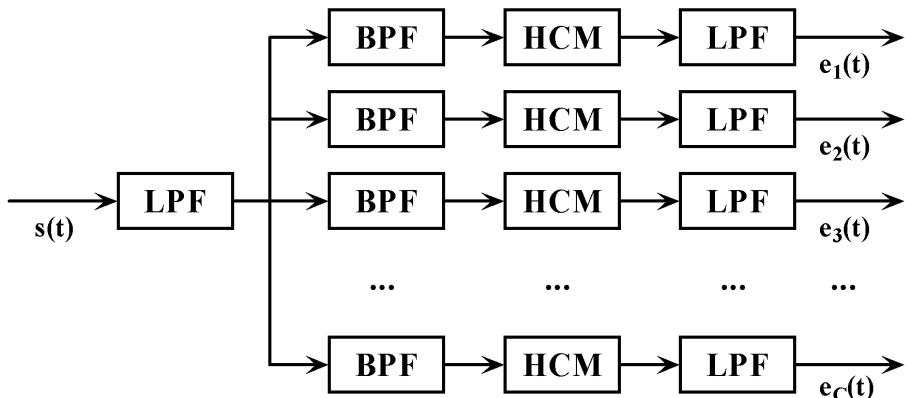


Figure 2.5: Schema of the Auditory Peripheral Module.

- The outer and inner ear filtering is implemented as a second-order low-pass filter (LPF) with a resonance frequency of 4 kHz. This accounts for the overall frequency response of the ear, a coarse approximation to the Fletcher-Munson curves.
- The filtering in the cochlea is implemented by an array of band-pass filters (BPF). In many of the simulations in this book, we use center frequencies that are spaced equidistantly on a critical band scale, having an overlap of half a critical band. Forty channels are used with center frequencies ranging from 141 to 8877 Hz. The filters have a 3 dB bandwidth of one critical band; a low-frequency slope of about 10 dB per critical band unit and a high-frequency slope of about 20 dB per critical band unit.

- The mechanical to neural transduction is performed by a hair cell model (HCM), which is assumed identical in all channels. The HCM is a forward-driven gain controlled amplifier that incorporates half-wave rectification and dynamic range compression. The HCM introduces distortion products that reinforce the low frequencies that correspond to the frequency of the beats.
- A low-pass filter at 1250 Hz does an envelope extraction of the patterns in each channel. This low-pass filter accommodates for the loss of synchronization observed in the primary auditory nerve. The filter can be argued to have an effect on the width of the critical band for high f_c although these may fall out of the scope of most musically relevant tones (Greenwood & Joris, 1996).

The output $e(t, c)$ for fixed c represents the rate-code of neural discharge in channel c . This pattern is also called the *auditory nerve pattern* in channel c .

2.2.4 Implementation

IPEMCalcANI	- Calculates auditory nerve image from signal
IPEMCalcANIFromFile	- Calculates auditory nerve image directly from sound file
IPEMLoadANI	- Loads auditory nerve image from .mat file
IPEMSaveANI	- Saves auditory nerve image to .mat file

2.2.5 Examples

In this section we provide some examples of how to transform musical signals into primary images (auditory nerve images). A straightforward function is **IPEMCalcANIFromFile** which allows you to read in a wav-file. If the wav-file *SchumannKurioseGeschichte.wav* is stored in the "Sounds" subdirectory of the default input directory (see also the **installation** instructions), and you agree with the default 40 subbands and the overlap of 1/2 critical band, then the function is simply:

```
IPEMCalcANIFromFile('SchumannKurioseGeschichte.wav', [], [], 1);
```

Similarly, if *ShepardCChord.wav* is stored in that same directory, you can write:

```
IPEMCalcANIFromFile('ShepardCChord.wav');
```

The results should be similar to what you see on these pages (except the score, of course):

1. Figure 2.6 shows the score of the used excerpt of Schumann's Kuriose Geschichte ¹.
2. Figure 2.7 shows the results of processing this short excerpt with the Auditory Peripheral Module. The upper panel shows the waveform, the lower panel shows the primary image.

¹Robert Schumann "Kinderszenen-Kreisleriana", played by Martha Argerich (Deutsche Grammophon 410 653-2, 1984)

3. Figure 2.8 shows the results of processing a signal consisting of the C major chord, build up from 3 Shepard tones ω with the Auditory Peripheral Module. Again, the upper panel shows the waveform, the lower panel shows the primary image.

To generate a Shepard-chord yourself, you can type (see [IPEMShepardToneComplex](#)):

```
MyShepardChord = IPEMShepardToneComplex([1 0 0 0 1 0 0 1 0 0 0 0],1,22050,1);
```

Now you generated a signal that is stored in `MyShepardChord`. And you can play it in the usual MATLAB way with:

```
sound(MyShepardChord,22050);
```

Pay attention to the fact that `MyShepardChord` is a signal in the MATLAB environment, not a wav-file. To process this signal with the auditory peripheral module you can do:

```
IPEMCalcANI(MyShepardChord,22050);
```

Note that the number 22050 is the sampling rate which should be specified, and `MyShepardChord` is a variable in MATLAB, therefore it should not be between quotes.

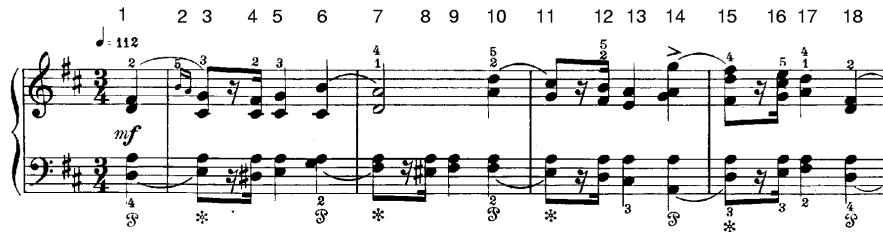


Figure 2.6: The first four measures of Schumann's piece "Kuriose Geschichte"

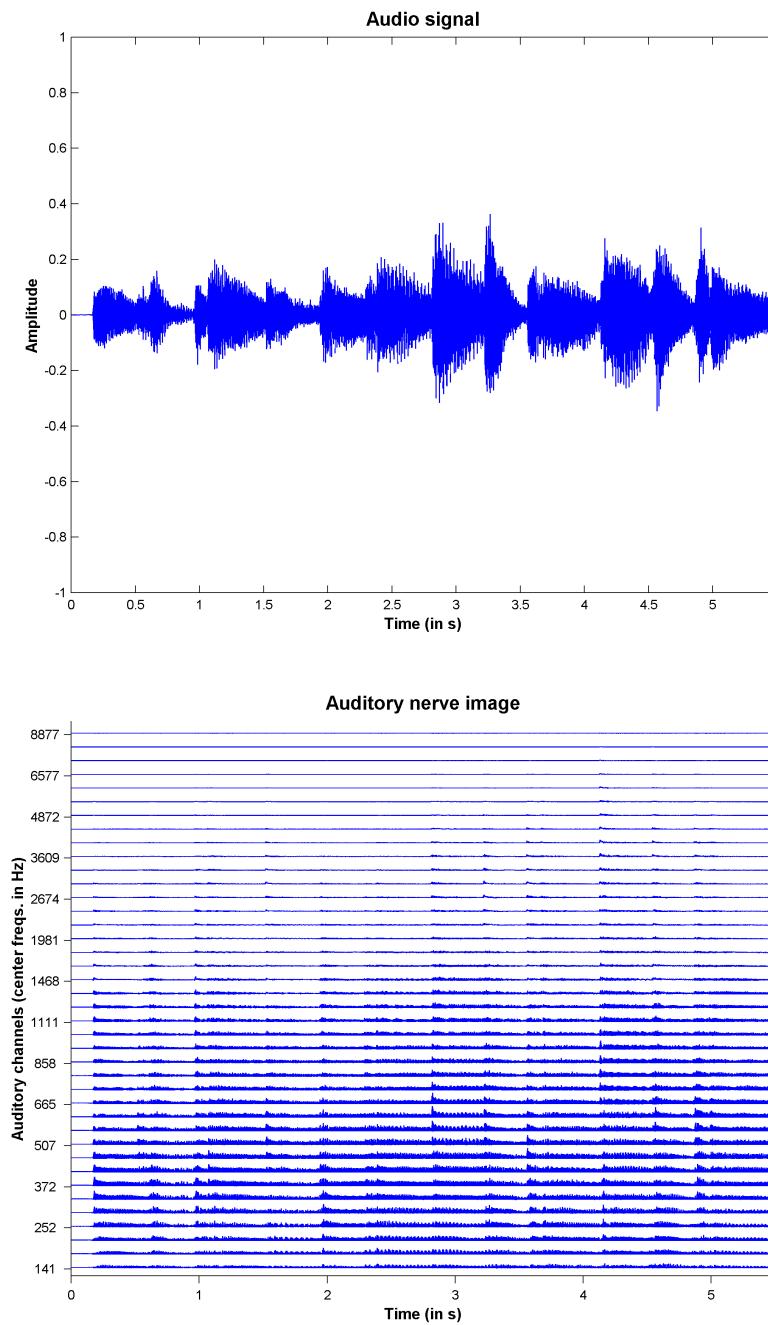


Figure 2.7: Top: waveform of a short excerpt of Schumann's Kuriose Geschichte. Bottom: primary image

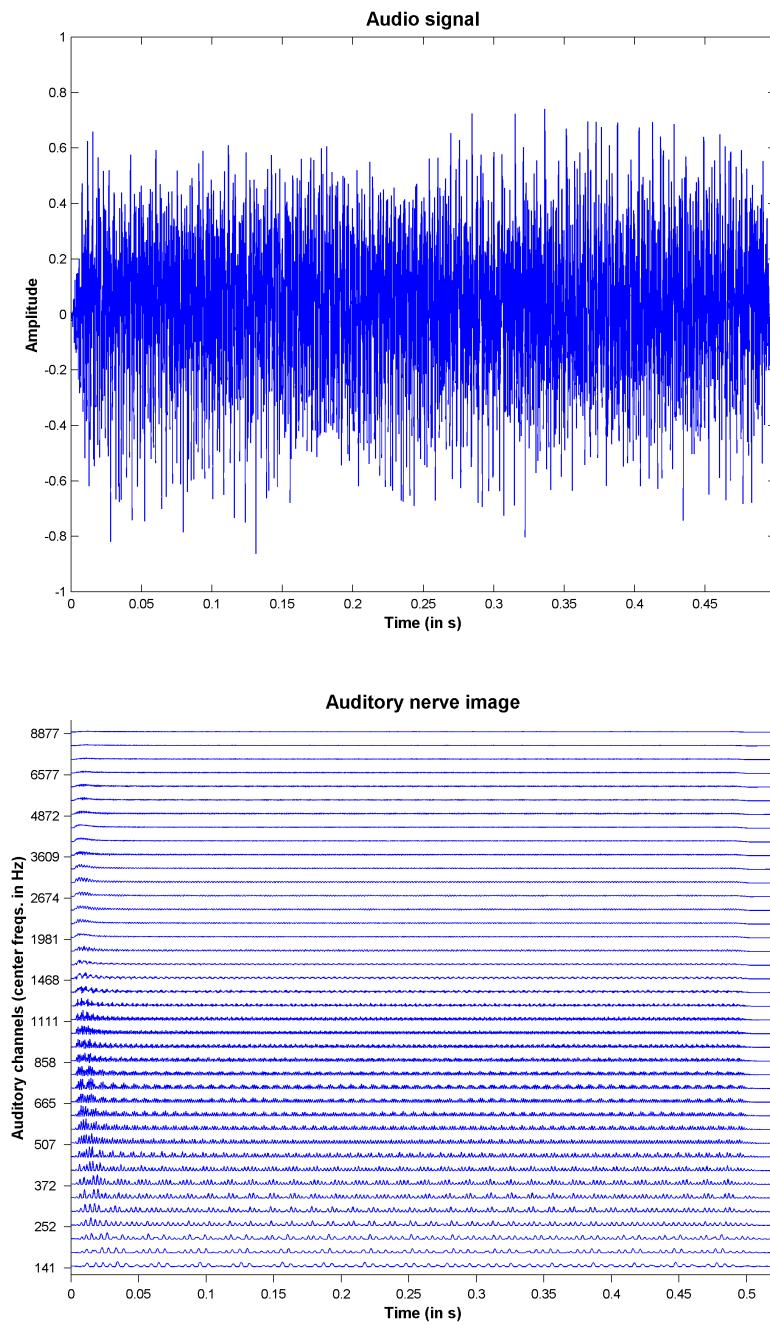


Figure 2.8: Top: waveform of a C major chord using Shepard tones. Bottom: primary image

2.3 Roughness Module

2.3.1 Introductory description

The Roughness Module (RM) calculates the roughness (or equivalently: the sensory dissonance) of a sound. Roughness is considered to be a sensory process highly related to texture perception. Hence, in our global chart of image transformation modules it is localized in the top left (texture and sensory based) section of figure 2.9.

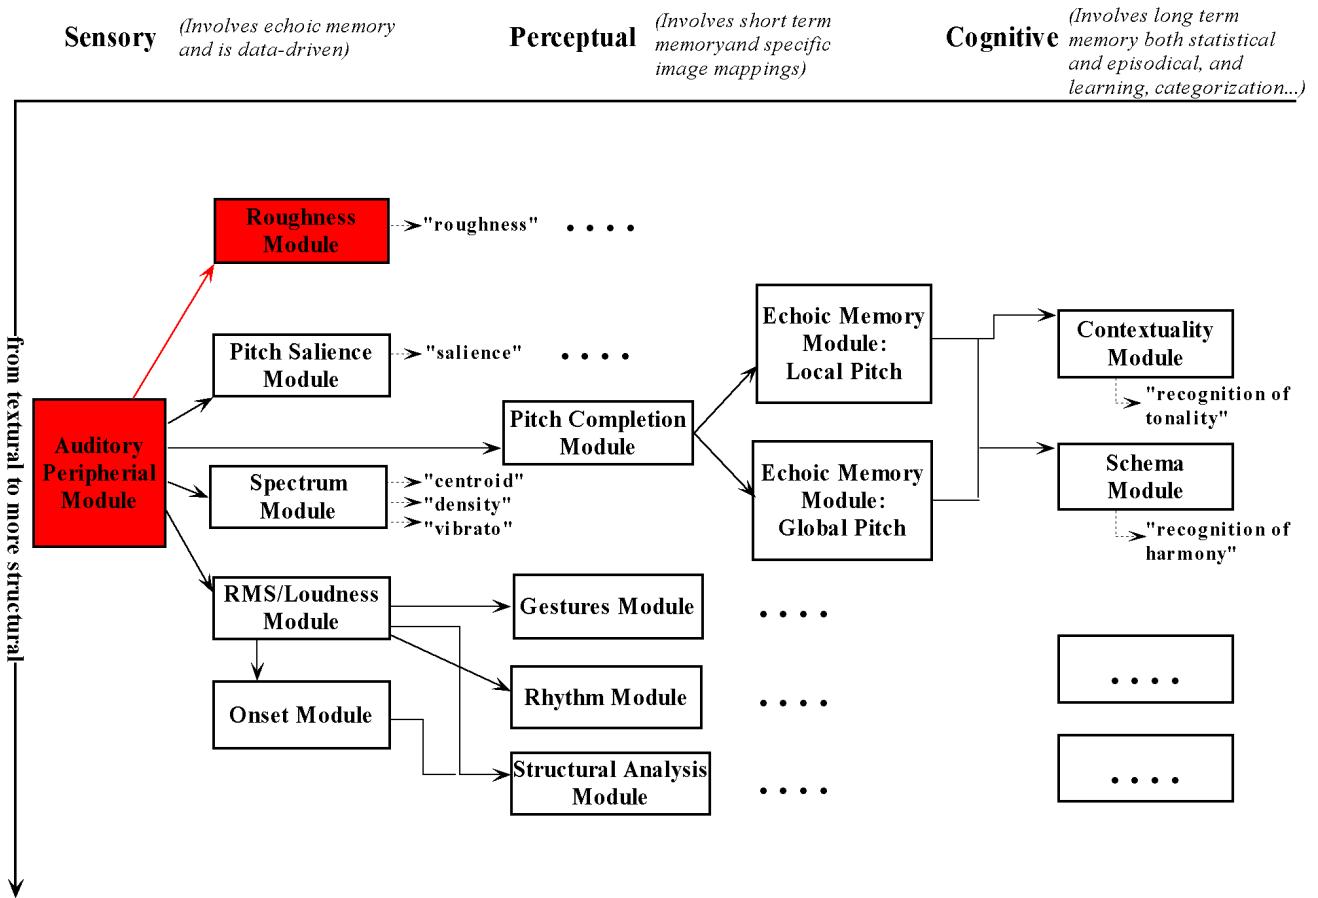


Figure 2.9: Chart of image transformation modules, with RM highlighted

The module takes sound as input (fig. 2.10) and produces a roughness estimation. The estimation should be considered an inference. However, the module offers more than just an inference. The visualization and calculation method of this module is based on Leman (2000b) where roughness is defined as the energy of the relevant beating frequencies in the auditory channels. The model is called the synchronization index model of roughness and it is based on phase-locking to frequencies that are present in the neural patterns. The synchronization index model allows a straightforward visualization of the energy components underlying roughness, in particular (i) with respect to auditory channels, and (ii) with respect to phase-locking synchronization (= the synchronization index for the relevant beating frequencies on a frequency scale). That's why it is more than just an inference. It assumes that neurons somehow extract the energy of the beating frequencies and

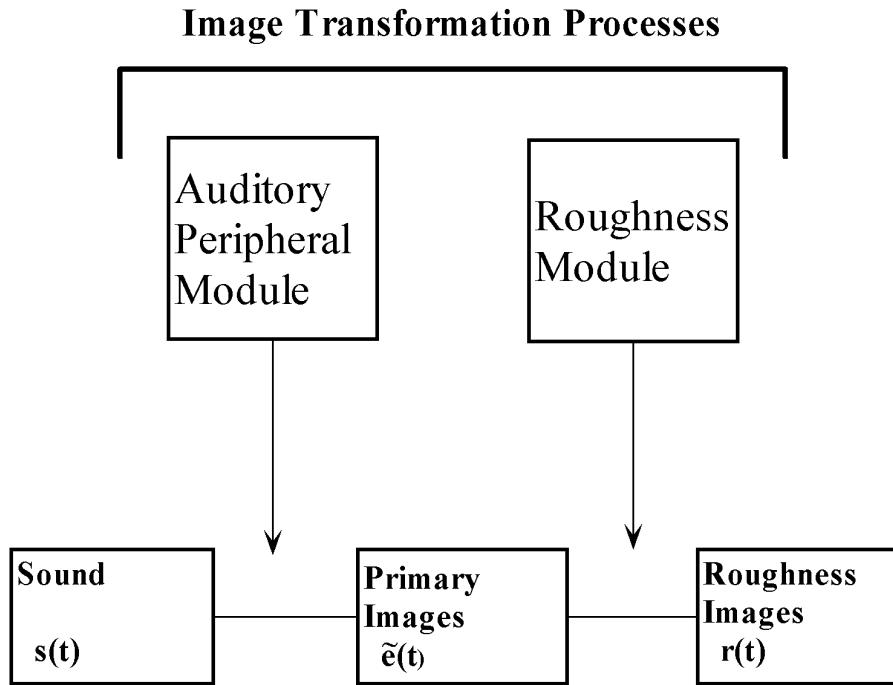


Figure 2.10: Image Transformation Process: Roughness Module

form internal images on which the inference is based.

But what are beating frequencies? Take an amplitude modulated sine wave with carrier frequency f_c and modulation frequency f_m . Such a signal has only three frequencies in the spectrum, in particular f_c , $f_c - f_m$, and $f_c + f_m$. The beating frequency is f_m . Now, in the auditory system, these frequencies are introduced as effective beating frequencies into the spectrum of the neural rate-code patterns. This is due to wave rectification in the cochlea where the lower part of the modulated signal is cut off, and as a result new frequencies are introduced of which the most important ones correspond with the beating frequency f_m and its multiples. As a matter of fact, neurons may synchronize with these frequencies provided that they fall in the frequency range where synchronization is physiologically possible. This mechanism forms a physiological basis for the detection of beats and hence, the sensation of roughness.

According to Javel, McGee, Horst, and Farley (1988, p. 520-521), the synchronization index represents the degree of phase-locking to a particular frequency that is present in the neural pattern. The index can be mathematically expressed as the (normalized) Fourier Series coefficient at the frequency of interest. Using this concept, the degree of roughness can be defined as the sum of the normalized magnitudes or 'energies' of the relevant beating frequencies in the Fourier spectrum.

2.3.2 Functional-logical description

Strictly speaking the roughness module realizes a transformation from auditory nerve images onto roughness values. In a broad sense the module starts from sound involving:

$$APM : s(t) \rightarrow \tilde{e}(t) \quad (2.2)$$

$$RM : \tilde{e}(t) \rightarrow r(t) \quad (2.3)$$

where APM is the Auditory Peripheral Module which is used as the first step, and RM is the transformation from the primary images $\tilde{e}(t)$ to the roughness values $r(t)$ (our Roughness Module in the strict sense). The latter are scalar values which we can compare with human behavioral outputs.

2.3.3 Signal processing description

The synchronization index model calculates roughness in terms of the 'energy' of neural synchronization to the beating frequencies. The 'energy' refers to a quantity which we derive from the magnitude spectrum.

Since the beating frequencies are contained in the lower spectral area of the neuronal pattern $e(t, c)$, the spectral part we are interested in is defined as:

$$B(t, f, c) = F(f, c)D(t, f, c) \quad (2.4)$$

where $F(f, c)$ is a filter whose spectrum is depending on the channel c , and $D(t, f, c)$ is the short-term spectrum in channel c which we define as:

$$D(t, f, c) = \int_{-\infty}^{+\infty} e(t, c)w(t' - t)e^{-j2\pi ft'} dt' \quad (2.5)$$

where $e(t, c)$ is the neural pattern in channel c , and $w(t' - t)$ is a (hamming) window. The *magnitude* spectrum is then defined as $|D(t, f, c)|$, and the *phase* spectrum as $\angle D(t, f, c)$.

In order to be able to reproduce the psychoacoustical data on roughness, the filters $F(f, c)$ should be more narrow at auditory channels where the center frequency is below 800 Hz, and the filters should be attenuated for high center frequencies as well. The proposed set of filters $F(f, c)$ are shown in figure 2.11. Small changes in parameters do not have a dramatic effect on the performance of the model.

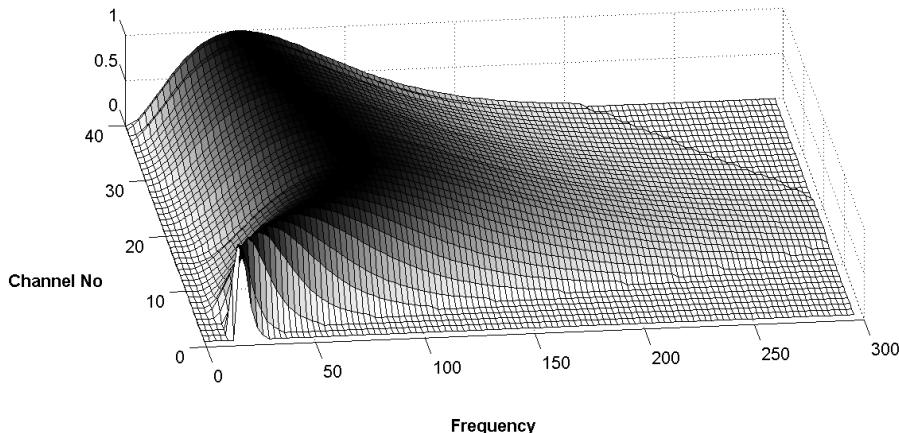


Figure 2.11: Filters become more narrow in the lower auditory channels.

$B(t, f, c)$ represents the *spectrum* of the neural synchronization to the beating frequencies in channel c . The *synchronization index* of the beating frequencies is then defined as the normalized magnitude:

$$I(t, f, c) = \left| \frac{B(t, f, c)}{D(t, 0, c)} \right| \quad (2.6)$$

where $I(t, f, c)$ is the normalized magnitude in channel c . $D(t, 0, c)$ is the DC-component of the whole signal in channel c .

The short-term '*energy*' spectrum of the neural synchronization to beating frequencies in a particular auditory channel c is defined as:

$$\mathbf{B}(t, f, c) = I(t, f, c)^{1.6} = \left| \frac{B(t, f, c)}{D(t, 0, c)} \right|^{1.6} \quad (2.7)$$

We then define the following relationships for the calculation of roughness:

$$\begin{aligned} R_{\mathbf{B}}(t) &= \int \mathbf{B}(t, f) df = \int \sum_{c=1}^C \mathbf{B}(t, f, c) df \\ &= \sum_{c=1}^C \mathbf{B}(t, c) = \sum_{c=1}^C \int \mathbf{B}(t, f, c) df \end{aligned} \quad (2.8)$$

where $R_{\mathbf{B}}(t)$ is the roughness at time t . The roughness is obtained by an integration of the '*energy*' over all frequencies, as well as over all channels. This definition implies a proper visualization, one along the axis of auditory channels and one along the axis of the (beating) frequencies.

2.3.4 Implementation

IPEMRoughnessFFT - Calculates roughness using synchronization index model

2.3.5 Examples

The function **IPEMRoughnessFFT** has a number of input parameters. To calculate the roughness of the wav-file *SchumannKurioseGeschichte.wav* do:

```
[ANI,ANIFreq,ANIFilterFreqs] = ...
IPEMCalcANIFromFile('SchumannKurioseGeschichte.wav');

IPEMRoughnessFFT(ANI,ANIFreq,ANIFilterFreqs,5,300,0.200,0.020,1);
```

Notice that the toolbox functions are linked through the variables **ANI**, **ANIFreq**, **ANIFilterFreqs**. These variables are the output of the Auditory Peripheral Module, and (part of) the input to the Roughness Module. The resulting figure should be similar to the one shown in figure 2.12.

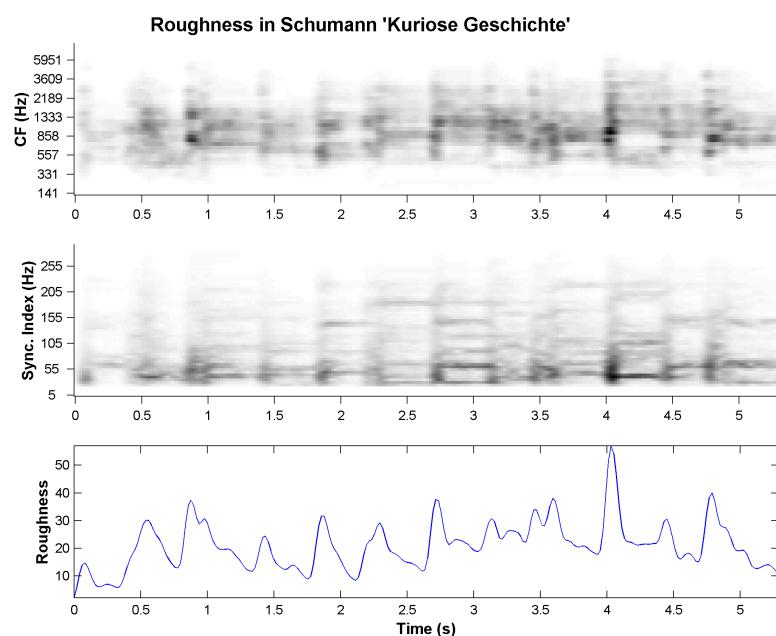


Figure 2.12: The upper panel shows the energy as distributed over the auditory channels, the middle panel shows the energy as distributed over the beating frequencies, the lower panel shows the roughness (which is the sum of either the upper or middle panel)

2.4 Onset Module

2.4.1 Introductory description

The onset module (OM) finds the onsets of sound events. The module takes sound as input and produces the moments where a new note, chord or sound event is triggered in the musical signal. Its place in the image transformation chart is shown in figure 2.13.

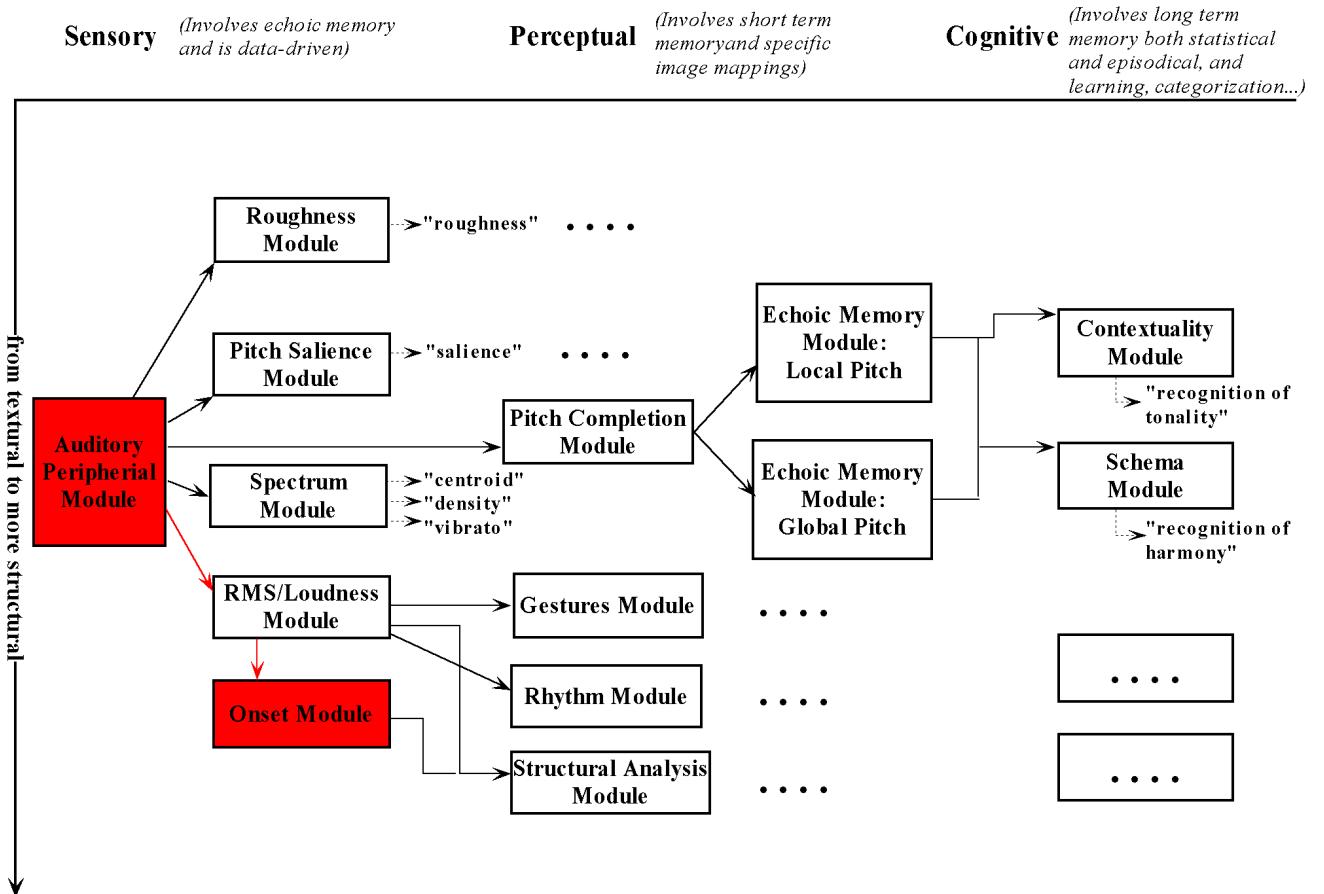


Figure 2.13: Chart of image transformation modules, with OM highlighted

The onset module analyzes the energy in the different auditory channels. It extracts the relevant peaks and combines the results for each channel to an overall onset estimation. Keep in mind that onset detection is a rather difficult enterprise due to the fact that some music instruments (violin, human voice, ...) have gliding onsets which are hard to detect. Our onset module can be used to chunk the frame-based representation into an event-based representation. Figure 2.14 shows the modules involved in the transformation process.

2.4.2 Functional-logical description

The onset module involves:

$$APM : s(t) \rightarrow \tilde{e}(t)$$

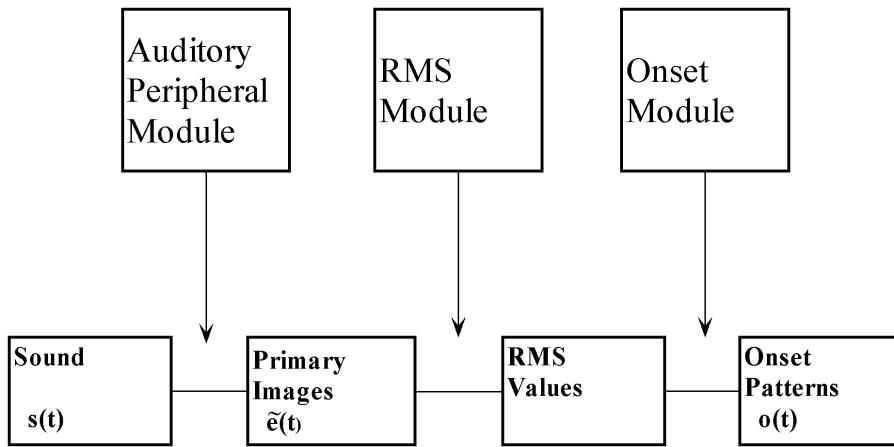


Figure 2.14: Modules involved for onset estimation.

$$RMS : \tilde{e}(t) \rightarrow \widetilde{rms}(t)$$

$$OM : \widetilde{rms}(t) \rightarrow o(t)$$

where APM provides the auditory nerve images and OM transforms these images into the signal $o(t)$. A non-zero value in $o(t)$ corresponds to an onset at that moment, the value itself being proportional to the estimated relevance of the onset.

2.4.3 Signal processing description

The algorithm used for detection of onsets consists of the following steps:

1. Calculation of auditory nerve image.

The auditory nerve image is calculated using the **Auditory Peripheral Module**.

2. Calculation of RMS.

For each of the channels in the ANI, the root mean square values are calculated using a frame size of 0.029 s and a frame step of 0.0058 s. For each frame F in each channel, an RMS value is calculated according to:

$$RMS = \sqrt{\frac{\sum_{i=1}^N F_i^2}{N}}$$

where N = frame width in samples, and F_i is the i-th sample of frame F . There is a toolbox function for the calculation of RMS values, see [IPEMCalcRMS](#).

3. Low pass filtering.

Each of these RMS-averaged channels is filtered with a second-order low pass Butterworth filter with a cutoff frequency of 15 Hz. This is done to reduce small, irrelevant peaks in the RMS signal.

4. Peak detection.

Detection of relevant peaks in each channel is the tricky part.... We start with all true peaks: moments where the first derivative goes from positive to negative. Then, the following approach is used for deciding whether a peak in a channel should be kept as an onset candidate for that channel:

- ignore everything below a certain threshold
- only keep peaks that are significantly bigger than their neighbors
- apply a mask so that following peaks due to irrelevant vibrations are thrown out
- for all remaining peaks, keep the biggest ones within a certain neighborhood, but only if it's big enough compared to the median value at that point

5. Application of integrate-and-fire neural network.

The contribution of this part is to cluster the peaks detected in the previous step over all channels and over time. This algorithm is based on an article by [L. S. Smith \(1996\)](#). One integrate-and-fire neuron is used per channel. Each neuron receives input from a set of adjacent channels, accumulates its input over time and its output is fed back to a number of adjacent neurons. The neuron dynamics are given by:

$$\frac{dA}{dt} = I(t) - diss \times A \quad \text{with: } \begin{aligned} A &= \text{neurons' accumulated value} \\ I &= \text{input} \\ diss &= \text{dissipation factor} \end{aligned}$$

If a neurons' accumulated value exceeds a certain threshold, it fires and the accumulated value is reset to zero. After firing, the neuron becomes insensitive to input for some period, called the refractory period. This configuration produces a sharp bursts of spikes when a number of channels provide evidence for onsets.

6. Final filtering.

The onset pattern is finally filtered using the following approach: an onset is detected at a certain moment

- if a certain number of channels detected an onset within a specific period of time, and
- if the moment falls at least a minimum period of time behind the last detected onset

The result of these steps is a signal where a non-zero value indicates the presence of an onset. To get an indication for the relevance of a specific onset, we have used the number of channels that detected an onset divided by the total number of channels.

2.4.4 Implementation

- | | |
|--------------------------------|--|
| IPEMCalcOnsets | - Calculates onsets of sound signal |
| IPEMCalcOnsetsFromANI | - Calculates onsets starting from an ANI |
| IPEMDOnsets | - Calculates onsets of sound file |
| IPEMOnsetPattern | - Integrate-and-fire neural net for onset detection |
| IPEMOnsetPatternFilter | - Filters onset detection peak pattern |
| IPEMOnsetPeakDetection | - Detects possible onset peaks in multi-channel signal |
| IPEMOnsetPeakDetection1Channel | - Detects possible onset peaks in single channel |
| IPEMSnipSoundFileAtOnsets | - Cuts sound file in segments at onsets |

2.4.5 Examples

To extract the onsets from Schumann's Kuriose Geschichte 7, use the following expression:

```
[Ts, Tsmp] = IPEMDOnsets('SchumannKurioseGeschichte.wav');
```

You can of course split the module in the several subparts. For example, if you first process the file with the Auditory Peripheral Module, and then apply the onsets module, you could use:

```
[ANI,ANIFreq,ANIFilterFreqs] = IPEMCalcANIFromFile('SchumannKurioseGeschichte.wav');
[OnsetSignal,OnsetFreq] = IPEMCalcOnsetsFromANI(ANI,ANIFreq);
```

Figures 2.15 to 2.20 illustrate the different steps for this particular example.

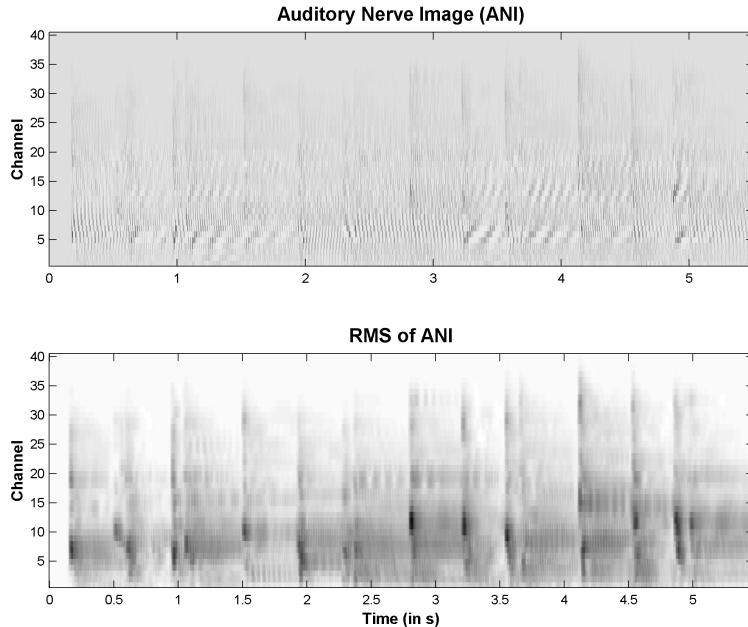


Figure 2.15: Top: ANI of the Schumann sound excerpt. Bottom: RMS values of the ANI.

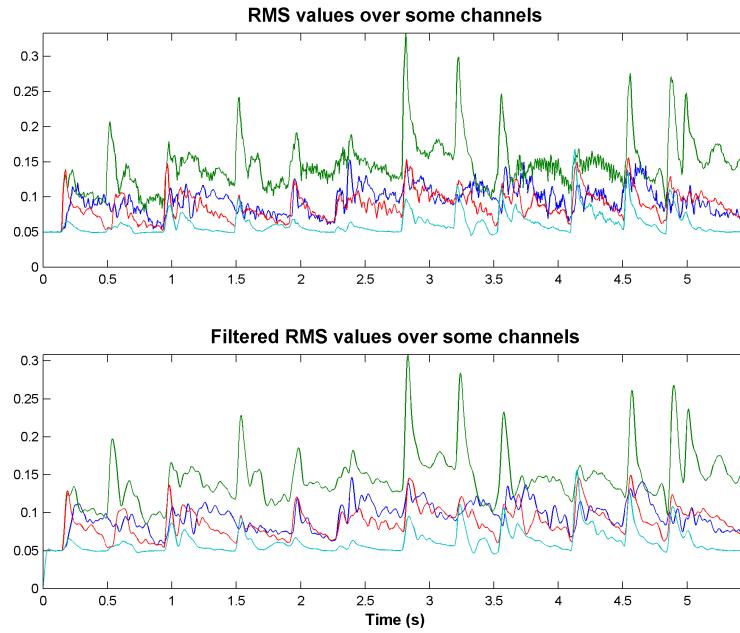


Figure 2.16: Top: RMS values for some channels of the ANI. Bottom: low pass filtered RMS values for the same channels.

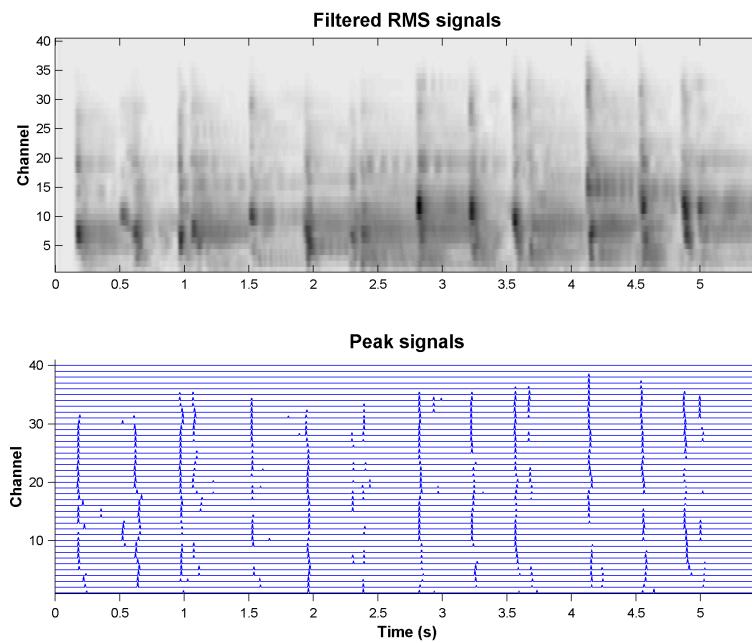


Figure 2.17: Top: filtered RMS values of the ANI. Bottom: peaks detected for each channel.

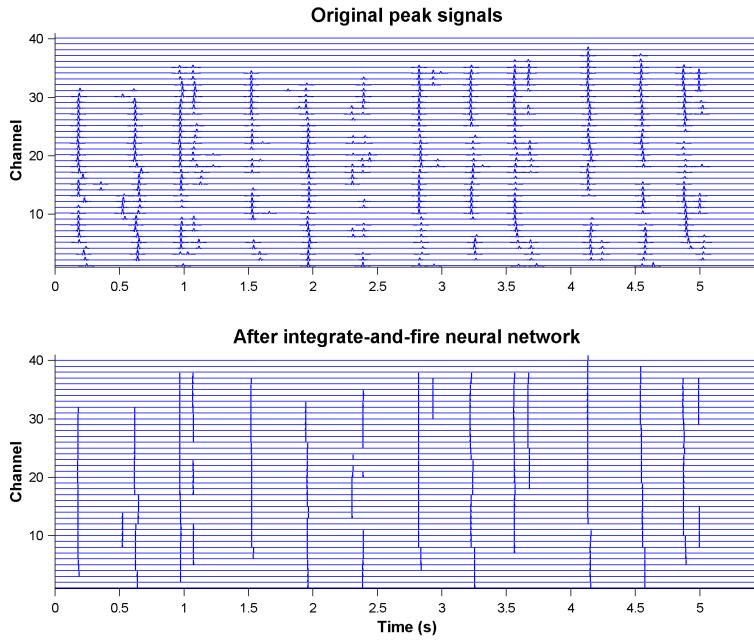


Figure 2.18: Top: detected peaks. Bottom: peaks after applying the integrate-and-fire neural network.

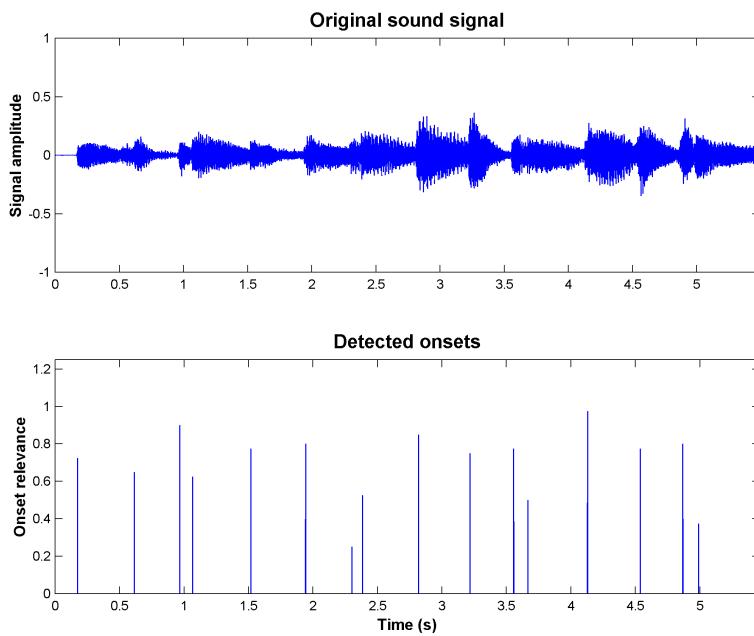


Figure 2.19: Top: original sound signal. Bottom: detected onsets and their relevance level.

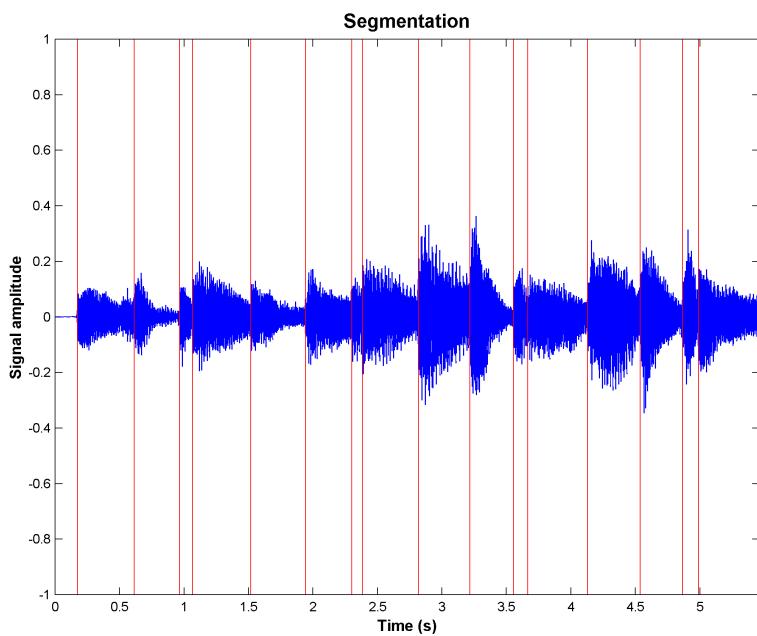


Figure 2.20: Segmentation of the original sound signal using the onset module.

2.5 Pitch Completion Module

2.5.1 Introductory description

The Pitch Completion Module (PCM) calculates the periodicity pitch image of a signal. Strictly speaking, the neural rate code of the auditory nerve images is taken as input and a periodicity analysis is the output. Otherwise stated, the inputs are primary images and the outputs are pitch images. In our global chart of image transformation modules PCM is localized in the section on perception (fig. 2.21).

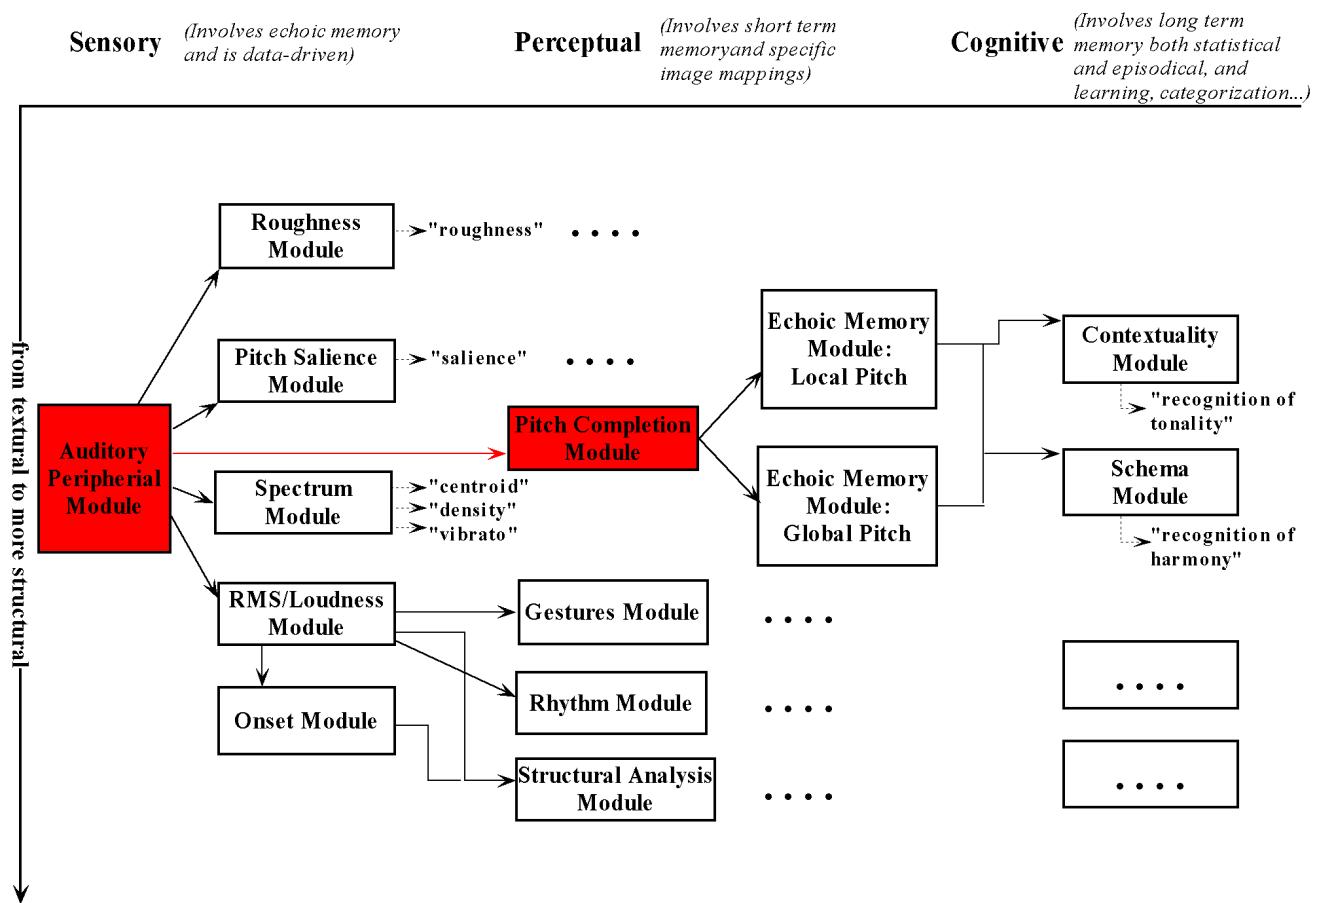


Figure 2.21: Chart of image transformation modules, with PCM highlighted

A periodicity analysis of primary images is dealt with at different occasions, such as in **rhythm**. The main difference is the focus of attention. In rhythm periodicity we look at larger periodicities than in pitch where the focus of our attention is between 80 and 1250 Hz.

- The lower limit of 80 Hz accounts for the fact that for smaller frequencies, the sensation of pitch becomes more a sensation of textural properties. This is a shift of perceptual categorization which is taken into account in our modeling. Indeed, our model of **roughness** took into account a range of frequencies between 5 and 300 Hz but the focus was on frequencies between 50 and 70 Hz (see the filter characteristics in figure 2.11).

- The higher limit of 1250 Hz is related to the limits of neural synchronization. Beyond about 1250 Hz, the neurons are no longer able to follow the exact period of the signal very accurately, and periodicity pitch becomes unreliable.

Pitch estimations at higher frequencies could rely on spectral information contained in $e(t, c)$. Such spectral images could be obtained by taking the RMS-values in each channel over short time periods using **IPEMCalcRMS**. However, the pitch information necessary to deal with harmonic relationships, is here assumed to be contained in the time-code of the auditory nerve images, and this within a frequency band of 80 to 1250 Hz.

From a conceptual point of view the Pitch Completion Module performs a transformation from time-code to place-code, i.e. from patterns whose frequency is contained within the temporal characteristics of the pattern to patterns whose frequency is encoded along a spatial array. Physiological evidence for periodicity pitch perception has been gathered by [Langner \(1992, 1997\)](#). Figure 2.22 shows the modules involved in the image transformation process.

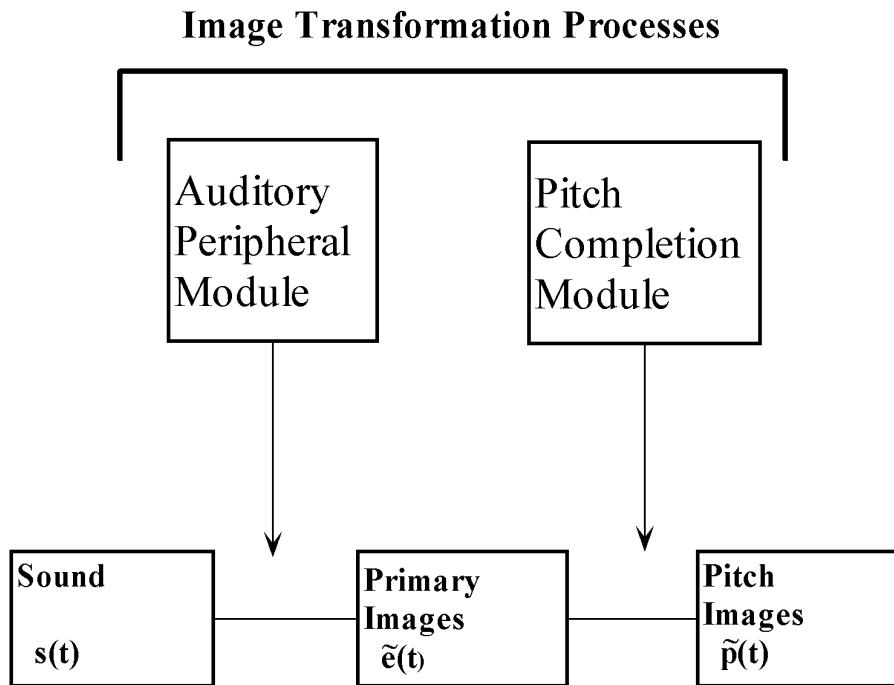


Figure 2.22: Image Transformation Process: Pitch Completion Module

2.5.2 Functional-logical description

First the auditory nerve images are calculated using the **Auditory Peripheral Module**. The Pitch Completion Module then transforms these images into pitch images.

$$APM : s(t) \rightarrow e(t, c)$$

$$PCM : e(t, c) \rightarrow p(t, \tau) (= \tilde{p}(t))$$

The Pitch Completion Module has two stages. First a periodicity analysis of each $e(t, c)$ pattern, band-pass filtered for 80-1250 Hz, is made, where $f[e(t, c)]$ is the band-pass filtered auditory nerve pattern for channel c , and $p(t, c, \tau)$ or $\tilde{p}(t, c)$ the periodicity analysis of that pattern, with τ denoting the period. The detected periods are registered along a spatial array of periods. We thus have:

$$PCM1 : \begin{bmatrix} f[e(t, 1)] \\ f[e(t, 2)] \\ \dots \\ f[e(t, C)] \end{bmatrix} \rightarrow \begin{bmatrix} \tilde{p}(t, 1) \\ \tilde{p}(t, 2) \\ \dots \\ \tilde{p}(t, C) \end{bmatrix}$$

Secondly, a *coincidence mechanism* sums up the $\tilde{p}(t, c)$ over all channels patterns and stores the result in a summary auto-correlation pattern called $\tilde{p}(t)$:

$$PCM2 : \begin{bmatrix} \tilde{p}(t, 1) \\ \tilde{p}(t, 2) \\ \dots \\ \tilde{p}(t, C) \end{bmatrix} \rightarrow \tilde{p}(t)$$

The resulting pattern \tilde{p} (in alternative notation $p(\tau)$) at time t is called the pitch image or *completion image* at time t . The pitch is represented along the τ axis. The completion image is related to the notion of virtual pitch pattern. It gives an account of the common periodicity along the auditory neurons in the frequency region of 80 Hz to 1250 Hz. Notice than the procedure outlined for PCM is very similar to the procedure used for the detection of periodicity in rhythm patterns, as described in the Example section of the [Rhythm Module](#).

2.5.3 Signal processing description

These are the steps of the periodicity pitch calculation:

- Filtering

The channels of the ANI are filtered between 80 and 1250 Hz. Due to the fact that the output of the auditory model gives the envelopes of the neural firing probabilities (< 1250 Hz), it suffices to first apply a low-pass filter and then subtract that from the original signal in order to obtain the pitch. The low-pass filter is a second order Butterworth filter with a cutoff frequency of 80 Hz.

$$\tilde{e}_{filt}(t) = f[\tilde{e}(t)]$$

- Auto-correlation

A frame-based auto-correlation analysis is performed on the filtered channels. A frame width FW and step size FS are chosen and then for each frame and each channel c , we perform an auto-correlation for each time-lag δ :

$$p(t, c, \delta) = \int_t^{t+FW} e_{filt,c}(\tau).e_{filt,c}(\tau + \delta).d\tau \quad \text{where } \delta \in [0, FW]$$

and, by combining the results for all time lags δ in one vector, this gives us $\tilde{p}(t, c)$

- Coincidence mechanism

Summation of the auto-correlation results over all channels:

$$\tilde{p}(t) = \sum_{n=1}^C \tilde{p}(t, c)$$

2.5.4 Implementation

IPEMPeriodicityPitch - Calculates periodicity pitch from nerve image

2.5.5 Examples

Figure 2.23 shows the results of first processing our familiar excerpt of Schumann's Kuriose Geschichte 2 with the APM and then processing the primary image with the Pitch Completion Module. The MATLAB code is:

```
[ANI,ANIFreq,ANIFilterFreqs] = ...
IPEMCalcANIFromFile('SchumannKurioseGeschichte.wav', [], [], 1);
[PP,PPFreq,PPPeriods,PPFANI] = IPEMPeriodicityPitch(ANI,ANIFreq,[],[],[],1);
```

where PPFANI are the filtered (between 80-1250 Hz) auditory nerve images.

Figure 2.24 shows the results of first processing the C major chord using Shepard tones 2 with the APM and then processing the primary image with the Pitch Completion Module.

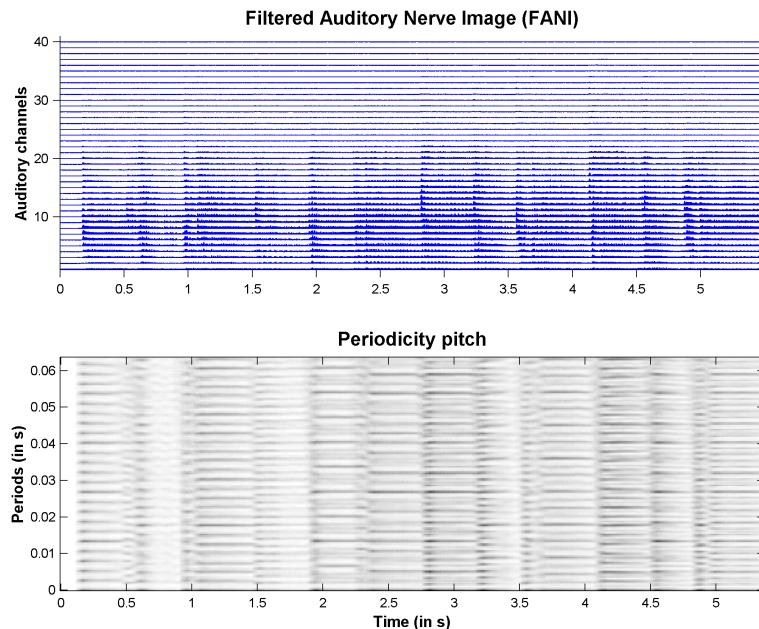


Figure 2.23: Periodicity Pitch for the excerpt of Schumann's Kuriose Geschichte

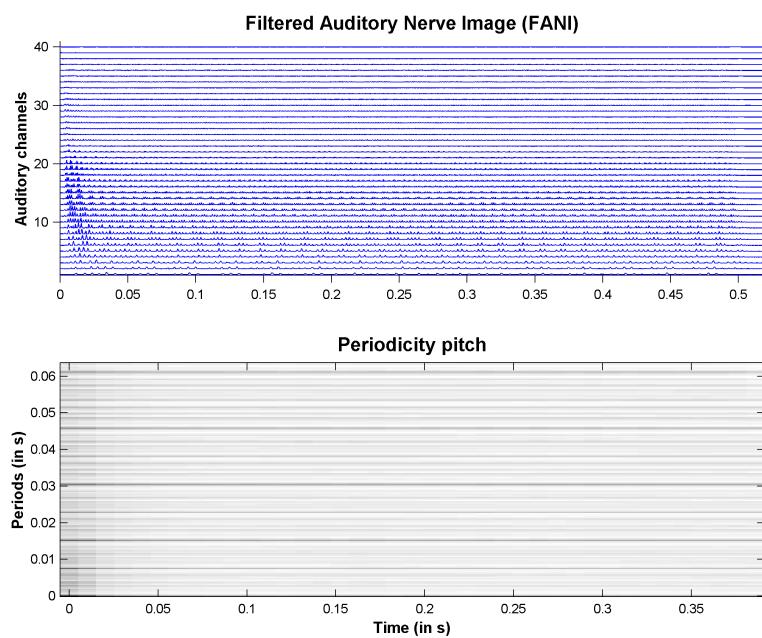


Figure 2.24: Periodicity Pitch for the C major chord using Shepard tones

2.6 Rhythm Module

2.6.1 Introductory description

The Rhythm Module (RhM) uses the Minimal Energy Change (MEC) algorithm to calculate the fundamental period of a signal. The input is a sound file and the output is an estimation of the fundamental period at each time step. The technique used is a generalization of the Average Magnitude Difference Function, known as AMDF (Leman & Verbeke, 2000). The basic idea is that:

- the energy calculated over the period of a repeating pattern is more or less the same at each moment in time
- minimal changes of this energy point to the period of the repetitive pattern.

In applying MEC to rhythm detection we perform the analysis on the energy patterns in the auditory nerve images. This module then has its place in the image transformation chart as shown in figure 2.25.

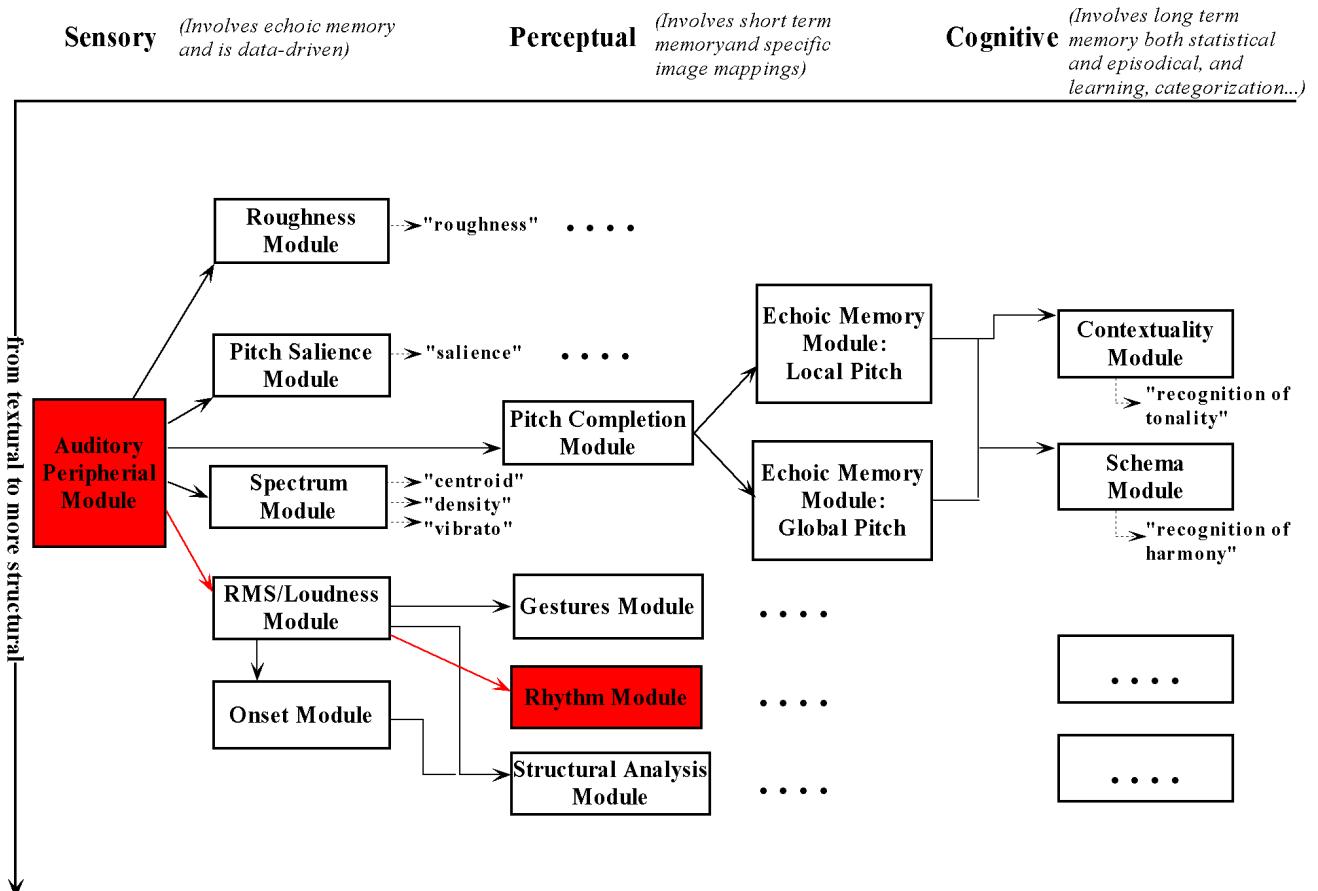


Figure 2.25: Chart of image transformation modules, with RhM highlighted

2.6.2 Functional-logical description

The Minimal Energy Change algorithm applied to the recognition of repetition in rhythm involves:

$$APM : s(t) \rightarrow e(t, c) \quad (2.9)$$

$$RMS : e(t, c) \rightarrow r(t, c) \quad (2.10)$$

where APM provides the auditory nerve images and where RMS considers the energy in these images. The application of MEC involves:

$$MEC1 : r(t, c) \rightarrow m(\tau, t, c) \quad (2.11)$$

where MEC1 performs a periodicity analysis of these energies and outputs the estimated periodicity patterns $m(\tau, t, c)$, where t is time as usual, τ is the period, and c is the auditory channel. To obtain the summary MEC-analysis, one has to sum over all auditory channels, which gives:

$$MEC2 : m(\tau, t) = \sum_{c=1}^C m(\tau, t, c) = \tilde{m}(t) \quad (2.12)$$

Hence, the MEC-function can be summarized as:

$$MEC : \tilde{r}(t) \rightarrow \tilde{m}(t) \quad (2.13)$$

In this notation, one should take into account that the tilde character, indicating a vector representation, applies to different dimensions. The components of $\tilde{r}(t)$ are the channels, while the components of $\tilde{m}(t)$ are the different periods considered in the analysis. A derived value is the so-called "best period" $p(t)$, which we obtain by taking the minimum in each $\tilde{m}(t)$ for fixed t , hence:

$$MEC : \tilde{r}(t) \rightarrow \tilde{m}(t) \rightarrow p(t) \quad (2.14)$$

2.6.3 Signal processing description

The MEC-algorithm takes a signal $s(t)$ and its shifted version $s(t - \tau)$. It then calculates the difference between both signals, takes the absolute value, and integrates the result:

$$a(\tau, t) = \int_0^t |s(t' - \tau) - s(t')| dt' \quad (2.15)$$

where $a(\tau, t)$ at fixed t contains an analysis for all τ . τ is typically chosen in the region of interest. For rhythm detection, this may be from 400 ms to 1 s. The obtained analysis $a(\tau, t)$ at fixed t is then searched for a minimum:

$$p(t) = \min_{\tau > 0} a(\tau, t) \quad (2.16)$$

which gives a value p for the best period at each time step t .

In practice, we can build in more stability by leaky-integrating the $a(\tau, t)$ over t :

$$a(\tau, t) = \int_0^t |s(t' - \tau) - s(t')| e^{\beta(t-t')} dt' \quad (2.17)$$

and apply Expression 2.16.

When applied to the signal in different auditory channels, the calculation will be applied to each channel, thus giving $a(\tau, t, c)$. To get the best period, one may calculate $p(t)$ using $\sum_{c=1}^C a(\tau, t, c)$.

2.6.4 Implementation

[IPEMMECAAnalysis](#)

- Performs a periodicity analysis of a (multi-channel) signal using the MEC model

[IPEMMECEExtractPatterns](#)

- Extracts the best pattern from the original signal using the results of an IPEMMECAAnalysis run

[IPEMMECReSynthUI](#)

- User interface callback function for interactively handling the resynthesis of MEC analysis results

[IPEMMECSaveResults](#)

- Utility function that saves the results of IPEMMECEExtractPatterns, so that resynthesis can still be done at a later date, after reloading this data

[IPEMMECSynthesis](#)

- Generates an AM modulated noise signal constructed from a repetition of the pattern found at the specified time moment

2.6.5 Examples

The first steps involve the preparation of the signal in which we want to look for periodicity. This is done by:

```
[ANI,ANIFreq,ANIFilterFreqs] = IPEMCALCANIFromFile('SchumannKurioseGeschichte.wav');
[RMS,RMSFreq] = IPEMCALCRMS(ANI,ANIFreq,0.050,0.020);
[Periods,Best,AnalysisFreq,Values] = IPEMMECAAnalysis(RMS,RMSFreq,0.5,3,[],1.5);
```

The MEC algorithm has stored the periods that were analyzed in `Periods`, the indices (into `Periods`) for the best periods found in each channel in `Best`, and all values $m(\tau, t, c)$ in `Values`. The data format of `Values` is described in [IPEMMECAAnalysis](#). Just to give you an example of how to work with MATLAB, we sum over all channels c in the following way:

```
SummaryPeriodicity = zeros(size(Values{1}));
for channel = 1:40;
SummaryPeriodicity = SummaryPeriodicity + Values{channel};
end
```

The values in `SummaryPeriodicity` contain the periodicity analysis over all channels which can be visualized using the MATLAB code:

```
figure;
imagesc((0:size(SummaryPeriodicity,2)-1)/AnalysisFreq,Periods,SummaryPeriodicity);
xlabel('Time (in s)'); ylabel('Period (in s)');
axis xy; colormap(1-gray);
```

We plot the best periods of the summary periodicity on top of the image.

```
[M,I]=min(SummaryPeriodicity); hold on;
plot((0:length(I)-1)/AnalysisFreq,Periods(I));
```

The resulting graph should be similar to the one in figure 2.26.

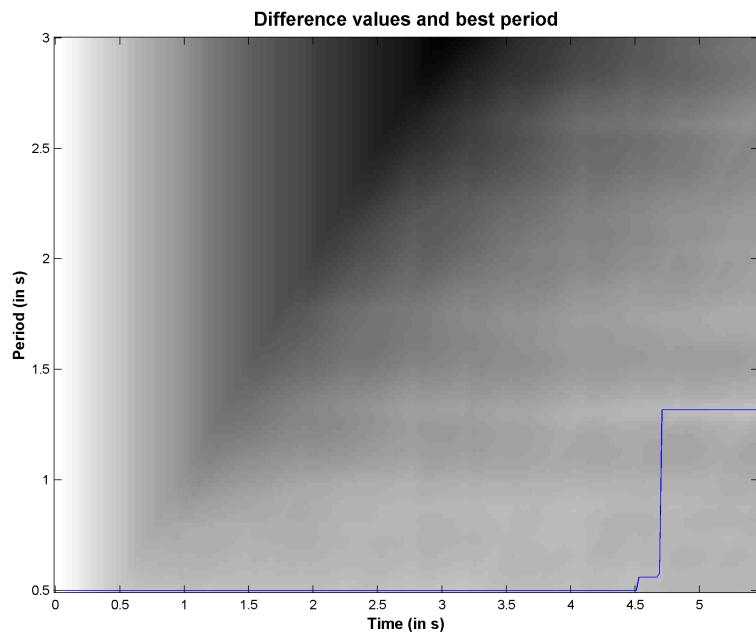


Figure 2.26: MEC analysis result of a short excerpt of Schumann's Kuriose Geschichte. Shown are the summed difference values (over all channels) and a plot of the best period on top of this.

2.7 Echoic Memory Module

2.7.1 Introductory description

An echo can be defined as a leaky integration with a given half decay time. At each moment in time, a leaky integrator adds the incoming signal value to an attenuated version of the previously calculated value to form the newly calculated value. The half decay time specifies the time it takes for an impulse signal to reach half its original value.

The Echoic Memory Module (EMM) takes an image as input and gives the leaky integrated image as output. The images are integrated so that at each time step, the new image is calculated by taking a certain amount of the old image which is then added with the new incoming image.

The Echoic Memory Module can be applied to pitch completion images, for example. We then start from our familiar APM, followed by PCM, and apply EMM. The echo which we specify defines the amount of context which we take into account. With very little (or almost no) context we speak about *local pitch images*. When context is taken into account we use a longer half decay time and call the images *global pitch images*. EMM has been used in [Leman \(2000a\)](#) to construct the pitch images of an echoic memory. In our global chart of image transformation modules EMM is localized in the middle part (perception section) of figure 2.27. The application is described in detail in the chapter on [tonality induction experiments](#). Figure 2.28 shows the modules involved

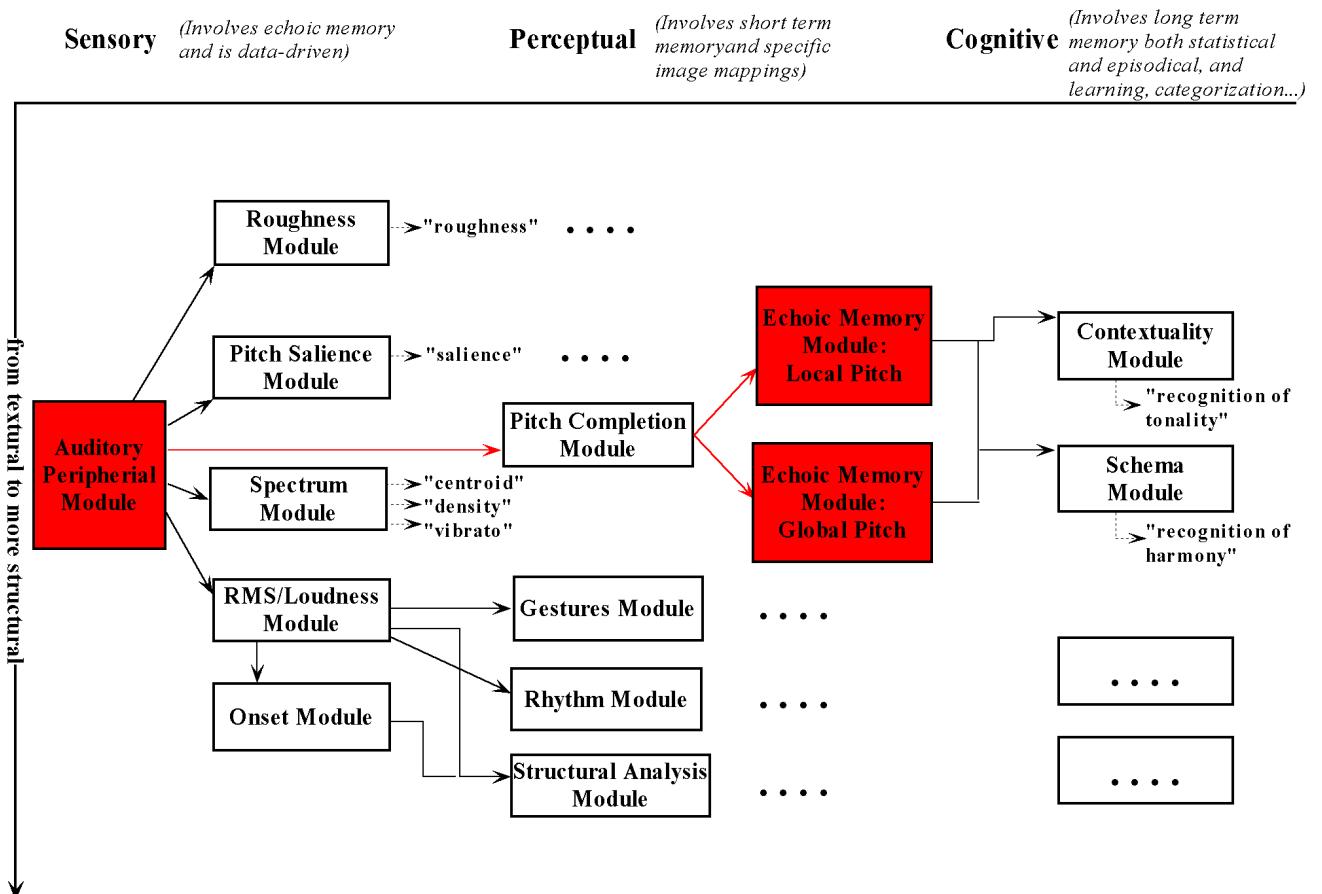


Figure 2.27: Chart of image transformation modules, with EMM highlighted

in the image transformation process.

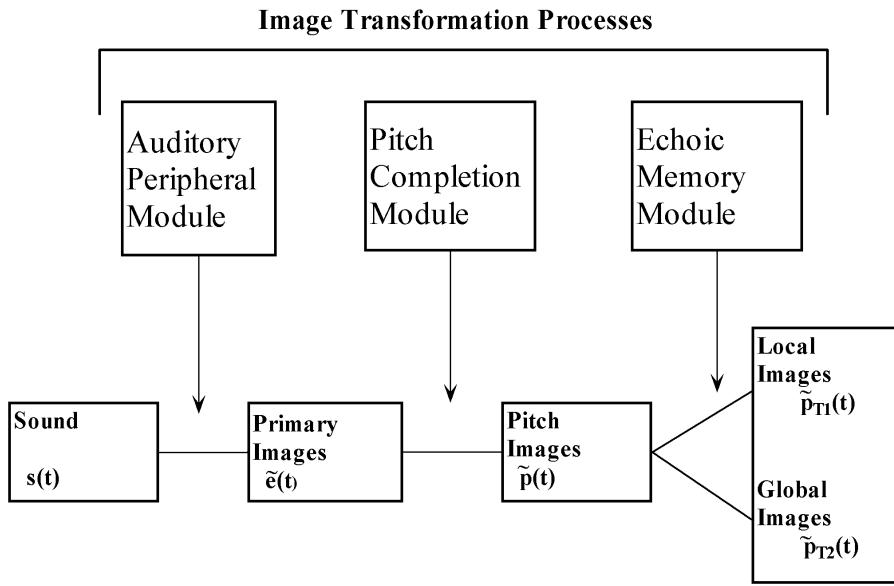


Figure 2.28: Image Transformation Process: Echoic Memory Module

2.7.2 Functional-logical description

A symbolic representation of the Echoic Memory Module may be written as:

$$EMM : \tilde{p}(t) \rightarrow \tilde{p}_T(t) \quad (2.18)$$

where T denotes the echo in seconds, $\tilde{p}(t)$ denotes a pitch completion image and $\tilde{p}_T(t)$ an echoed image.

2.7.3 Signal processing description

$$\begin{aligned} \tilde{p}_T(0) &= \tilde{p}(0) \\ \tilde{p}_T(t) &= \tilde{p}(t) + \tilde{p}_T(t-1) * 2^{-\frac{1}{T}} \quad , t \neq 0 \end{aligned}$$

2.7.4 Implementation

IPEMLeakyIntegration - Calculates leaky integration with specified half decay time

2.7.5 Examples

The Echoic Memory Module can be applied to a pitch completion image to obtain a local and a global pitch image. A local pitch image can be defined as $\tilde{p}_{T=0.1}(t)$, with an echo of 0.1 s, while a global pitch image can be defined as $\tilde{p}_{T=1.5}(t)$, with the echo of 1.5 s. Local images are built up by echoes which, by definition, are smaller than echoes for the global images. Use the following

MATLAB functions to create those pitch images:

```
[ANI,ANIFreq,ANIFilterFreqs] = IPEMCalcANIFromFile('SchumannKurioseGeschichte.wav');
[PP,PPFreq,PPPeriods,PPFANI] = IPEMPeriodicityPitch(ANI,ANIFreq);
LocalIntegration = IPEMLeakyIntegration(PP,PPFreq,0.1,0.1,1);
GlobalIntegration = IPEMLeakyIntegration(PP,PPFreq,1.5,1.5,1);
```

Figure 2.29 shows the pitch images for the excerpt of Schumann's Kuriose Geschichte ②.

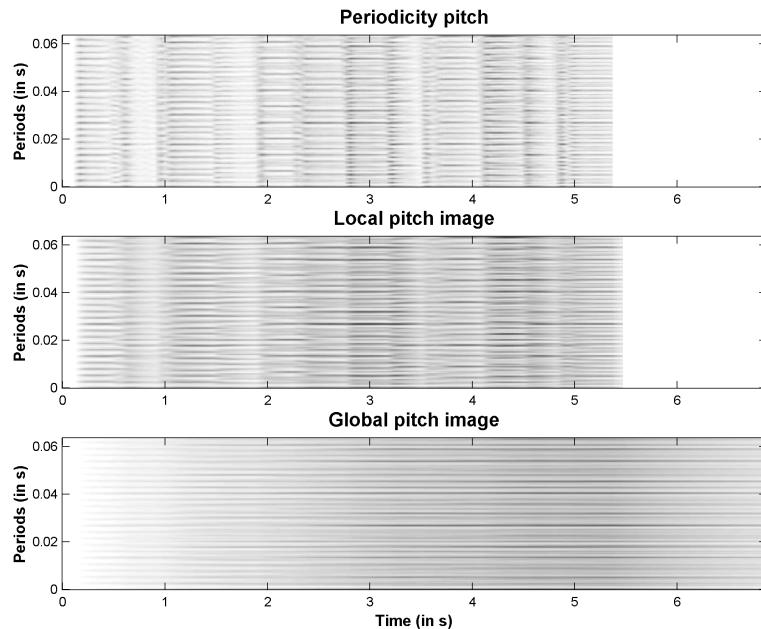


Figure 2.29: Top: periodicity pitch for the excerpt of Schumann's Kuriose Geschichte Middle: local pitch image of Schumann's Kuriose Geschichte (echo of 0.1 s). Bottom: global pitch image (echo of 1.5 s)

2.8 Contextuality Module

2.8.1 Introductory description

Contextuality measures the pitch commonality between two running pitch images of the same sound, each having a possible different echo. In Leman (2000a), contextuality is used to measure the pitch commonality of local (=short echo) images versus global (=long echo) images. Two types of measurement are taken into account:

- The first method is based on an *inspection* of the pitch sequence by means of a fixed image (or *probe*).
- The second method is based on a *comparison* of the pitch images (each having a possible different echo) at running time. This amounts to the calculation of the differences, due to the echo, between both running images.

Inspection and comparison of echoic pitch images are useful methods for studying tonal tensions in terms of figure/ground relationships.

In our global chart of image transformation modules CM is localized in the cognitive section of figure 2.30.

Figure 2.31 shows the modules involved in the image transformation and inference process.

2.8.2 Functional-logical description

Contextuality realizes a mapping of two images onto a scale:

$$\text{Contextuality} : \tilde{p}_{T1} \otimes \tilde{p}_{T2} \rightarrow [-1...1] \quad (2.19)$$

where \tilde{p} is a pitch image, $T1$ and $T2$ are standing for the respective echoes (with $T1 \leq T2$), and \otimes represents the calculation of a correlation coefficient. The two methods then correspond to the following mappings:

- Inspection:

$$\text{Contextuality}_I : \tilde{p}_{T1}(\tau) \otimes \tilde{p}_{T2}(t) \rightarrow [-1...1]$$

where τ represents a fixed time instance, typically situated at the end of the sound sequence, and t represents the running time.

- Comparison:

$$\text{Contextuality}_{II} : \tilde{p}_{T1}(t) \otimes \tilde{p}_{T2}(t) \rightarrow [-1...1]$$

where t represents the running time for both images.

2.8.3 Signal processing description

The Contextuality Module involves several modules.

- The auditory nerve image is calculated using the **Auditory Peripheral Module**
- The **Pitch Completion Module** transforms the auditory nerve image into pitch images

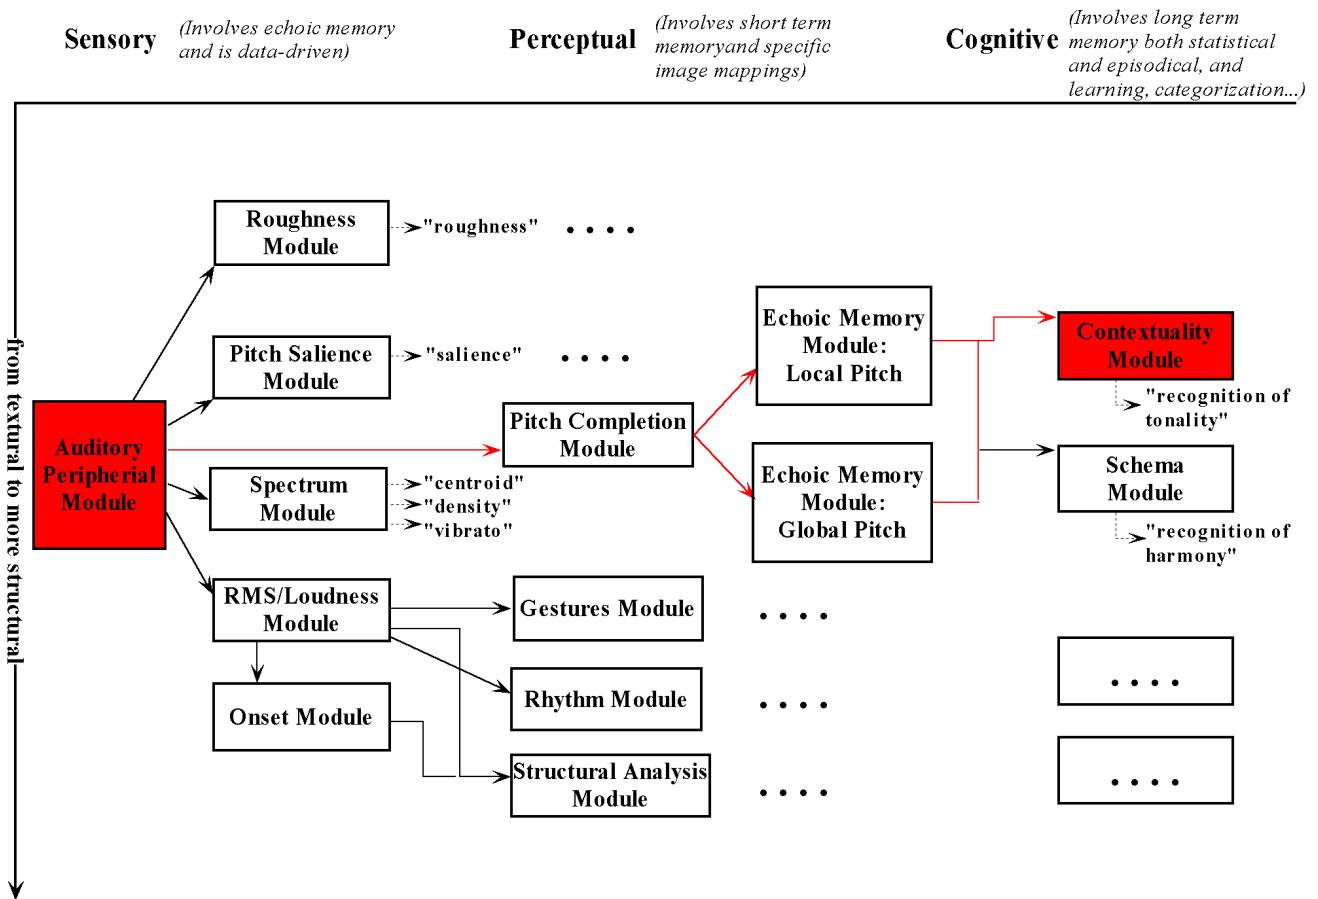


Figure 2.30: Chart of image transformation modules, with CM highlighted

- Through the **Echoic Memory Module** the pitch images undergo the effect of an echo and are processed into local and global images
- The Contextuality Module considers two types of contextuality

2.8.4 Implementation

IPEMContextualityIndex - Calculates the contextuality index. Two methods are used: the method of comparison compares running chords with running tone center images, while the method of inspection compares a fixed chord with running chords and running tone center images.

2.8.5 Examples

Contextuality has been used to simulate the probe-tone experiments of Krumhansl and Kessler (Leman, 2000a; Krumhansl & Kessler, 1982).

In what follows, contextuality is applied to the short excerpt of Schumann's Kuriose Geschichte 2 followed by a 0.1 s period of silence and an f♯ Shepard tone. Use the following MATLAB code

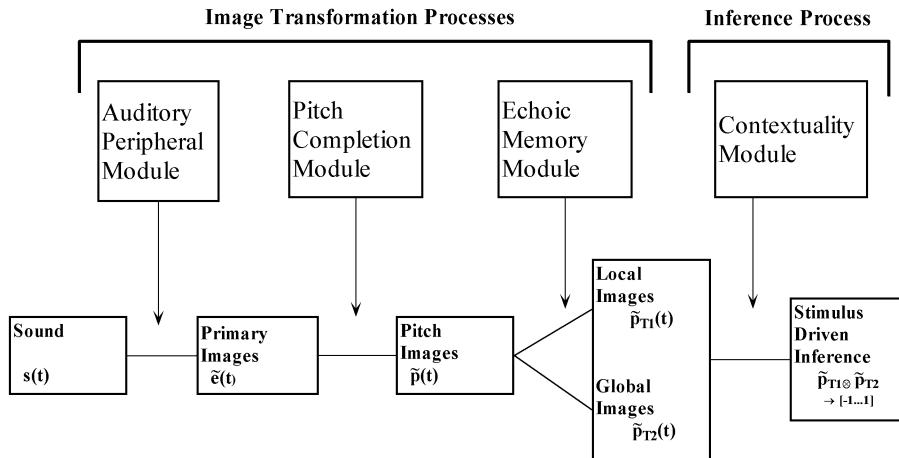


Figure 2.31: Image Transformation and Inference Process: Contextuality Module

to read in the sound file, generate the silence and generate the Shepard tone f♯:

```
[s1,fs] = IPEMReadSoundFile('SchumannKurioseGeschichte.wav');
s2 = zeros(1,round(0.1*fs));
NoteFreq = IPEMCalcNoteFrequency('F#4');
s3 = IPEMShepardTone(NoteFreq,0.5,fs,1,-20);
```

Concatenate the 3 sound signals, calculate the auditory nerve image and from there the periodicity pitch (see also **Echoic Memory Module**) using the following Matlab code:

```
s = [s1 s2 s3];
[ANI,ANIFreq,ANIFilterFreqs] = IPEMCalcANI(s,fs);
[PP,PPFreq,PPPeriods,PPFANI] = IPEMPeriodicityPitch(ANI,ANIFreq,[],[],[],1);
```

Then finally, the function **IPEMContextualityIndex** is applied for the context analysis.

```
[Chords,ToneCenters,LocalInspection,GlobalInspection,LocalGlobalComparison] = ...
IPEMContextualityIndex(PP,PPFreq,PPPeriods,[],0.1,1.5,[],1);
```

In this example, PP contains the periodicity pitch image of the concatenated sounds, PPFreq contains the sampling rate, PPPeriods contains the analyzed periods, and the fourth argument specifies the location of the fixed image. Here, it is left blank, which defaults to the end of the signal. The local and global echoes are set to 0.1 s and 1.5 s, respectively, and the plot flag is set to 1.

Figure 2.32 shows the periodicity pitch image of the Schumann excerpt followed by the Shepard tone, the local pitch image, and the global pitch image.

The method based on inspection has been applied to the local and the global pitch images.

LocalInspection (fig. 2.33) contains the results of inspecting the local images with the fixed

(local) image. The graph shows the similarity of the Shepard probe tone of f♯ within the Schumann excerpt at a local level.

GlobalInspection (fig. 2.34) contains the results of inspecting the global images with the fixed (local) image. The graph shows the similarity of the Shepard probe tone of f♯ within the Schumann excerpt at a more global pitch level.

LocalGlobalComparision (fig. 2.35) contains the results of comparing the local images with the global images over the whole piece. It shows the degree in which the local pitch images (= the level of individual tones and chords) fit with the global pitch images (= the level of the tone context). At points where the local image deviates from the global image, such as in a tonal movement, there will be a low correlation. Inversely, at points where the local image confirms the global image, there will be a high correlation.

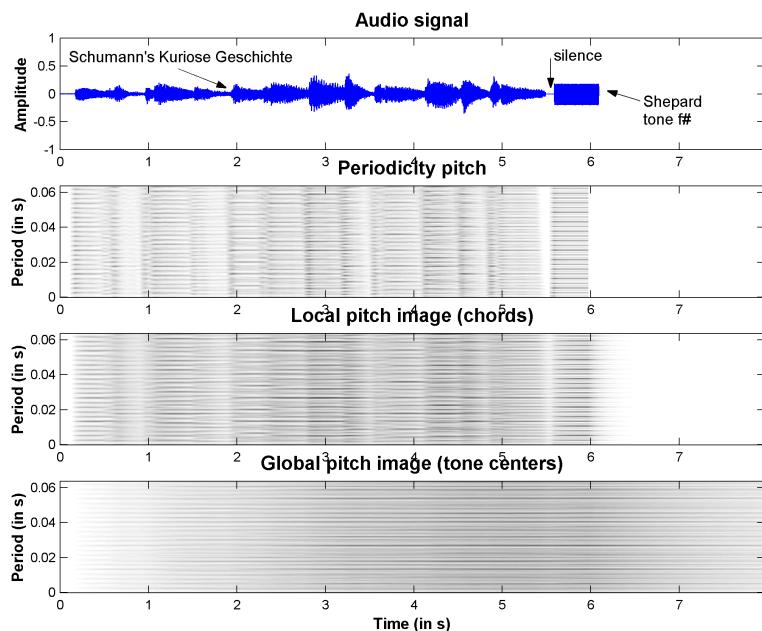


Figure 2.32: Pitch images for the excerpt of Schumann's Kuriose Geschichte followed by a small period of silence and a Shepard probe-tone of f♯. From top to bottom: the analyzed audio signal, the periodicity pitch image, the local pitch image and the global pitch image.

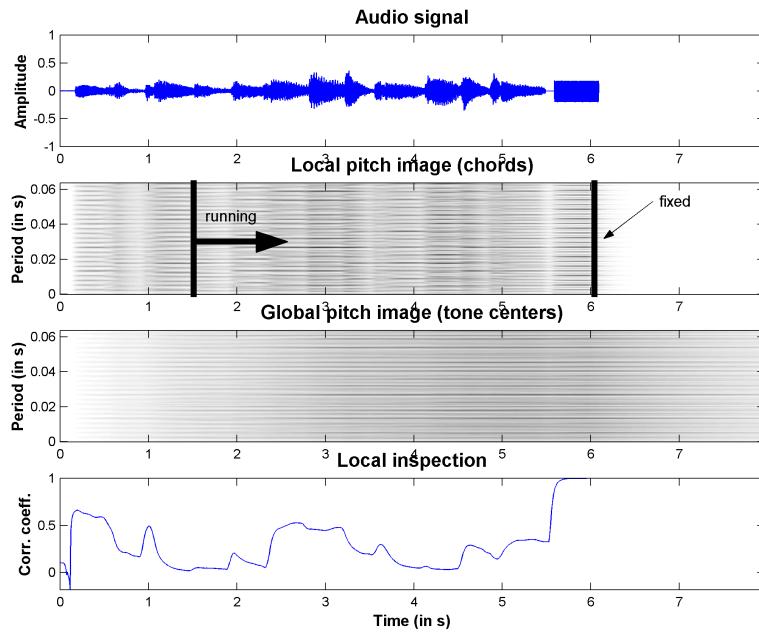


Figure 2.33: Inspection of Schumann's Kuriose Geschichte followed by a small period of silence and a Shepard probe-tone of $f\#$. Inspection of local images. From top to bottom: audio signal, local pitch image, global pitch image and local inspection.

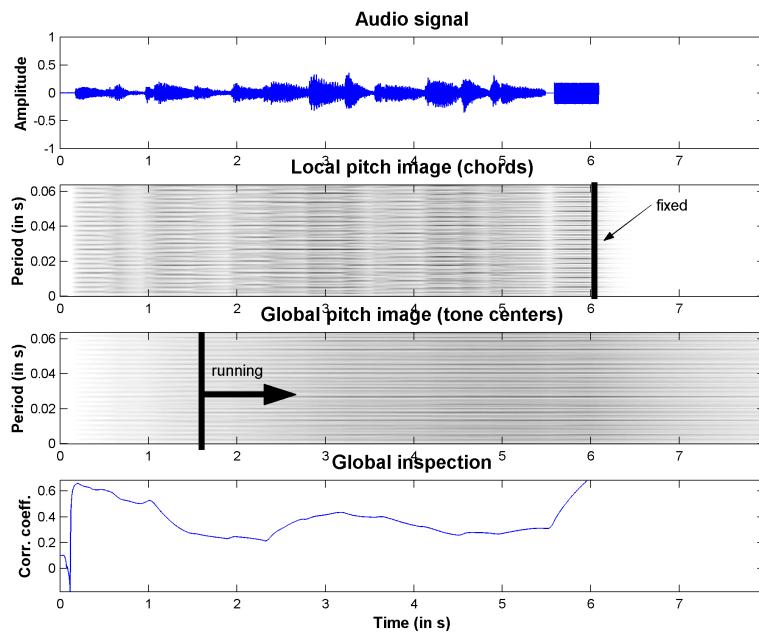


Figure 2.34: Inspection of Schumann's Kuriose Geschichte followed by a small period of silence and a Shepard probe-tone of $f\#$. Inspection of global images. From top to bottom: audio signal, local pitch image, global pitch image and global inspection

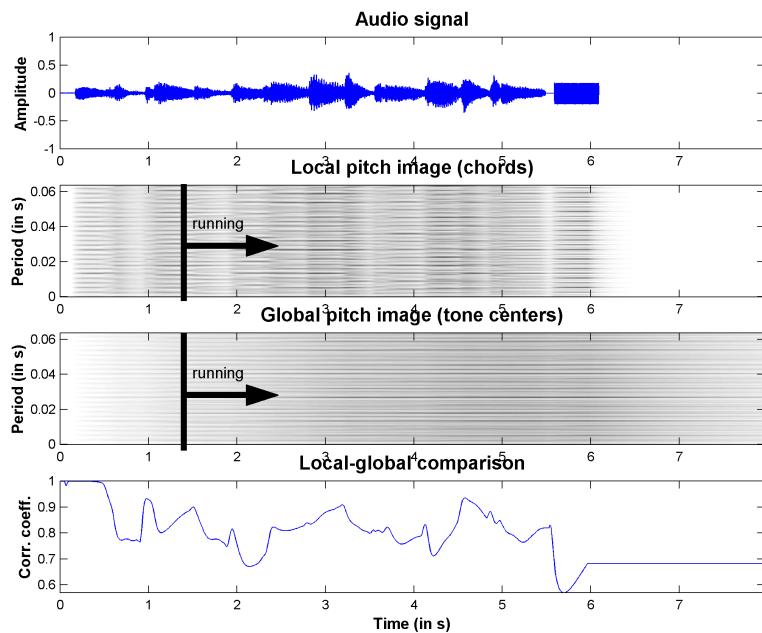


Figure 2.35: Comparison of the local and global pitch images of Schumann's *Kuriose Geschichte* followed by a small period of silence and a Shepard probe-tone of $f\sharp$. From top to bottom: audio signal, local pitch image, global pitch image and comparison between local and global pitch images.

Chapter 3

Experiments

3.1 Introduction

... To Be CompletedDoublecheck this...

Whereas the previous chapter provided a background and description of the basic modules contained in the IPEM Toolbox, this chapter focuses on a comparison of the results obtained with our tools with those obtained from psychoacoustical experiments. Each of the sections in this chapter contains:

- a small introduction in which the purpose of the experiment is explained
- a description of the method that is used
- a summary of how to run the experiment in practice using the toolbox and its functions
- an overview of the obtained results and a discussion thereof

The experiments are explained from a conceptual viewpoint, and each time a link is made between these concepts and the practical tools (MATLAB functions) available in the toolbox.

The following experiments are included:

- Simulations of roughness in comparison with psychoacoustical data.
- Short-term model of tonal induction experiments which apply contextuality to simulate two psychological experiments.

Each experiment can be seen as a comparison of one (or more) module(s) that were handled in the previous chapter with data obtained from experiments in the psychoacoustic research field.

3.2 Roughness experiments

3.2.1 Introduction

The [Synchronization Index Model](#), which forms the kernel of the Roughness Module (RM), is based on the idea that roughness depends on the degree in which neurons synchronize with the beating frequencies. The model provides a straightforward visualization of the components contributing to roughness.

In this demonstration the Roughness Module is compared with behavioral data in psychoacoustics and music psychology. Due to the fact that a lot of sounds are processed, this demonstration consumes time and hard disk memory. The output needs about 253MBytes of your hard disk. Performing the full demonstration takes about 29 minutes on a Windows NT 4.0 system with Pentium III 600 MHz processor, 256Mb RAM PC130.

3.2.2 Method

... To Be Completed ...

3.2.3 Application

Some Practical Considerations

The contents and the functions of the demonstration package for the roughness experiments can be found in the directory IPEM\Demos\Roughness. To perform the experiments execute the main script [IPEMRoughnessDemo](#) for demonstrating the calculation methods of roughness. The function [IPEMGenerateANIForRoughnessDemo](#) generates all the files for the tests contained in the roughness demo:

- [IPEMGenerateANIForRoughnessTest](#) ([IPEMGenerateANIForRoughnessDemo](#) Part I) generates sounds and auditory nerve images for the psychoacoustical tests. The sounds and images are further processed with [IPEMRoughnessDemo](#) (Part II).
- [IPEMGenerateANIForScales](#) ([IPEMGenerateANIForRoughnessDemo](#) Part II) generates an ANI and soundfile of several harmonic and inharmonic sounds for the musical tests.

The output, a number of mat-files containing the required auditory nerve images and the soundfiles, is stored in the directory RoughnessDemo relative to IPEMRootDir ('output'). The function [IPEMRoughnessRun](#) then calculates roughness of all the specified auditory nerve images using different models and parameters.

Psychoacoustical tests

This section compares the model's output with psychoacoustical data. We start with amplitude modulated signals because they are characterized by a few parameters whose effect of roughness has been studied by many authors. We consider the effects of modulation depth (m), amplitude (A), carrier frequency (f_c), and modulation frequency (f_m).

1. **Modulation depth** The dependency of the roughness R on the degree of modulation m has been expressed as a power law $R \sim m^n$. Depending on the experimental setup, n can slightly vary from 1.2 up to about 2 (Terhardt, 1997). Figure (... To Be Completed ... ask Marc) 3.1 shows the roughness of an amplitude modulated tone with f_c 1000 Hz, f_m 70 Hz, 60 dB in function of the modulation depth m linearly increasing from 0 to 1. The

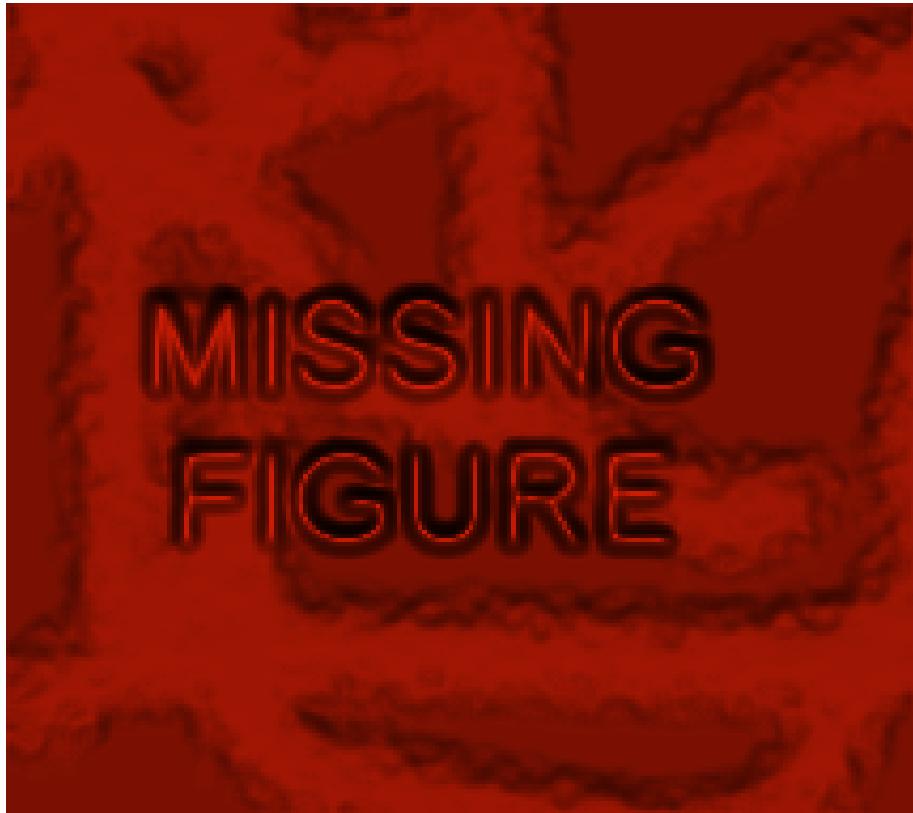


Figure 3.1: Calculated roughness as a function of modulation index $m = 0 \dots 1$ for a 1000 Hz AM tone

dotted line shows the relation $R \sim m^{1.3}$ for $m \leq 1$. The slope of this curve depends on the parameter δ (see Equation ??) which was set to 1.5 to obtain this result. A change of δ has a slight effect on the slope of the curve in Fig. 3.1.

2. **Amplitude** The effect of amplitude is illustrated in Fig. ?? using an AM tone with f_c 1000 Hz, $m = 1$, f_m 50 Hz, and a linear decrease of SPL dB from 70 to 0. The dynamic range of neurons is known to be limited in the order of 30-50 dB. However, the large dynamic range of the human auditory system is explained by the spread of excitation due to the resonance characteristics of the basilar membrane (R. Smith, 1988). The spread of excitation is reflected in figure 3.2a where a broad excitation range can be noticed at a level of 70 dB. This range narrows down as the level decreases. In figure 3.2b, however, one may notice that the excitations contribute to the same frequency of 50 Hz. From 70 to about 53 dB the roughness halves.

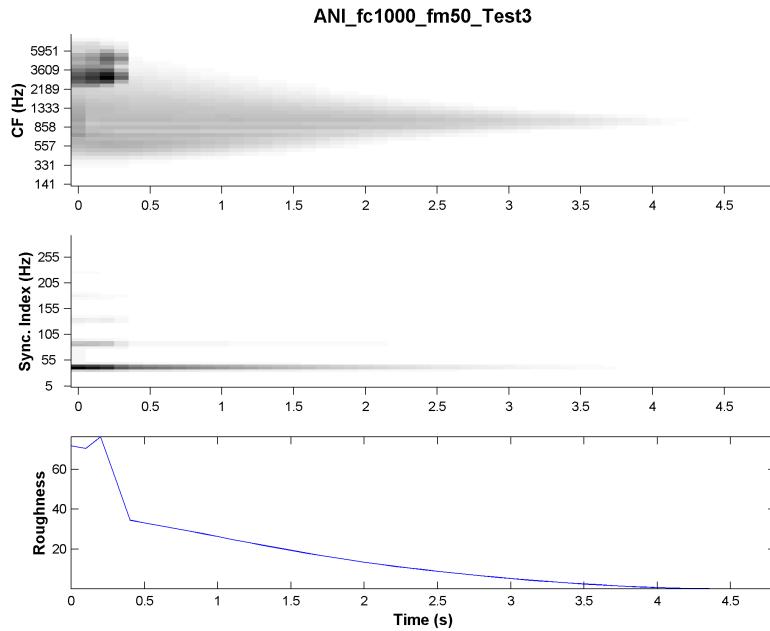


Figure 3.2: Calculated roughness as a function of amplitude.

3. Frequency Modulation The formula used to calculate an FM tone is (Daniel & Weber, 1997):

$$s(t) = \hat{s} \cdot \sin \left(2\pi f_c t - \frac{\Delta f}{f_{mod}} \cdot \cos(2\pi f_{mod} t) \right)$$

where \hat{s} is the effective amplitude. The effects of frequency modulation on roughness show a band-pass characteristic with maximal values from 40 Hz to 70 Hz (Kemp, 1982) figure 3.3, full line), which our model (dotted line) approximates except at f_{mod} 20 Hz, where the estimated roughness is clearly too high. The other values fall within the quartile range. The signal is a 1.6 kHz FM tone at 60 dB with Δf 800 Hz whose roughness is measured at f_{mod} 1 10 20 40 60 80 100 200 300 400 500 Hz.

In addition, the roughness of a FM tone with f_c 1600 Hz, f_{mod} 70 Hz, at 60 dB is measured in function of the frequency deviation Δf at 10 15 30 60 120 160 340 580 880 1130 1400 1580 1720 Hz (figure 3.4). The results correspond rather well with the available psychoacoustical data, as shown in figure 3.3. (All values fall within the quartile range).

Musical tests

The music theoretical interest of the concept of roughness is demonstrated by taking a harmonic tone complex and playing it together with a pitch shifted version thus specifying different musical intervals of the timbre over a defined range. As an illustration, we took a harmonic tone complex consisting of a fundamental (f_0) at 500 Hz and 5 harmonics with equal amplitude. This tone is played together with a pitch shifted copy. The shift over 5 seconds is linear in frequency up to the upper octave (f_0 1000 Hz). The roughness as calculated with the synchronization index model is shown in figure 3.5.

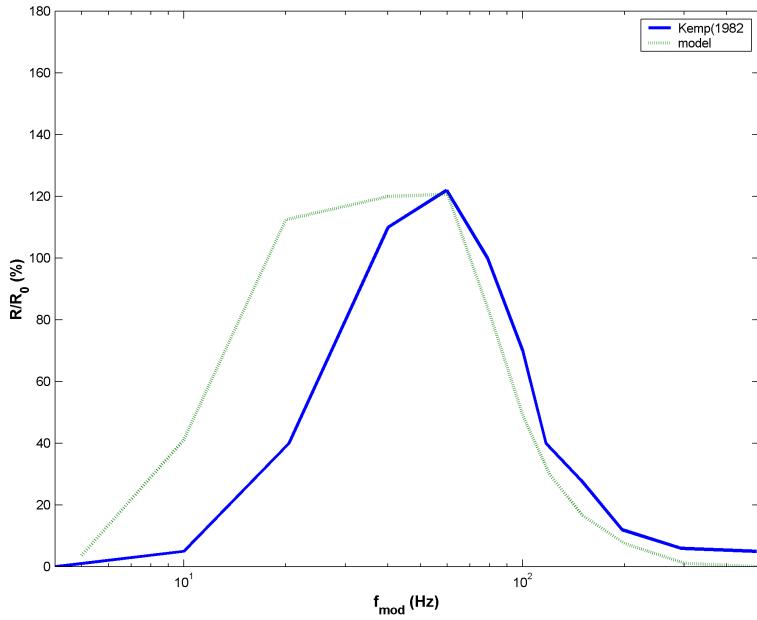


Figure 3.3: Calculated roughness as a function of frequency modulation. The signal is a 1.6 kHz FM tone at 60 dB with $\Delta f = 800$ Hz, and f_{mod} from 1 10 20 40 60 80 100 200 300 400 500 Hz. The peaks are generated at the onsets

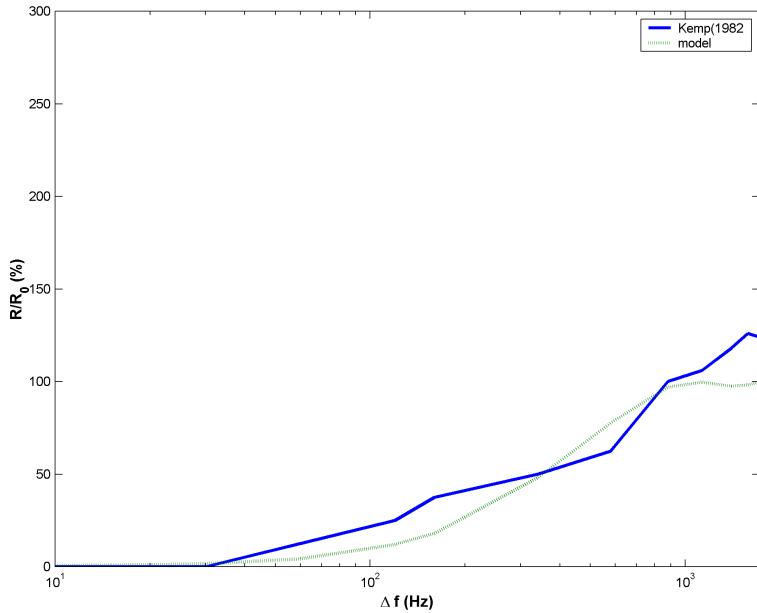


Figure 3.4: Comparison of the model's output (dotted line) with psychoacoustical data (full line). The signal is a 1.6 kHz FM tone at 60 dB with $\Delta f = 800$ Hz measured at f_{mod} from 1 10 20 40 60 80 100 200 300 400 500 Hz.

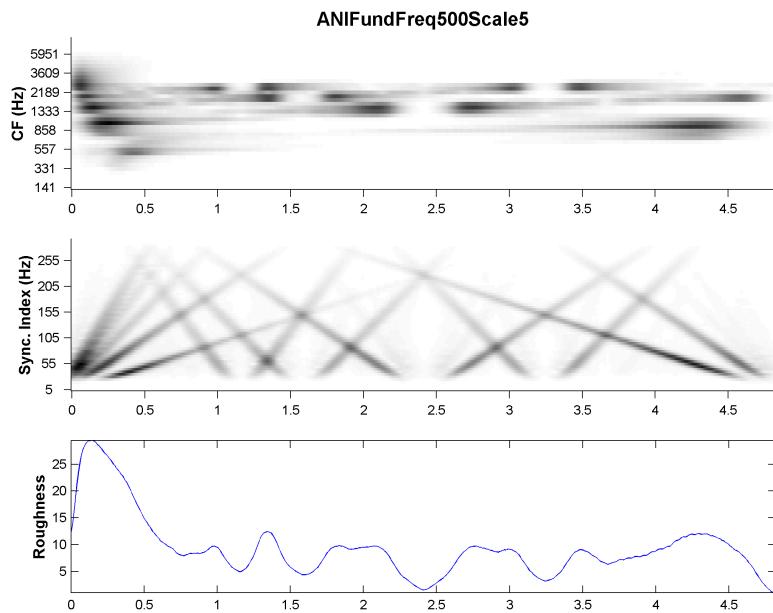


Figure 3.5: A harmonic tone complex with f_0 500 Hz is played together with a pitch shifted copy (up to f_0 1000 Hz). The harmonics have equal amplitudes. Top panel: excitation in the auditory channels. Middle panel: synchronization index. Lower panel: roughness as calculated with the synchronization index model.

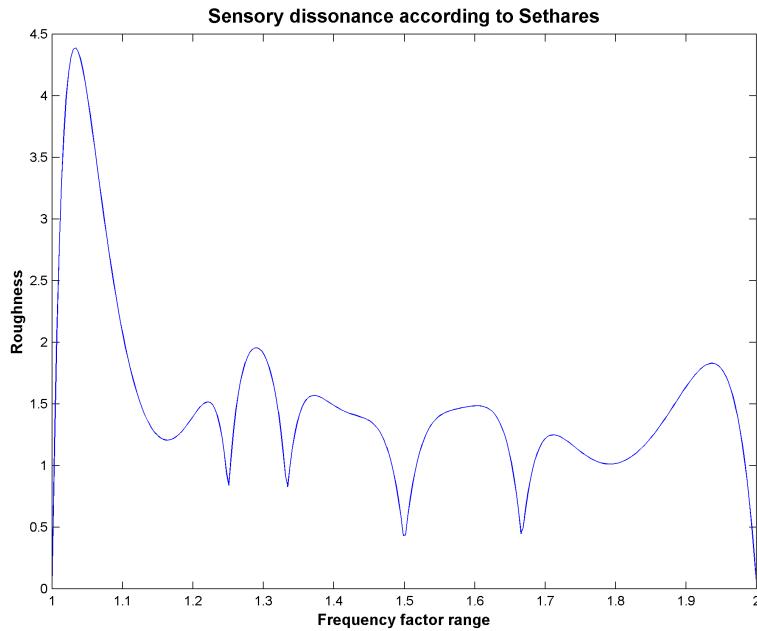


Figure 3.6: Sensory dissonance as calculated with the curve-mapping model of Sethares

The roughness curve in figure 3.5 can be compared with the results of the curve-mapping model of roughness of Sethares (1998)(figure 3.6). This model takes the frequency-amplitude values as input (no sound!) and calculates the curve using the psychoacoustical curve of ? (?). The Synchronization Index Model (figure 3.7), apart from its good agreement with this theoretical model, provides an additional cue in showing the 'spectral' (=excitation in the auditory channels) as well as 'temporal' (=synchronization index profile) factors that contribute to roughness. Note that the dips are less deep than in Sethares' model, which is partly due to the fact that the model works on a glissando of harmonic tone complexes, rather than on separate intervals. The points of minimal roughness indicate a hierarchical order of intervals in terms of roughness. This hierarchy can be musically exploited as the points of minimum roughness may indicate candidates for a musical scale.

3.2.4 Results and discussion

Dependency

The dependency of roughness on m , f_m and f_c provides evidence for the theory that roughness may be directly related to the degree in which neurons synchronize with the beating frequencies of the AM tones. Roughness should then depend on two basic effects:

- The 'spectral' filtering resulting from the resonance properties of the basilar membrane.
- The neuronal synchronization and its possible limitations.

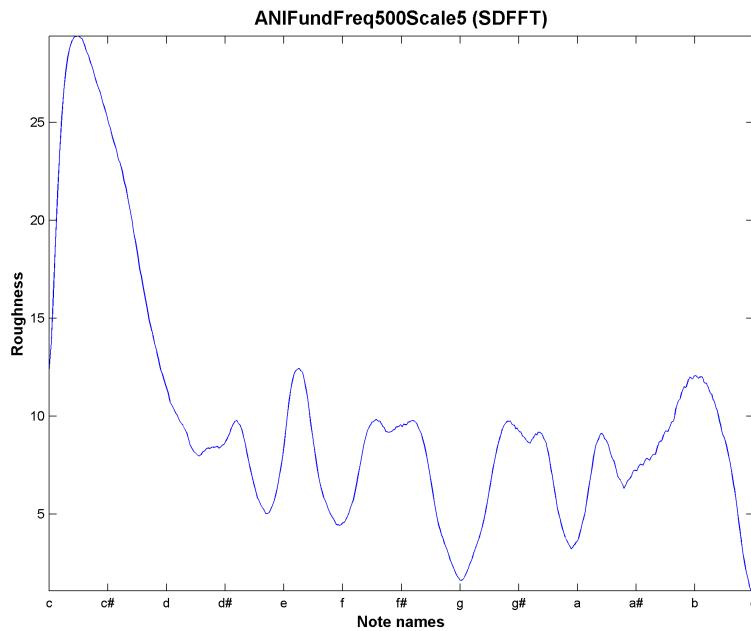


Figure 3.7: Roughness as calculated with the Synchronization Index Model

Modulation frequency smaller than 25 Hz

For f_m smaller than about 25 Hz, the sensation of roughness goes over into a sensation of fluctuation and loudness. Roughness starts at about $f_m > 25$ Hz but then the effects of spectral filtering and synchronization should be taken into consideration.

Frequencies outside the filter bandwidth

AM tones have spectral energy at both sides ($f_c \pm f_m$) of a carrier frequency f_c . An increase of the modulation frequency f_m results in a widening of the intervals. Consequently, when the side frequencies ($f_c \pm f_m$) fall outside the central scope of the f_c filter bandwidth, their energy becomes more and more attenuated and the beating decreases in relation to that. The bandwidth of the auditory filters corresponds to resonance regions of equal distance on the basilar membrane (Greenwood & Joris, 1996). Due to the quasi logarithmic relationship between distance and frequency, the filter bandwidth of this filter becomes broader in the higher frequency regions. The 'spectral' filtering that results from the resonance properties of the basilar membrane accounts for the fact that the roughness decreases when f_m exceeds the bandwidth of the filtering associated with f_c . As the frequency components exceed the bandwidth of an auditory filter, they will no longer interfere and will cease to produce beating frequencies.

Frequency components within the filter bandwidth

In the case, when the frequency components fall within the bandwidth of the spectral filter, the neurons will attempt to synchronize to the beating frequencies provided that those frequencies are not too fast. Analysis of neuronal responses in the auditory nerve indicates that neuronal synchro-

nization has an upper limit at about 1000 Hz which implies an absolute limit for inferences based on synchronization (Joris & Yin, 1992). But fast beating frequencies, however, will be picked up as pitches rather than roughness (Langner & Schreiner, 1988). Only the beating frequencies (e.g. $25\text{Hz} < f_m < 300\text{ Hz}$), will contribute to roughness, suggesting that synchronization is actually subjected to a band-pass filtering.

Conclusion

The synchronization index model offers a standard method in the calculation of the modulation depth for roughness. The method allows a proper visualization of the features that contribute to roughness both in terms of fluctuations in the auditory channels, and in terms of phase-locking synchrony. The visualization provides interesting cues for analyzing the factors that contribute to roughness.

3.3 Tonality induction experiments

3.3.1 Introduction

In this demonstration contextuality is applied to simulate two psychological experiments described by Krumhansl and Kessler (1982). The application is based on an auditory model testing the role of short-term memory in probe-tone experiments ? (?). The paper investigates the relationship between short-term memory and the probe-tone technique by means of auditory modeling. The results of the simulation of the probe-tone Experiment I and Experiment II of Krumhansl and Kessler show that a short-term memory model, based on echoic images of periodicity pitch patterns, gives a good fit with the probe-tone ratings.

3.3.2 Method

The probe-tone technique inspects a tonal context by means of probe-tones. The listener hears a chord or a scale in order to establish a key, then a probe tone is presented and the listener is asked to rate how well that tone fits within the key being examined. Krumhansl and Kessler (1982) take the technique as an experimental paradigm for the study of context-dependent pitch perception. The tonal context is generated by a so-called *context-defining sequence*, which is usually a harmonic progression of chords expressing a stable major or minor key, but a major or minor scale of pitches can be used as well. After presentation of the context-defining sequence followed by a probe, the listener is asked to rate the degree of fit between probe and sequence on a scale. The probe is varied over all 12 pitches of the chromatic scale hence 12 trials are needed to probe a context-defining sequence. The responses obtained for those 12 trials specify a (tonal hierarchy) *profile*.

The simulations are done using an auditory pitch model. They aim to process the stimuli given to human subjects in the psychological experiments. The simulated probe-tone ratings are focussed on stimulus-driven inferences. The *Contextuality Module* is aimed at simulating the degree-of-fit task of probe-tone experiments. Simulation I and II cover all essential steps in the probe-tone experiments of Krumhansl and Kessler.

Framework

The auditory workspace for the simulation of the probe-tone experiments draws on the distinction between *auditory images, processes and inferences*. The modules involved in the auditory model are the *Auditory Peripheral Module*, the *Pitch Completion Module*, the *Echoic Memory Module* and the *Contextuality Module*. An auditory peripheral system transforms sound waveforms into auditory patterns or primary images. The Pitch Completion Module then transforms the primary images into pitch images. The stimulus-driven inference is based on a comparison of images having undergone the effect of an echo. To measure pitch commonality between pitch images Leman uses two types of stimulus-driven inference called Method I and Method II.

- Inspection:

In the first method stimulus-driven inference has been interpreted as an inspection of the profile build-up during the temporal deployment of the context-defining sequence. A local

image is fixed at the end of the probe-tone. Running images over the whole stimulus are inspected by means of this image.

- Comparison:

In the second method pitch images with possible different echo are compared at running time. Method II assumes two memory buffers whose size corresponds to the size of a local plus a global image. Both buffers are updated at run time and operate as echoic memories.

3.3.3 Application

Simulation of probe-tone experiments

The short-term memory models related to the methods of inspection and comparison for measuring the pitch commonality of local versus global pitch images are used in Simulation I and II of Krumhansl and Kessler's Experiment I and II ([Krumhansl & Kessler, 1982](#)).

Performance

The contents of the demonstration package for the contextuality experiments can be found in the directory IPEM\Demos\Contextuality. To perform the simulation of Experiment I and II described by Krumhansl and Kessler first create the periodicity pitch patterns of all pitches and chords used in order to construct the sounds. The calculations are done using the function **IPEM-GenerateProbes**. The periodicity pitch results and the signals for the sound sequences are stored in the directory PitchImagesShepardTone, relative to IPEMRootDir ('output')\ContextualityDemo. The content of this directory takes 8,44Mb. Then perform the two simulations with the functions **IPEMProbeToneExperimentKK1** and **IPEMProbeToneExperimentKK2**.

Remark: Note that performing this demonstration is time-consuming for it takes around 2 hours and 30 minutes on a Windows NT 4.0 system with Pentium III 600 MHz processor, 256Mb RAM PC130.

Simulation I

Simulation I corresponds to Experiment I that probed for the major and minor key templates. In Experiment I, Krumhansl and Kessler used twelve different context-defining elements, called here sequences. (In our terminology, a sequence may have one single tonal element.) The sequences are shown in Table 3.1.

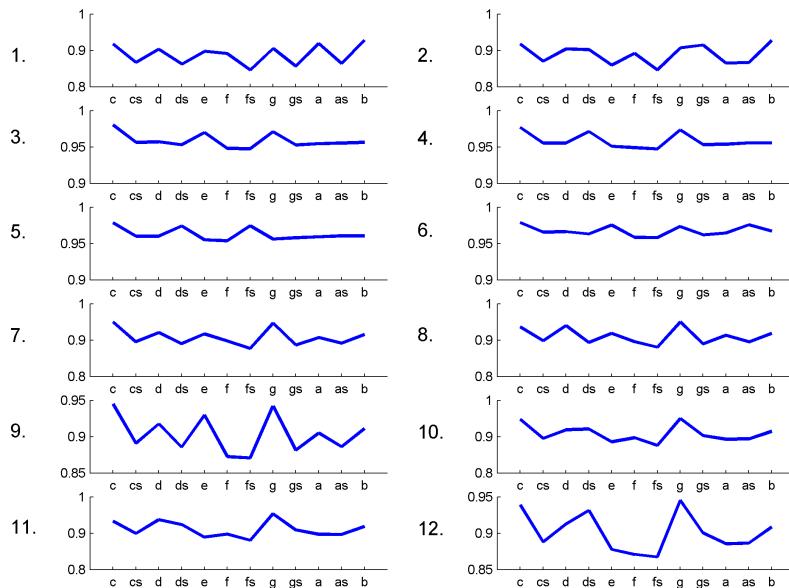
Each trial consists of a context-defining sequence followed by a probe tone. Given 12 sequences and 12 probes, this gives $12 \times 12 = 144$ trials. On each trial, the context-defining sequence was heard first, followed by a one-second pause, and then the probe sounded for 0.5 seconds. The timings of the different context-defining sequences were as follows: the tonic tones of the two scales (major and minor) were sounded for 0.5 seconds, and the remaining scale tones for 0.25 seconds, with a 0.19-seconds pause between scale tones. The single-chord sequences were played for 0.5 seconds, as were the chords in the three-chord cadences. In the cadences there was about

Context-Defining Element	Example
1. Ascending major scale	c d e f g a b c
2. Ascending harmonic minor scale	c d e \flat f g a \flat b c
3. Major chord	c e g
4. Minor chord	c e \flat g
5. Diminished chord	c e \flat g \flat
6. Dominant seventh chord	c e g b \flat
7. IV-V-I cadence in major	F maj chord G maj chord C maj chord
8. II-V-I cadence in major	D min chord G maj chord C maj chord
9. VI-V-I cadence in major	A min chord G maj chord C maj chord
10. IV-V-I cadence in minor	F min chord G maj chord C min chord
11. II-V-I cadence in minor	D dim chord G maj chord C min chord
12. VI-V-I cadence in minor	A \flat maj chord G maj chord C min chord

Table 3.1: Context Defining Elements Used in Experiment 1

0.25 seconds between chords. In order to achieve stylistic neutrality, the tones used in all trials were Shepard tones.

Simulation I uses the same 144 trials. First the sounds are processed with APM, PCM, and EMM, into local and global images. The echos are respectively set to $T = 0.1$ and $T = 1.5$. The stimulus-driven inference is calculated for each trial and a snapshot is taken at the end of the trial. (Method I and Method II deliver the same results at this point.) The snapshots of all 12 different probe tones for one single context-defining sequence then represents the profile of that sequence. At the end, there are 12 profiles, corresponding to the 12 context-defining sequences of Table 3.1. Figure 3.8 shows the twelve profiles that result from this simulation. The number to the left of each graph corresponds with the number of the context-defining sequence in Table 3.1.

Figure 3.8: Profiles of 12 context-defining sequences for tone center $T = 1.5$.

For each figure, the vertical axis represents the correlation between the local image and the global image at the end of the trial (12 trials are represented in each figure). What matters is the relative relationships between the correlation coefficients within a profile. In fact, when profiles are compared with each other, using the correlation coefficient as a measure of comparison, then the relative values are important. Table 3.2 shows the similarities among the twelve profiles.

element	1.	2.	3.	4.	5.	6.	7.	8.	9.	10.	11.	12
1.	1.000	0.457	0.421	0.219	-0.226	0.336	0.709	0.720	0.665	0.428	0.408	0.315
2.		1.000	0.208	0.439	-0.024	0.090	0.404	0.384	0.289	0.698	0.708	0.671
3.			1.000	0.670	0.163	0.867	0.865	0.724	0.916	0.647	0.492	0.589
4.				1.000	0.466	0.531	0.647	0.525	0.588	0.887	0.756	0.923
5.					1.000	0.053	0.000	-0.140	0.010	0.224	0.065	0.332
6.						1.000	0.677	0.577	0.770	0.457	0.334	0.435
7.							1.000	0.938	0.940	0.794	0.712	0.659
8.								1.000	0.910	0.727	0.768	0.629
9.									1.000	0.662	0.609	0.611
10.										1.000	0.936	0.943
11.											1.000	0.917
12.												1.000

Table 3.2: The correlation coefficients represent the similarity between the profiles, for the 12 context-defining sequences used in Simulation I

The following observations can be made:

- Sequence 1 (major scale) correlates best with sequences 7, 8, and 9 (the cadences in major). The average correlation between the major scale (1) and the three cadences in major (7, 8, 9) is 0.698. It was 0.796 in the data of Krumhansl and Kessler.
- Sequence 2 (minor scale) correlates best with sequences 10, 11, 12 (the cadences in minor). The average correlation between the minor scale (2) and the three cadences in minor (10, 11, 12) is 0.6923. It was 0.727 in the data of Krumhansl and Kessler.
- Sequence 3 (major chord) correlates best with sequences 6, 7, 8, and 9 (dominant seventh chord and major cadences). The average correlation between the major chord (3) and the three cadences in major (7, 8, 9) is 0.835. It was 0.896 in the data of Krumhansl and Kessler.
- Sequence 4 (minor chord) correlates best with sequences 10, 11, 12 (minor cadences). The average correlation between the minor chord (4) and the three cadences in minor (10, 11, 12) is 0.855. It was 0.910 in the data of Krumhansl and Kessler.
- Sequence 5 (diminished chord) has a low correlation with all other sequences.
- Sequence 6 (dominant seventh chord) has the highest correlation with sequence 9 (VI V I cadence in major).
- All cadences in major are strongly connected, and all cadences in minor are strongly connected.

Krumhansl and Kessler concluded that the major chord element and the three cadences in major indicate a consistent pattern of ratings. In a similar way, the major scale was somewhat less similar. They therefore took the major key profile as the average ratings, given the 12 probe tones, for the major chord and the three cadences in major. In correspondence with this idea, and justified by the results shown in Table 3.2, the major key profile of Simulation I is defined as the average of the profiles that correspond to the sequences 3, 7, 8, and 9. The minor key profile of Simulation I is determined according to the same reasoning, taking the average of the profiles that correspond to the sequences 4, 10, 11, and 12. The similarity of the major key profile of Experiment I with the major key profile of Simulation I has correlation coefficient of 0.848. The similarity of the minor key profile of Experiment I with the minor key profile of Simulation I is 0.825. This is a significant correlation showing that the short-term memory model gives a good account of the data in Experiment I.

An important parameter of the model pertains to the echo of the global images. In order to clarify the role of the echo, a series of simulations are carried out in which the echo of the global image is systematically varied from $T = 0.2$ to $T = 5$ in steps of 0.2. Twenty-five figures are plotted that show the effect of changing the echo (T) of the global images on the profiles of the 12 context-defining sequences. Each time the 144 trials are processed. Figures 3.9 to 3.12 show the 12 profiles for halfdecay tone centers with the echo set to $T = 0.4$, $T = 1.4$, $T = 2.4$ and $T = 3.4$

Simulation II

Simulation II corresponds to Experiment II that aimed at sensing directly how the listener's sense of key develops and changes, providing a quantitative measure of the relative strengths of tonal interpretations at each point in time. The test uses 10 different chord sequences containing modulations from one key to another key. Table 3.3 gives an overview of the 10 different chord sequences. The sequences, ordered from 1 to 10 are each composed of 9 chords. They have a starting key and an ending key, and the modulation from one key to the other key can be direct, indirect, close, or remote.

Sequence 1 has no modulation, starting and ending key are in C major.

Sequence 2 has no modulation, starting and ending key are in C minor.

Sequence 3 has a direct and close modulation, starting in C major and ending in G major.

Sequence 4 has an indirect and remote modulation, starting in C major and ending in B \flat major.

Sequence 5 has a direct and close modulation, starting in C major and ending in A minor.

Sequence 6 has an indirect and remote modulation, starting in C major and ending in D minor.

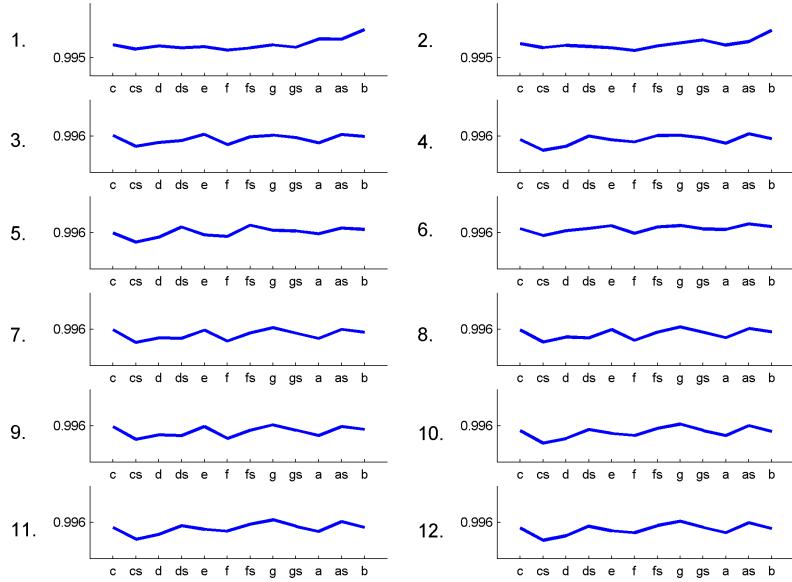
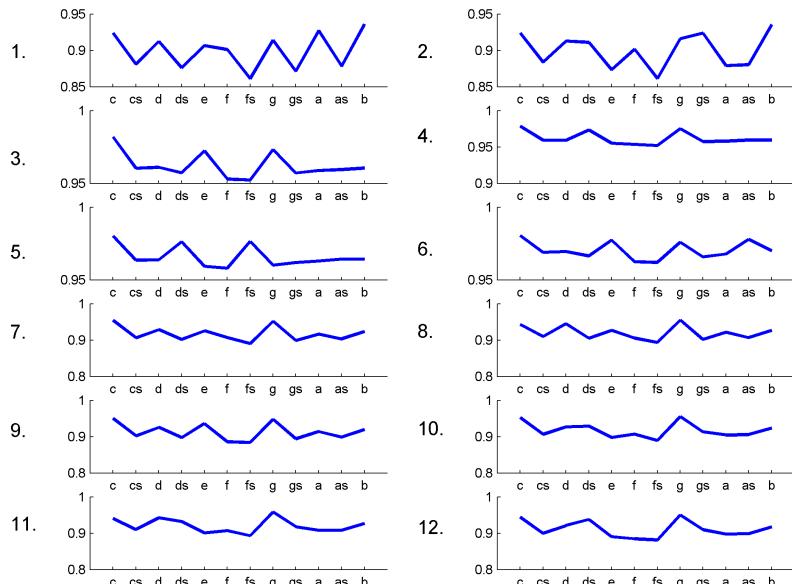
Sequence 7 has a close modulation, starting in C minor and ending in F minor.

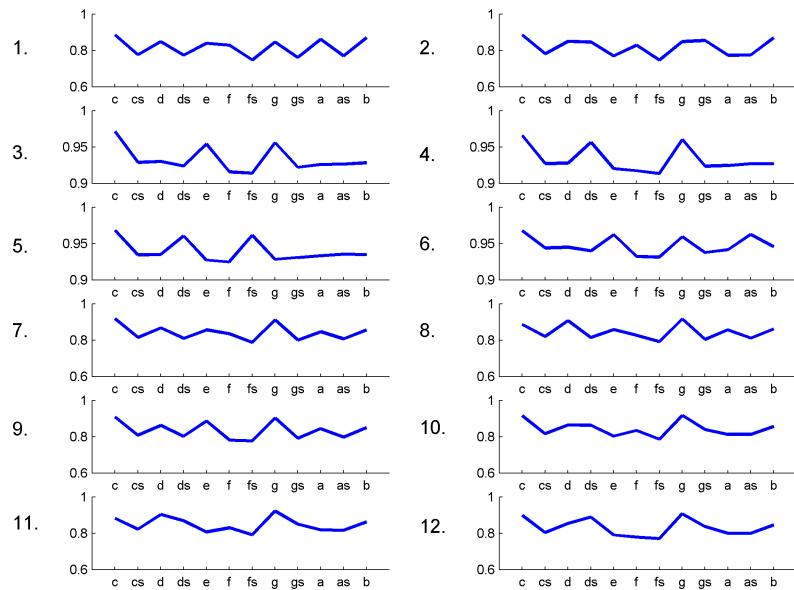
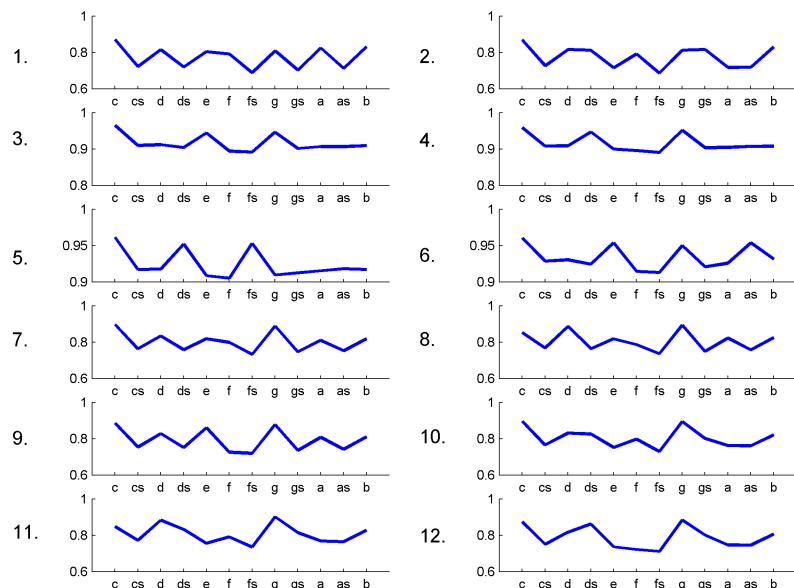
Sequence 8 has a remote modulation, starting in C minor and ending in C \sharp minor.

Sequence 9 has a close modulation, starting in C minor and ending in C major.

Sequence 10 has an indirect and close modulation, starting in C minor and ending in A \flat major.

A single chord sequence consists of 9 chords and for each sequence, there were 9 trials. In the first trial, the context-defining sequence was limited to the first chord and the listeners were asked to probe the chord with the 12 probe tones. In the second trial, the context-defining sequence was made of the first two chords of the sequence and the listeners were asked to probe this two-chord sequence with the 12 probe tones. In the third trial, the context-defining sequence was made of

Figure 3.9: Profiles of 12 context-defining sequences for tone center $T = 0.4$.Figure 3.10: Profiles of 12 context-defining sequences for tone center $T = 1.4$.

Figure 3.11: Profiles of 12 context-defining sequences for tone center $T = 2.4$.Figure 3.12: Profiles of 12 context-defining sequences for tone center $T = 3.4$.

	Chords								
	1	2	3	4	5	6	7	8	9
Seq. 1:	Fmaj	Gmaj	Amin	Fmaj	Cmaj	Amin	Dmin	Gmaj	Cmaj
Seq. 2:	Fmin	Gmaj	A♭maj	Fmin	Cmin	A♭maj	Ddim	Gmaj	Cmin
Seq. 3:	Fmaj	Gmaj	Cmaj	Amin	Emin	Bmin	Emin	Dmaj	Gmaj
Seq. 4:	Fmaj	Gmaj	Cmaj	Amin	Fmaj	Gmin	E♭maj	Fmaj	B♭maj
Seq. 5:	Fmaj	Gmaj	Cmaj	Fmaj	Dmin	Emaj	Bdim	Emaj	Amin
Seq. 6:	Fmaj	Gmaj	Cmaj	Fmaj	Dmin	B♭maj	Edim	Amaj	Dmin
Seq. 7:	Ddim	Gmaj	Cmin	A♭maj	Fmin	D♭maj	B♭min	Cmaj	Fmin
Seq. 8:	Ddim	Gmaj	Cmin	Gmaj	A♭maj	Amaj	F♯min	G♯maj	C♯min
Seq. 9:	Ddim	Gmaj	Cmin	A♭maj	Gmaj	Amin	Fmaj	Gmaj	Cmaj
Seq. 10:	Ddim	Gmaj	Cmin	A♭maj	Fmin	E♭maj	B♭min	E♭maj	A♭maj

Table 3.3: Chord Sequences Used in Experiment II. Cmaj means the C major chord, Cmin the C minor chord, and Cdim, the C diminished chord.

the first three chords of the sequence and the listeners were asked to probe the sequence with the 12 probe tones, and so on. In the ninth trial, the context-defining sequence was equal to the complete sequence. For each chord sequence, one thus obtained 9 profiles containing information about the gradual temporal deployment of the tonal sequence. The 10 different chord sequences thus correspond to 10 different *profile-sequences*, each containing 9 profiles.

The analysis of a profile-sequence had two parts. In the first part the analysis was done in reference to the profile of the first key of the corresponding chord sequence. In the second part, the analysis was done in reference to the second key of the corresponding chord sequence. Thus, if the chord sequence started in C major, and modulated to G major, then the first analysis correlated each profile of the sequence with the profile of the C major key. The second analysis correlated each profile of the sequence with the profile of the G major key. One thus obtained for each chord sequence two (nine-step) *correlation-sequences*, one correlation-sequence representing the sense of key with respect to the first key, another correlation-sequence representing the sense of key with respect to the second key.

Simulation II has then been set up as follows. The 10 chord sequences with 9 subsequences each, define 90 context-defining sequences. Each sequence was processed as follows:

1. The sound $s(t)$ of the chord sequence is processed with the auditory model (APM, PCM, and EMM) into local and global images. The echos are set to $T = 0.1$ and $T = 1.5$.
2. The global images $\tilde{p}_{T=1.5}(t)$ are inspected by the 12 different probe tones using the stimulus-driven inference Method I. This was done in order to get an insight into the way in which the profile-sequence is built up at run-time. Figure 3.13 shows the profile-sequence for the global images inspected by the probe tones. Chord sequence 4 is used as an example.
3. In a similar way, the local images $\tilde{p}_{0.1}(t)$ are inspected by the 12 different probe tones using the stimulus-driven inference Method I. Figure 3.14 shows the profile-sequence of chord sequence 4 for the local images inspected by the probe tones.
4. The profiles are then correlated with the major and minor key profiles of Simulation I, in agreement to the setup of Experiment II. The key profiles are obtained by rotation of the C major and C minor key profiles of Simulation I, similar to the approach taken by Krumhansl and Kessler.
5. Finally, the averaged correlations are taken according to the series specified in the original simulation, that is: (i) average chord sequences 1 and 2, (ii) average chord sequences 3, 5, 7, 9, 10 for the first key, and for the second key, (iii) average chord sequences 4, 6, and 8 for the first key, and for the second key.

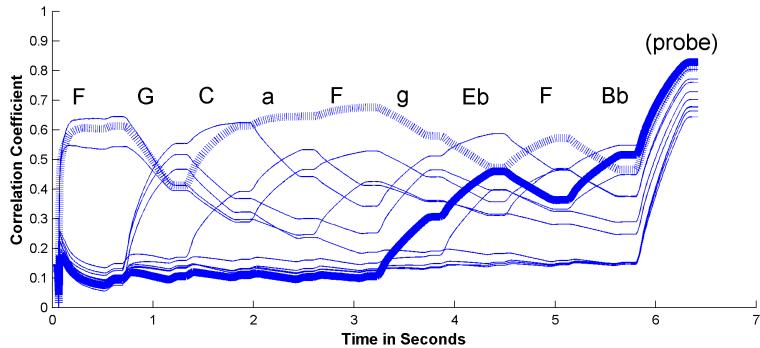


Figure 3.13: Stimulus-driven inferences (Method I) of chord sequence 4 used in Simulation II. The labels specify major and minor chords and are located at the time instances where they are introduced. The context-defining sequence is probed with 12 different probe tones. The stimulus-driven inference for each trial is plotted. The buildup and decay of the inferences depends on the echo of the global images ($T = 1.5$). The dotted line represents the probe tone c , the solid line represents the probe tone bb

The effects of echo are clearly visible in Fig. 3.13. At the beginning of the sequence, one gets the same profile for the F major chord as in Fig. 3.14. But the next chords, G major and then C major lead to quite different profiles, owing to the effect of the echo. The dotted line shows

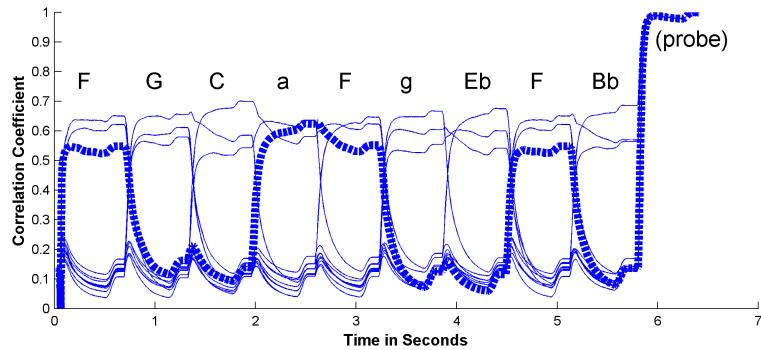


Figure 3.14: Stimulus-driven inferences (Method I) of chord sequence 4 used in Simulation II. The global images have the same echo as the local images ($T = 0.1$). The dotted line represents the probe tone f .

the correlation with the probe tone c , the tonic of the first key, the full line shows the correlation with the probe tone bb , the tonic of the second key. Shuffling chords around may have an effect on the resulting profile. Certain chord progressions will prevent certain pitches to pop up, while other chord progressions will favor other pitches in the profile. But the echo takes into account the temporal order of the chords, provided that it is within the limits of the half decay time.

3.3.4 Results and discussion

Using the same stimuli as in Experiment I and Experiment II of Krumhansl and Kessler ([Krumhansl & Kessler, 1982](#)), the short-term memory model produces a good fit with the data. A comparison of the key profiles of Simulation I with those obtained in Krumhansl and Kessler's Experiment I gives a correlation of 0.848 for the major key, and 0.825 for the minor key. The echo, or half decay time, of the global image, which determines the content of the context-defining sequence, was determined to be about 1.5 seconds. Using this echo, Simulation II aimed at repeating Krumhansl and Kessler's Experiment II in which the dynamic aspects of key sensitivity were investigated.

Simulation I

The results of Simulation I are summarized in Table 3.4.

The left column contains the half decay time T of the global images (the local images have a constant echo of $T = 0.1$), the first column reports the averaged correlation of the profile of sequence 1 with the profile of sequences 7, 8, and 9. The second column contains the average correlations of the profiles of sequence 2 with those of sequences 10, 11, and 12, the third column of 3 with 7, 8, 9, and the fourth of 4 with 10, 11, 12. The last two columns contain the correlations of the major and minor key profiles of Experiment I with those of Simulation I.

At $T = 0.2$ the effect of the tonal context can be neglected. Hence one gets high correlations between the profiles, but very low correlations with the key profiles. At $T = 0.4$ the effect of context is growing, but the local effects are still dominating. For $T > 0.4$ the effect of context becomes visible and three general trends can be distinguished:

- The scale becomes more similar to the cadences, both for major and minor.

T (=echo)	1 vs 7+8+9	2 vs 10+11+12	3 vs 7+8+9	4. vs 10+11+12	KKMajor	KKMinor
0.2	0.989	0.987	1.000	1.000	-0.519	-0.087
0.4	0.454	0.531	0.895	0.891	0.264	0.493
0.6	0.258	0.309	0.907	0.917	0.771	0.808
0.8	0.390	0.408	0.891	0.909	0.809	0.812
1	0.509	0.514	0.872	0.891	0.827	0.817
1.2	0.603	0.601	0.855	0.875	0.838	0.821
1.4	0.672	0.667	0.841	0.861	0.845	0.824
1.6	0.720	0.714	0.829	0.850	0.851	0.826
1.8	0.755	0.748	0.820	0.840	0.855	0.828
2	0.780	0.772	0.813	0.832	0.858	0.829
2.2	0.799	0.790	0.806	0.826	0.860	0.830
2.4	0.813	0.803	0.801	0.820	0.863	0.831
2.6	0.824	0.814	0.796	0.815	0.864	0.831
2.8	0.831	0.821	0.792	0.811	0.866	0.832
3	0.838	0.827	0.789	0.807	0.867	0.832
3.2	0.843	0.832	0.786	0.804	0.868	0.833
3.4	0.847	0.836	0.783	0.801	0.869	0.833
3.6	0.851	0.839	0.780	0.798	0.870	0.833
3.8	0.853	0.842	0.778	0.796	0.870	0.834
4	0.856	0.844	0.776	0.794	0.871	0.834
4.2	0.858	0.846	0.775	0.792	0.871	0.834
4.4	0.860	0.847	0.773	0.790	0.872	0.834
4.6	0.861	0.849	0.771	0.788	0.872	0.834
4.8	0.862	0.850	0.770	0.787	0.873	0.835
5	0.863	0.851	0.768	0.786	0.873	0.835

Table 3.4: The effect of changing the echo (T) of the global images on profiles

- The chords become less similar to the cadences, both for major and minor.
- The key profiles become more similar.

The trend can be explained in terms of the increasing contribution of the tonic of the scale when T increases. Taking into account the data of Krumhansl and Kessler, who observed a less good correlation with the scales than with the chords, the optimal echo for global images is taken to be $T = 1.5$. The effect is shown in Fig. 3.15, where $T = 4$, all other parameters being equal.

Simulation II

The results of Simulation II Method I are summarized in five graphs shown in Fig. 3.16.

Method I gives a useful insight in how the profiles are build up during the temporal deployment of the chord sequence but the actual build up of the profile should not be taken into consideration when data are compared with the probe tone ratings. What actually matters in that case is the snapshot taken at the end of each trial of the chord sequence. This can be obtained with either Method I or Method II. Fig. 3.16 shows the results when the average correlations are taken at the end of each chord, for the nine chords in the sequence.

The data of Experiment II are summarized as follows:

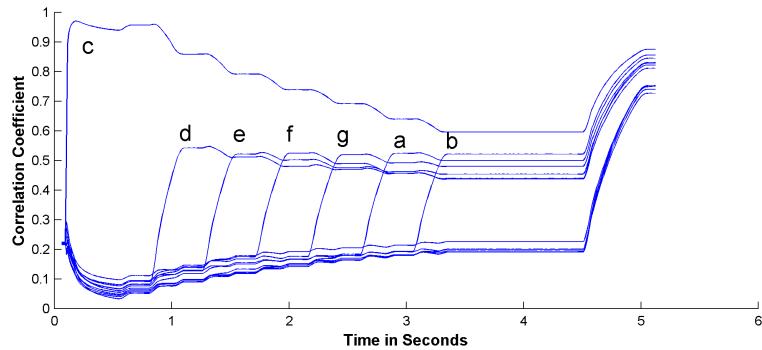


Figure 3.15: Stimulus-driven inference of the ascending C-major diatonic, using an echo of $T=4$ for the global images.

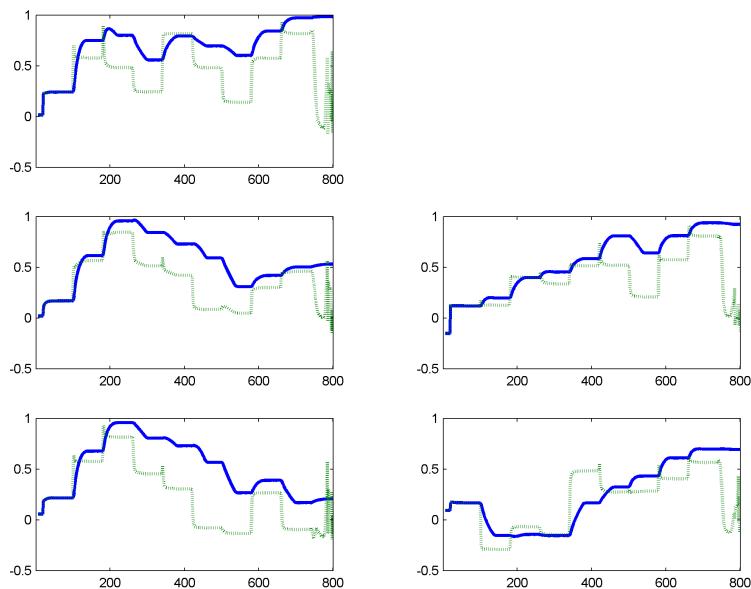


Figure 3.16: Results of Simulation II. The horizontal axis represents time in hundreds of seconds. The vertical axis is the correlation. The average correlations are taken at the end of each chord, for the nine chords in the sequences. The buildup effect of the local profiles now disappears.

- The full line in the top left figure represents the average correlation-sequence for chord sequences 1 and 2. These chord sequences have no modulation, hence, the correlation-sequence for the first key is identical to the correlation-sequence for the second key. The data therefore collapse to a single figure.
- The full lines in both middle figures represent the average correlation-sequences for chord sequences 3, 5, 7, 9 and 10. These chord sequences have a modulation between closely related keys. Given the two different keys, the results are split up in two figures, one left and one right figure. The full line on the left figure contains the averaged correlation-sequence with

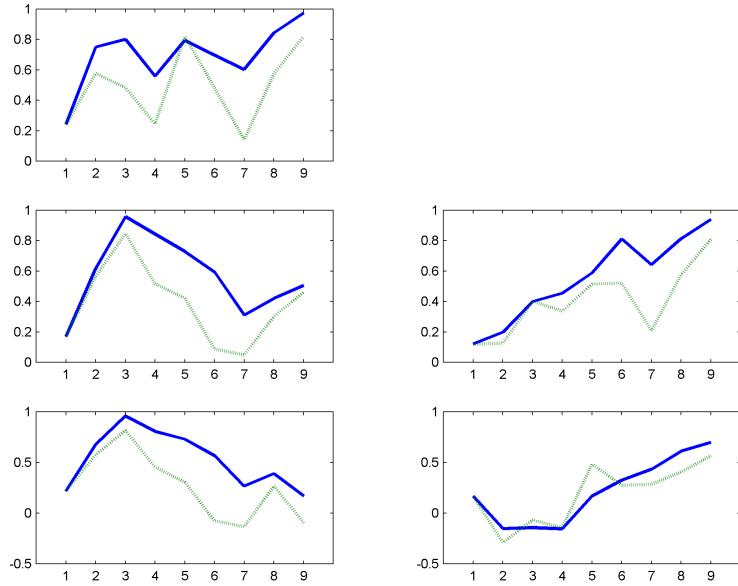


Figure 3.17: Results of Simulation II. Same graphs as previous figure, but the points at the end of each chord have been selected

the first key. The full line on the right figure contains the averaged correlation-sequence with the second key.

- The full lines of the bottom figures represent the average correlation-sequences for the chord sequences 4, 6, and 8. These chord sequences have a modulation between distant keys. The full line on the left figure contains the averaged correlation-sequence with the first key. The full line on the right figure contains the averaged correlation-sequence with the second key.

The dashed line in each of the left-sided figures represents the averaged correlation between the profile of the individual chords and the profile of the first key, at each of the nine steps. The dashed line in each of the right-sided figures represents the averaged correlation between the profile of the individual chords and the profile of the second key, at each of the nine steps. For example, in the top left figure, the dashed line of the first point is the average of two correlations, in particular the correlation between the profile of the F major chord, and the C major key profile, and the correlation between the profile of the F minor chord and the C minor key profile. (The chord profiles and the key profiles used to calculate the dashed line were all obtained in Experiment I.) The correlations between the individual chord profiles and the key profiles predict how strongly each chord in isolation would be expected to be related to the intended key of the sequence and, thus, serve as a baseline against which the strength of the intended key can be compared. There is a striking similarity with the results of Krumhansl and Kessler Experiment II. The correlation ranges show a good agreement for both lowest and highest values. There is a small difference in the comparison of profiles to the first key (=left column figures). The listeners seem to sense the key already after the second chord, which is not the case in the simulation, where the correlations corresponding to the global and local images are almost similar for the first two chords of the

sequence.

Chapter 4

Applications

4.1 Introduction

Whereas the previous chapters handled the basic modules and a comparison of their results with psychoacoustic data, this chapter focuses on some more elaborated applications where these modules are used in practice. Each of the sections in this chapter contains:

- a small introduction in which the purpose of the application is explained
- a description of the method that is used
- a summary of how to run the application in practice using the toolbox and its functions
- an overview of the obtained results and a discussion thereof

The applications are explained from a practical viewpoint, and each time a link is made between the related concepts and the practical tools (MATLAB functions) available in the toolbox.

The following applications are included:

- Roughness applications: relationships between timbre, scales and roughness.
- A rhythmic pattern extraction demonstration for extracting repetitive rhythm patterns from a sound fragment.

Each application can be seen as a demonstration of one (or more) module(s) that were handled in the first chapter. You can of course freely experiment with the modules yourself and build some new applications from the basic functions. If you think that you've come up with something interesting, we would love to hear about it!

4.2 Roughness applications

... To Be Completed ... (ask Marc for latest version since Koen sent it for that paper)

4.2.1 Introduction

4.2.2 Method

4.2.3 Application

4.2.4 Results and discussion

4.3 Rhythmic pattern extraction

4.3.1 Introduction

In this demonstration, the MEC algorithm from the [Rhythm Module \(RhM\)](#) is used to detect and extract rhythmic patterns from a sound fragment.

The extracted patterns are then used to synthesize sound using modulated AM noise, which allows for an auditive comparison between the original sound and the sound resynthesized from what the Rhythm Module has extracted.

4.3.2 Method

For the extraction of rhythmic patterns, we are interested in periods ranging from roughly 50 ms to 5 seconds.

The signal that will be analyzed is a root-mean-square (RMS) signal calculated every 10 ms over a frame of 20 ms. We can calculate the RMS:

- of the signals in the different channels of the output of an auditory model (which yields a multi-channel RMS signal)
- of the sound signal itself (which yields a 1-dimensional RMS signal)

This energy signal is then first analyzed using the MEC algorithm and the most prominent periods occurring in the RMS signal is detected by finding the minimum in the calculated difference values. In the first case, this can be done by detecting a minimum for each channel separately (yields best period for each channel) or by summing the difference values from all channels first and then picking the minimum from there (yields a single best period for all channels).

Once the best periods are known, the patterns corresponding to that period can be extracted from the RMS signal.

Then, a specific moment in time is chosen, and the extracted patterns at that moment are used for a resynthesis in order to listen to what the MEC algorithm has found as rhythmic pattern.

For more detailed information on the Rhythm Module itself, see page [43](#).

4.3.3 Application

The contents and the functions for this demo on rhythmic pattern extraction using MEC can be found in the directory IPEMToolbox\Demos\MECPatternExtraction.

The main functions are:

- [IPEMDemoMECRhythmExtraction](#)

This demo function starts an interactive session in which you can select a sound file, set some basic parameters, analyze the sound using MEC and then resynthesize using the analysis results.

- **IPEMDemoStartMEC**

This demo function starts an entire MEC analysis run in one go (without interaction). It allows to set more parameters and to save the analysis results to a file.

Both demo functions are based on the core functions related to the **Rhythm Module (RhM)**. You can have a closer look to the Matlab code of the demo functions to get to know how these core functions are used, or just check out the functions in the **Reference Manual**.

Interactive demo

Figure 4.1 shows the execution flow of IPEMDemoMECRhythmExtraction.

1. First of all, the user has to select the sound file he/she wants to analyze.

Three example sound files are located in the directory containing the demo:

- Example 1: BartokScherzoSuiteOp14.wav 

This is a short fragment from the beginning of the second movement "Scherzo" from "Suite op. 14" by Béla Bartók (piano).

- Example 2: TomWaitsBigInJapan.wav 

This is an excerpt from "Big in Japan" taken from "Mule Variations" by Tom Waits (voice, bass, drums, electric guitar).

- Example 3: PhotekTheLightening.wav 

This is a short fragment from "The Lightening (Digital Remix)" by Photek (electronic, drum and bass).

2. Next, a few choices have to be made for setting the analysis parameters:

- Period range

Specify the range of periods for which periodicity should be checked. For rhythmic patterns reasonable values are from 0.050 to 5 seconds. The bigger the range, the more periods will be checked, so the slower the analysis will be.

- Type of analysis

Finding the "best periods" can be done in three ways:

- using the difference values calculated from the RMS of the sound signal itself: this yields a single "best period" at each moment in time
- using the summed difference values (over all channels) calculated from the RMS of the ANI: this yields a single "best period" for all channels together at each moment in time
- using the difference values calculated from the RMS of the ANI separately for each channel: this yields a "best period" for each channel at each moment in time

- Speed-up factor

Instead of calculating difference values for every sample in the RMS signal, you can specify a speed-up factor n (a positive integer) so that values are calculated only every n samples.

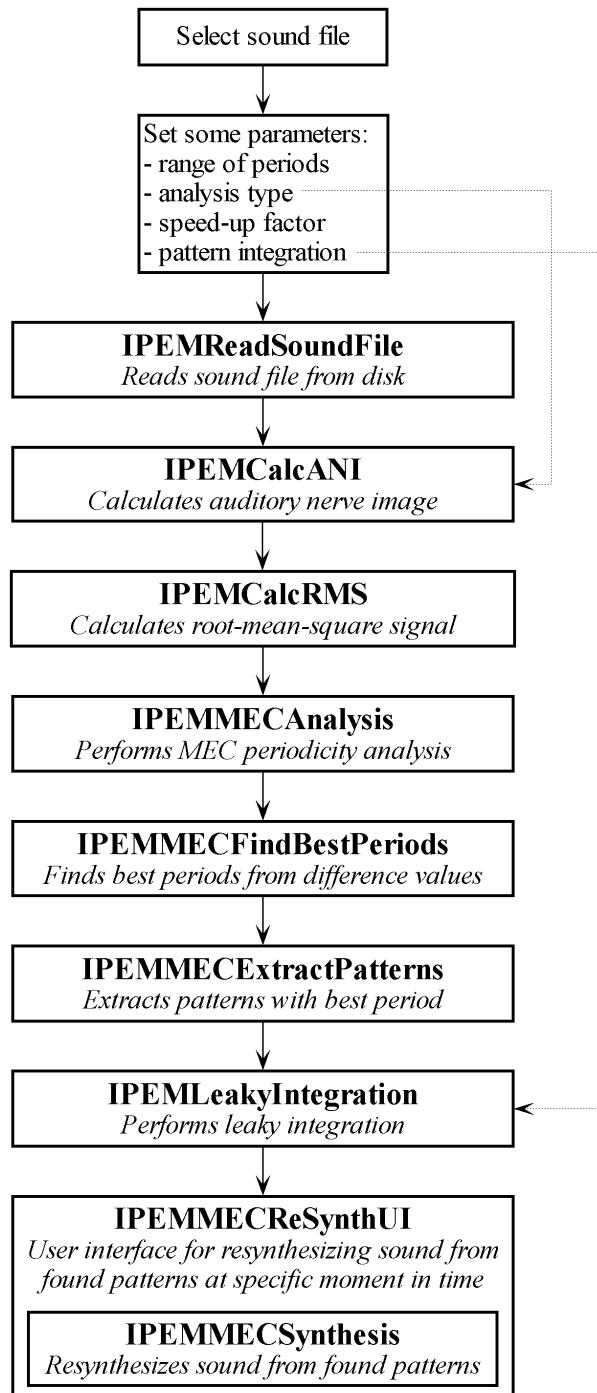


Figure 4.1: Execution flow of IPEMDemoMECRhythmExtraction

- Pattern integration

After extracting the patterns using the found best period at a certain moment, the patterns can be integrated over time. This can be done to decrease the effect of small local variations.

3. Then, the sound file is read and the RMS is calculated with **IPEMCalcRMS**, either on the auditory nerve image of the sound, or on the sound itself, depending on what was specified in the previous step. The auditory nerve image is calculated with **IPEMCalcANI** using 10 bands with a bandwidth of 1.5 critical band units (cbu), starting at 3 cbu. This corresponds to center frequencies of 215 Hz to 3266 Hz.
4. The MEC algorithm then starts analyzing the RMS signal (**IPEMMECAnalysis**) by calculating difference values.
5. Then the most prominent period(s) at each moment in time are selected. A figure with the best period(s) over time is generated (for RMS of ANI, this is a multi-channel plot, for RMS of sound signal, this is a single-channel plot).
6. The results of the previous analysis (the best period(s) at each moment) is then used to extract the corresponding pattern(s) from the RMS signal with **IPEMMECExtractPatterns**. Depending on the specified parameters, these patterns can be further integrated over time with **IPEMLeakyIntegration**.
7. Finally, resynthesis of the extracted patterns can be done using **IPEMMECReSynthUI** (which calls **IPEMMECSynthesis** internally).

The presented user interface consists of two figures: one is the figure showing the (multi-channel) best periods over time and the other is a "control palette" with some push buttons. In the first figure, a moment in time (and a specific channel) can be selected by pressing the left mouse button.

Once this selection is made, a resynthesis can be done by pressing the "Resynthesize for selected time" button. This button is only enabled when the selected time is different from the time for which resynthesis was last done.

In order to evaluate the resynthesized sound, it is combined with the original sound into a stereo sound file, where the left channel plays the original sound and the right channel plays the synthesized sound. By pressing the "Play all channels" button, the last resynthesized sounds for all channels are played together, while the "Play selected channel" only plays the selected channel.

Pressing a key (rather than clicking with the mouse) in the first figure ends the resynthesis. Figure 4.2 shows a snapshot of the resynthesis user interface.

One run demo function

This demo function does essentially the same as the above interactive demo, but it runs without any user intervention. All parameters have to be specified when calling the function, and then everything is calculated in one go, without interruption.

Figure 4.3 shows the execution flow of **IPEMDemoStartMEC**.

The results of the analysis can be saved to a .mat file so that resynthesis can be performed at a time later using either **IPEMMECSynthesis** or **IPEMMECReSynthUI** after reloading the data.

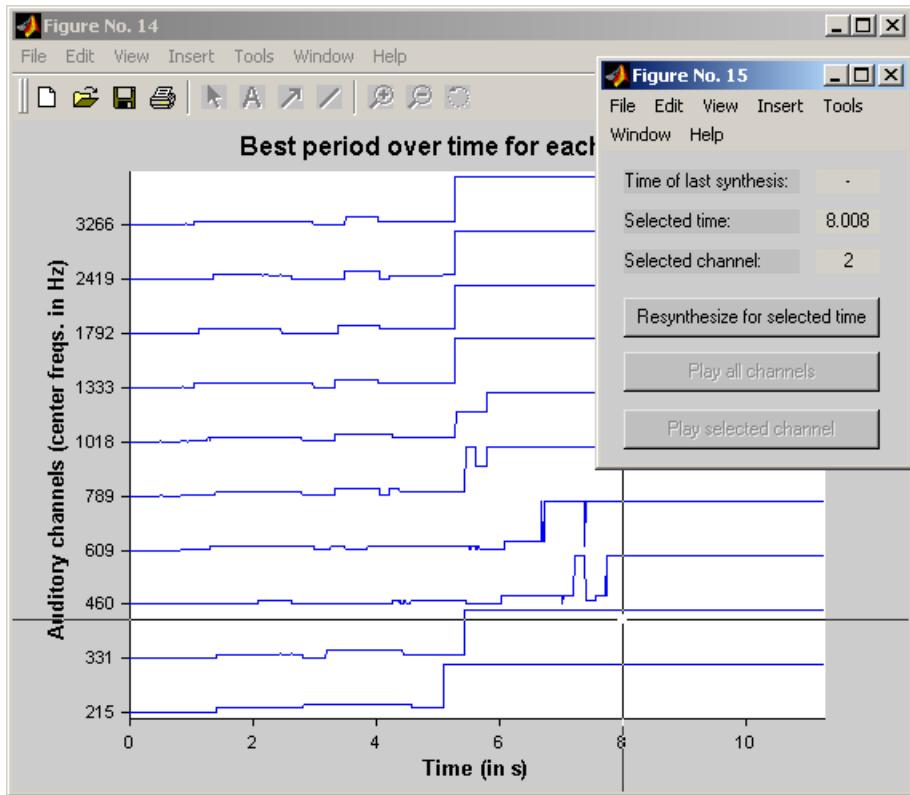


Figure 4.2: Resynthesis user interface

Remarks about the method of "resynthesis"

The resynthesis is done by amplitude modulation of specific noise bands with the extracted patterns. In case the RMS was calculated on the auditory nerve image, the noise bands have a center frequency equal to the center frequency of the corresponding auditory channel. In case the RMS was calculated directly on the sound signal, broad band noise is used. The resynthesized sound pattern is then repeated until the total sound duration matches that of the original sound. Both are finally combined into a stereo sound where one channel plays the original sound and the other plays the resynthesized sound.

Note that this "resynthesis" is just a "quick and dirty" way to quickly hear the results of the pattern extraction. It assumes that the entire sound is a constant repetition of the selected pattern and this has a few important limitations:

1. drift

When the pattern is selected for the specified moment in time, it is inserted before and after that moment as many times as needed to obtain the total duration of the original sound. This means that if there are small tempo fluctuations, the sound will only match very well in the neighborhood of the selected moment in time, but maybe not elsewhere. You might hear this "drifting" when playing back the stereo sound.

2. gaps

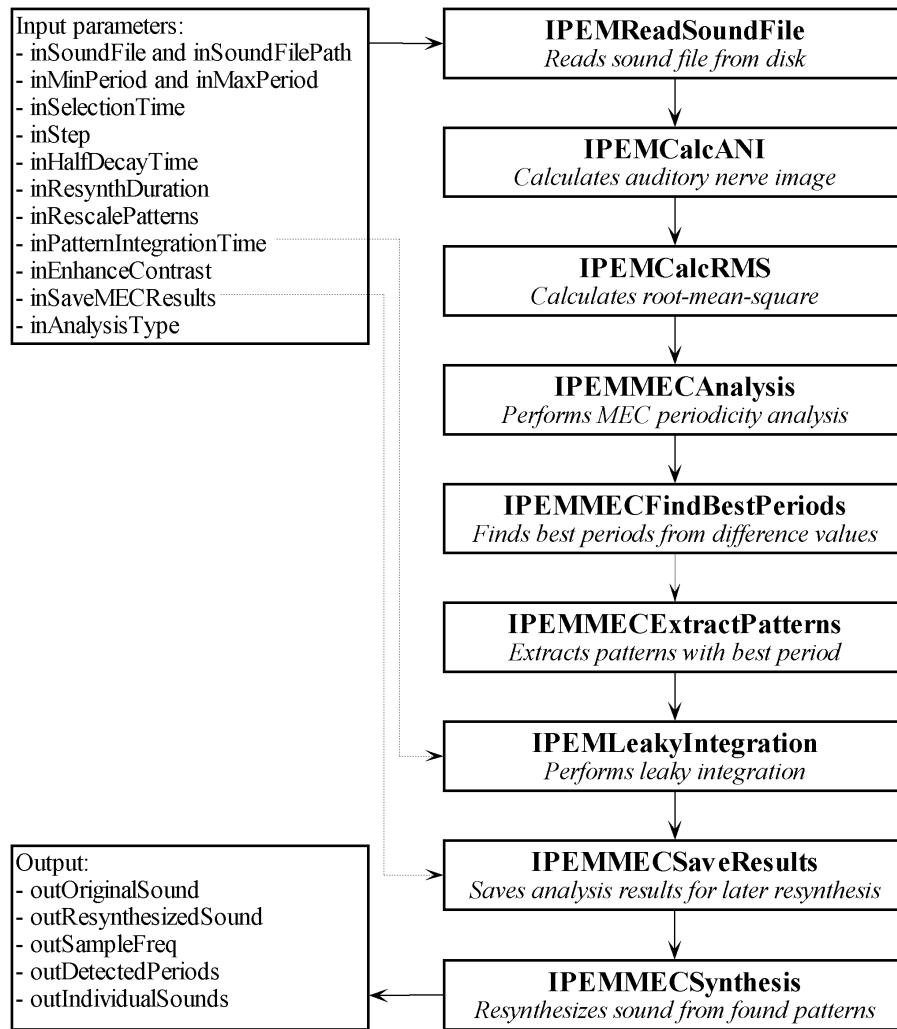


Figure 4.3: Execution flow of IPEMDemoStartMEC

This is a more extreme case of the above: when there is a gap in the rhythmic flow (even if the tempo is constant), the resynthesized sound will be out of sync with the original sound. Again, only at the selected moment in time, there will be perfect synchronization.

3. separated channels

When "separated channels" are used, a best period is found for each channel, so it is perfectly possible that in the higher channels a period was detected different from the one in the lower channels. The resynthesis uses these patterns and you might thus hear different patterns at the same time (each in a different frequency range). This might be good for drum sequences where you have different patterns for hihats and bass drum. If the pattern lengths only differ slightly, again drift might be hearable.

4.3.4 Results and discussion

Sound example 1

Figure 4.4 shows graphic results for the sound example BartokScherzoSuiteOp14.wav ②.

The *top left graph* shows the original sound and the calculated RMS signal levels. It is this RMS signal that is examined by the MEC analysis for repeating patterns.

The *top right graph* shows the difference values calculated by the MEC analysis together with the best period (corresponding to the period for which the difference value is minimal).

The *lower left graph* shows the effect of pattern integration. A leaky integration with half decay time 0.5 s was used, which has the effect that small differences between successive patterns are eliminated in favor of a more stable pattern. Of course, this also "flattens" the detected patterns and "smears" them over time, so an appropriate choice for the half decay time should be chosen.

The *lower right graph* shows the graphical interface that can be used for resynthesizing the rhythm starting from the extracted patterns at a selected moment in time. To hear the result of this "rhythm resynthesis" for this example, you can listen to the following sound:

- Resynthesis from RMS of sound ②

This sound was resynthesized using repetitions of the pattern extracted at 4.5 s, where a best period of 1.604 s was detected. The left channel plays the original sound, while the right channel plays the synthesized sound. The patterns in each channel were rescaled between 0 and 1 (this is a parameter of [IPEMMECEExtractPatterns](#)).

The signals of the synthesized sounds were further processed using a third power function to enhance contrast in the sounds ($s_{enh}(t) = s(t)^3$, where $s(t)$ was first normalized between -1 and 1). This is a parameter of the base function for resynthesis of MEC results [IPEMMECSynthesis](#).

Sound example 2

Figure 4.5 shows graphic results for the sound example TomWaitsBigInJapan.wav ②.

The *top left graph* shows the auditory nerve image (ANI) calculated for 10 bands having a configuration as explained in the interactive demo.

The RMS signal of this ANI is shown in the *top right graph*. It is this signal that is examined by the MEC analysis for repeating patterns.

The *middle left graph* then shows the difference values calculated by the MEC analysis together with the best period (corresponding to the period for which the difference value is minimal). Difference values were calculated for each channel and then summed together to result in the

shown overall difference values.

The *middle right graph* shows the effect of pattern integration. Again, a leaky integration with half decay time 0.5 s was used.

The *lower left graph* again shows the graphical interface that can be used for resynthesizing the rhythm starting from the extracted patterns at a selected moment in time. This resynthesis can be done for each channel separately, but they all use the same "best period". To hear some examples of resynthesis, you can listen to the following sounds:

- Resynthesis for all channels ↗
- Resynthesis for channel 2 ↗
- Resynthesis for channel 9 ↗

The sounds were in this case resynthesized using repetitions of the patterns extracted in each channel at around 12 s, where a best period of 2.133 s was detected. The left channel plays the original sound, while the right channel plays the synthesized sound. Again, re-scaling of patterns and contrast enhancement for the synthesized sounds was applied.

Sound example 3

Figure 4.6 shows graphic results for the sound example PhotekTheLightening.wav ↗.

The *top left graph* shows the auditory nerve image (ANI) calculated for 10 bands having a configuration as explained in the interactive demo.

The RMS signal of this ANI is shown in the *top right graph*. It is this signal that is examined by the MEC analysis for repeating patterns.

The *middle left graph* then shows the difference values calculated by the MEC analysis together with the best period (corresponding to the period for which the difference value is minimal). These results are obtained for each channel, but here only the values for channel 5 are shown as an example. In contrast to the previous example where the difference values were summed together, here each channel is treated separately.

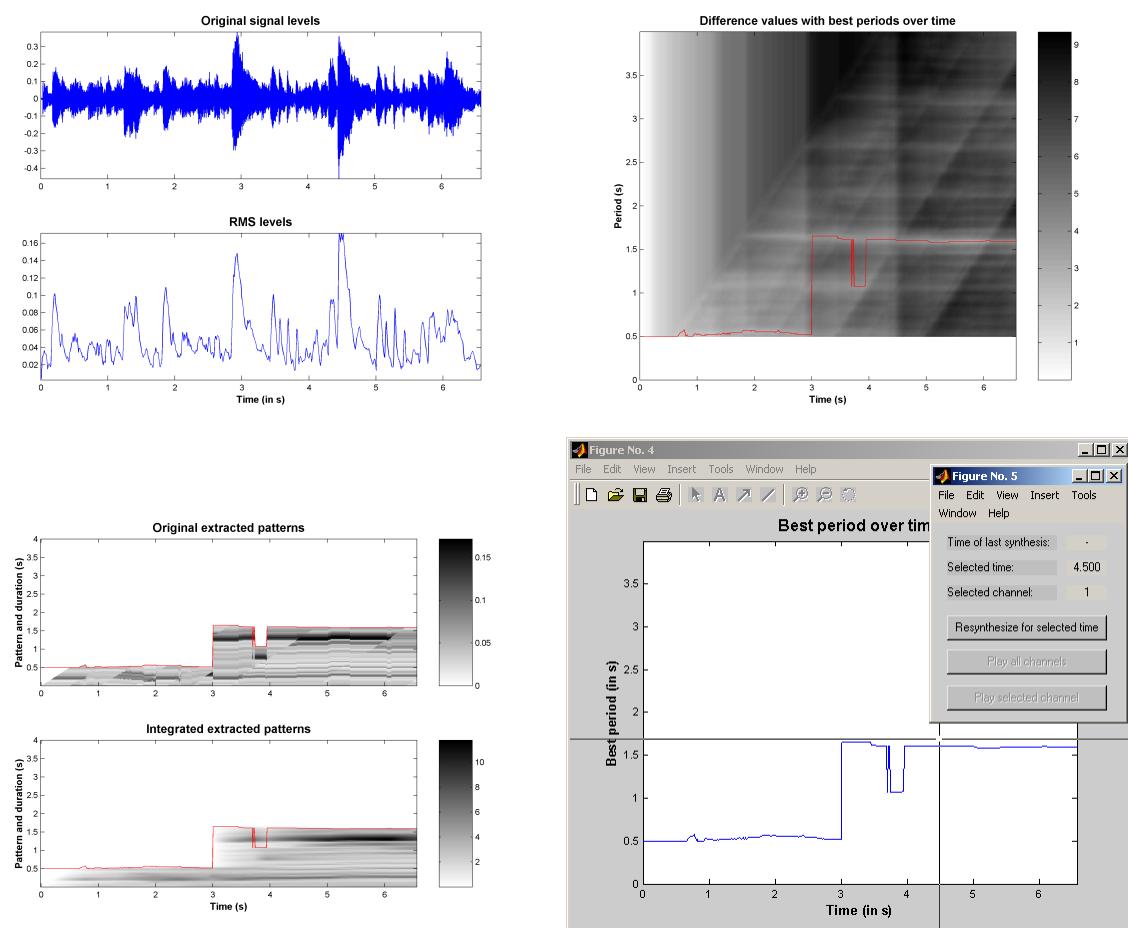
The *middle right graph* shows the best periods over time for all channels. No absolute values are shown, but rather an overview of the stability of the detected periods.

The *lower left graph* shows the effect of pattern integration. As in the other examples, a leaky integration with half decay time 0.5 s was used.

The *lower right graph* again shows the graphical "resynthesis" interface, this time with (possibly) different best periods per channel. The following sounds were synthesized from the extracted patterns:

- Resynthesis for all channels ↴
- Resynthesis for channel 2 ↴
- Resynthesis for channel 9 ↴

These sounds were resynthesized using repetitions of the patterns extracted in each channel at around 8 s. Left and right channel again play original and resynthesized sounds respectively, and pattern scaling and "contrast" enhancement were applied as well.

Figure 4.4: Results for `BartokScherzoSuiteOp14.wav`

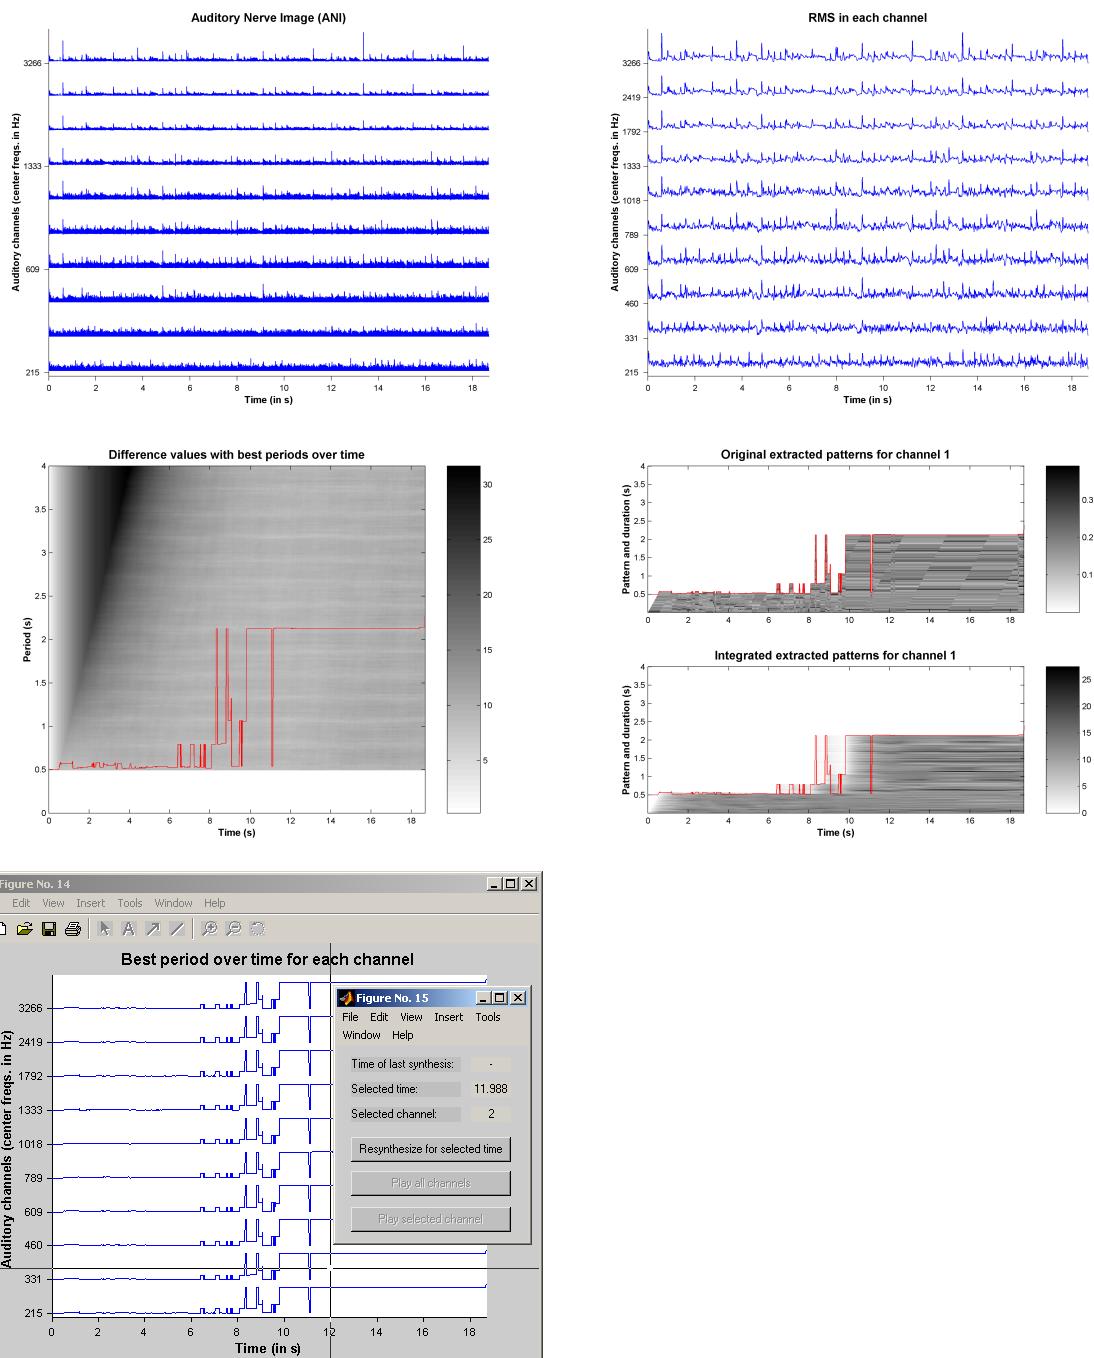


Figure 4.5: Results for TomWaitsBigInJapan.wav

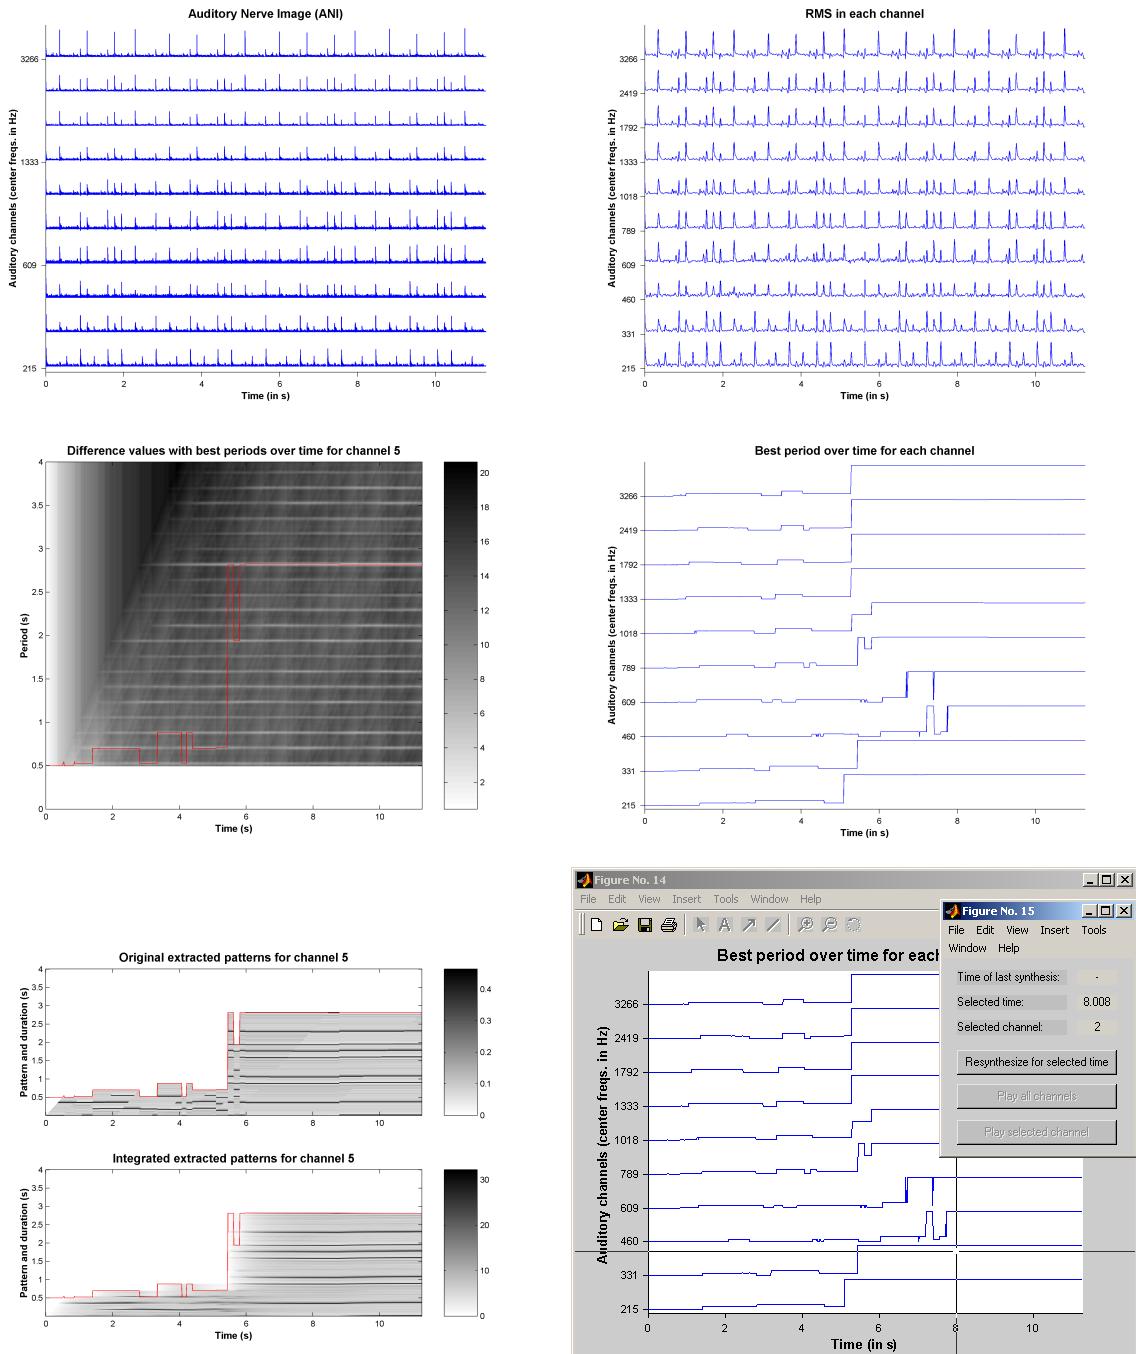


Figure 4.6: Results for PhotekTheLightening.wav

Acronyms

ANI	Auditory Nerve Image
AMDF	Average Magnitude Difference Function
APM	Auditory Peripheral Module
BPF	Band Pass Filter
CM	Contextuality Module
EMM	Echoic Memory Module
FFT	Fast Fourier Transform
HCM	Hair Cell Model
LPF	Low Pass Filter
MEC	Minimal Energy Change
OM	Onset Module
PCM	Pitch Completion Module
RhM	Rhythm Module
RM	Roughness Module
RMS	Root Mean Square

References

- Daniel, P., & Weber, R. (1997). Psychoacoustical roughness: Implementation of an optimized model. *Acustica*, 83, 113-123.
- Greenwood, D., & Joris, P. (1996). Mechanical and 'temporal' filtering as codeterminants of the response by cat primary fibers to amplitude-modulated signals. *Journal of the Acoustical Society of America*, 99(2), 1029-1039.
- Javel, E., McGee, J., Horst, J., & Farley, G. (1988). Temporal mechanisms in auditory stimulus coding. In G. Edelman, W. Gall, & W. Cowan (Eds.), *Auditory function: neurobiological bases of hearing*. New York: John Wiley and Sons.
- Joris, P., & Yin, T. (1992). Responses to amplitude-modulated tones in the auditory nerve of the cat. *Journal of the Acoustical Society of America*, 91(1), 215-232.
- Kemp, S. (1982). Roughness of frequency-modulated tones. *Acustica*, 50, 126-133.
- Krumhansl, C., & Kessler, E. (1982). Tracing the dynamic changes in perceived tonal organization in a spatial representation of musical keys. *Psychological Review*, 89, 334-368.
- Langner, G. (1992). Periodicity coding in the auditory system. *Hearing Research*, 60, 115-142.
- Langner, G. (1997). Temporal processing of pitch in the auditory system. *Journal of New Music Research*, 26, 116-132.
- Langner, G., & Schreiner, C. (1988). Periodicity coding in the inferior colliculus of the cat. Part I. Neuronal mechanisms. *Journal of Neurophysiology*, 60(6), 1799-1822.
- Leman, M. (1999). Naturalistic approaches to musical semiotics and the study of causal musical signification. In I. Zannos (Ed.), *Music and signs – semiotic and cognitive studies in music* (p. 11-38). Bratislava: ASKO Art & Science.
- Leman, M. (2000a). An auditory model of the role of short-term memory in probe-tone ratings. *Music Perception*, 17(4), 481-509.
- Leman, M. (2000b). Visualization and calculation of the roughness of acoustical musical signals using the synchronization index model (SIM). In D. Rochesso (Ed.), *Proceedings of the COST G-6 Conference on Digital Audio Effects (DAFX-00)*, December 7-9. Verona, Italy: University of Verona.
- Leman, M., & Schneider, A. (1997). Origin and nature of cognitive and systematic musicology: An introduction. In M. Leman (Ed.), *Music, Gestalt, and computing - studies in cognitive and systematic musicology* (p. 13-29). Berlin, Heidelberg: Springer-Verlag.
- Leman, M., & Verbeke, B. (2000). The concept of minimal 'energy' change (mec) in relation to fourier transform, auto-correlation, wavelets, amdf, and brain-like timing networks – application to the recognition of repetitive rhythmical patterns in acoustical musical signals. In K. Jokinen, D. Heylen, & A. Nijholt (Eds.), *Proceedings of the workshop on Internalising Knowledge, Ieper 22-24 november* (p. pp. 191-200). Ieper, Belgium: Cele-Twente.
- Marr, D. (1982). *Vision: a computational investigation into the human representation and processing of visual information*. San Francisco: W. H. Freeman and Company.
- Sethares, W. (1998). *Tuning, timbre, spectrum, scale*. Berlin Heidelberg, New York: Springer-Verlag.
- Smith, L. S. (1996). Onset-based sound segmentation. *Neural Information Processing Systems* 8.
- Smith, R. (1988). Encoding of sound intensity by auditory neurons. In G. Edelman, W. Gall, & W. Cowan (Eds.), *Auditory function: neurobiological bases of hearing*. New York: John Wiley and Sons.

- Terhardt, E. (1997). *Akustische Kommunikation*. Berlin, Heidelberg: Springer.
- Van Immerseel, L., & Martens, J. (1992). Pitch and voiced/unvoiced determination with an auditory model. *The Journal of the Acoustical Society of America*, 91, 3511–3526.

Part III

Reference manual

Chapter 5

Introduction

The "IPEM Toolbox, Toolbox for perception-based music analysis" is a toolbox for MATLAB developed by the IPEM as part of a research methodology in the domain of musical content extraction.

It contains functions for characterizing and examining features present in acoustical signals and prerecorded musical sounds (audio signals).

The IPEM Toolbox can be useful for researchers who are interested in musical content analysis and automatic extraction and description of musical features.

It is our aim to provide these people with functionality for dealing with different aspects of feature extraction in the field of music perception such as: chords and tonality, pitch, roughness, onset detection, beat and meter extraction, timbre characteristics.

Some of the functions in this toolbox are general purpose signal processing functions, whereas others are implementations of specific algorithms developed by the IPEM.

They all obey the 'black box' principle: input → process → output, and they are kept as general as possible (hence the sometimes great number of input arguments). Scripts are only provided for demonstrating the use of some functions.

Apart from a set of functions for general use, the toolbox will also contain some functions which are written specifically for a particular demo. You can find the entire list of functions in the [Function Reference](#).

Chapter 6

General information

6.1 Conditions/Disclaimer

All of the following applies to the "IPEM Toolbox, Toolbox for perception-based music analysis" (Version: 1.02 (beta)), of which this manual is a part.

The copyright holder of this toolbox and its manual is Ghent University.

The code is distributed under the conditions of the GNU GPL license (see the ReadMe.txt and gpl.txt file in the code tree for details, including conditions and disclaimer).

This manual is distributed under the conditions of the GNU FDL license for which the full verbatim text is shown in the following section.

Full contact info can be found on the front page of this manual.

6.2 GNU Free Documentation License

GNU Free Documentation License
Version 1.2, November 2002

Copyright (C) 2000,2001,2002 Free Software Foundation, Inc.
59 Temple Place, Suite 330, Boston, MA 02111-1307 USA
Everyone is permitted to copy and distribute verbatim copies
of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document "free" in the sense of

freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or noncommercially. Secondarily, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of "copyleft", which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The "Document", below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as "you". You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A "Modified Version" of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A "Secondary Section" is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the

subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The "Invariant Sections" are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For

works in formats which do not have any title page as such, "Title Page" means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

A section "Entitled XYZ" means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as "Acknowledgements", "Dedications", "Endorsements", or "History".) To "Preserve the Title" of such a section when you modify the Document means that it remains a section "Entitled XYZ" according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and

Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title

distinct

from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.

B. List on the Title Page, as authors, one or more persons or entities

responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.

C. State on the Title page the name of the publisher of the Modified Version, as the publisher.

D. Preserve all the copyright notices of the Document. E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.

F. Include, immediately after the copyright notices, a license notice

giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.

G. Preserve in that license notice the full lists of Invariant Sections

and required Cover Texts given in the Document's license notice.

H. Include an unaltered copy of this License. I. Preserve the section Entitled "History", Preserve its Title, and add

to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.

J. Preserve the network location, if any, given in the Document for

public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.

K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.

L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers

- or the equivalent are not considered part of the section titles.
- M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements"
 - or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties--for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your

combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements".

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an "aggregate" if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which

are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled "Acknowledgements", "Dedications", or "History", the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided for under this License. Any other attempt to copy, modify, sublicense or distribute the Document is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation.

ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

Copyright (c) YEAR YOUR NAME.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled "GNU Free Documentation License".

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the "with...Texts." line with this:

with the Invariant Sections being LIST THEIR TITLES, with the Front-Cover Texts being LIST, and with the Back-Cover Texts being LIST.

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License,

to permit their use in free software.

6.3 System requirements

Hardware

- same as the hardware requirements for your Matlab version on PC
- as much memory and processing speed as you can get (preferably 128 MB and 500 MHz or better)

Software

- same as the software requirements for your Matlab version on PC (only tested on Windows 95/98/NT with SP4 and Windows2000 though...)
- Matlab 6.0 (R12)
- Matlab Signal Processing Toolbox 5.0

Note:

The Matlab 5.3.1 (R11.1) version is no longer maintained. We shipped the code with the Matlab 6.0 version of IPEMPeriodicityPitch.m . If you want to use the package with Matlab 5.3.1, you'll need to edit that file on lines 170 and 173, and use the appropriate auditory model dll for 5.3.1. There is a separate package containing all the necessary files to build that dll.

6.4 Installation

To install the IPEM Toolbox on your computer, perform the following steps:

1. Unzip all files in the IPEMToolbox.zip file to a directory called IPEMToolbox.
2. Start Matlab, go to 'Set Path...' in the 'File' menu and add the IPEMToolbox directory to the Matlab search path.
3. Type the following line at the Matlab command prompt:
`IPEMSetup;`
 You should now see: `Initializing IPEM Toolbox...` and then: `Done.`
4. Set your preferred 'input' and 'output' root directories by typing:
`setpref('IPEMToolbox','RootDir_Input','your input directory');`
`setpref('IPEMToolbox','RootDir_Output','your output directory');`
 where 'your input directory' is the complete path to your preferred input directory between single quotes (same for output directory).
 If you don't do this, IPEMToolbox\Temp will be used as both input and output directory.

Remark: Each user on your Windows machine needs to perform steps 3. and 4. once to initialize the IPEM Toolbox and set his/her preferred input and output directory.

6.5 Reporting problems

Although this toolbox has been used by the IPEM for quite some time now, the possibility of running into bugs always exists. Therefore, if you ever run into a bug or malfunction, we would highly appreciate it if you would let us know.

However, before doing so, please read the following lines:

- Make sure you're using the latest version of the IPEM Toolbox !
You can download it from: www.ipem.ugent.be/Toolbox
- Check whether you're using the function in the correct way !
For more info, look up the function in the **Function Reference** or type `help` followed by the function's name at the Matlab command prompt.
- If you just started to use a newer version of the IPEM Toolbox, and something that was previously working fine doesn't work any longer, then make sure to read the 'Changes.txt' file and check for changes that could possibly affect your programs/functions as far as they are based on the IPEM Toolbox functions. Normally, this type of problem should rarely occur, since we always try to keep the code backwards compatible between different versions...

To report a problem, please send an email to [toolbox\[at\]ipem\[dot\]ugent\[dot\]be](mailto:toolbox[at]ipem[dot]ugent[dot]be) containing at least the following information:

1. General info

Your coordinates, affiliation, ...

2. IPEM Toolbox version

The version of the IPEM Toolbox can be found by typing `IPEMToolboxVersion` in Matlab, or by looking at the first two lines of the file 'ReadMe.txt'.

3. System specifications

This means: processor type, amount of RAM, operating system, Matlab version

4. Function name

The name of the function in which the bug or malfunction occurs.

5. Error or warning message

If applicable, add the text of the error or warning message to your email.

6. Description

Describe the bug or malfunction in your own words. Try to be as clear and complete as possible:

- what exactly is going wrong ?
- are you using any special/extreme input arguments ?
- how are you using the IPEM function ?

- can you reproduce the bug/malfunction ? if so, how ?

If possible, save your Matlab environment to a .mat file and attach this to your email together with your Matlab code and any input files you are processing (only do this if your attachments are smaller than 1 MB in total, otherwise send us a short email first to receive more instructions).

Note that the IPEM Toolbox project has finished now and that we do not officially provide any support at all, however we might do an effort to answer your questions if you ask politely ;-)

6.6 Updates, ideas and remarks

Since we are constantly expanding this toolbox to suit our research needs, many functions will be added in the future. As already mentioned above, you can find the latest version of the toolbox at: www.ipem.ugent.be/Toolbox

If you ever feel the need for new functionality, please tell us about it, so that we can take your ideas into account for newer versions. Of course, if we incorporate your ideas or algorithms into a new function, you will be listed in the 'Authors' section of the function's documentation.

Also, if you have any remarks about existing functions, or about the way we are doing things, please let us know.

6.7 About the IPEM Toolbox concepts

In the first half of this document there is a part called '**Concepts**', in which you will find more fundamental documentation about the concepts behind the Matlab functions defined in the IPEM Toolbox.

Whereas this Reference Manual is meant as a practical guide for working with the IPEM Toolbox itself, the 'Concepts' part of this document deals with the fundamental ideas that led to the implementation of the toolbox. It defines a conceptual framework based on auditory modeling principles, psychoacoustics and experimental data. You should definitely read this if you want to fully understand the models used in this toolbox !

Chapter 7

About the documentation

7.1 Documentation structure

The description for each function is always structured in the following way:

Function's name

Usage:

This shows the syntax for using the function: what and how many input arguments and outputs are there, and what is their order ?

Description:

This shows the semantics of the function:
what is it intended for and/or how does it work ?

Input arguments:

A description of the input arguments together with a specification of whether they are optional or not, and what their defaults are.

Output:

A description of the outputs of the function.

Remarks:

Any additional remarks that are not in the other sections.

Example:

A 'copy-and-paste' example of how to use the function.

Authors:

A list of the people who wrote (parts of) the code of the function, or have made some changes to it. Accompanied by the date of their most recent changes. As always, dates are written in the following format: YYYYMMDD (year,month,day)

7.2 Used terminology

7.2.1 Prefixes

In the function declarations, input arguments are always starting with the prefix 'in', while output arguments are always starting with the prefix 'out'. In some cases 'io' is used for arguments that are used for both input and output.

7.2.2 Signal and SampleFreq

Signal is a general name for any matrix representing a time-varying entity in possibly more than one dimensions.

Single-channel signals are represented by a row-vector. Multi-channel signals are represented by a matrix in which each row represents one channel. Each column thus represents a moment in time.

A signal is always accompanied by a sample frequency, called something like SampleFreq and is specified in Hertz (Hz).

7.2.3 ANI, ANIFreq and ANIFilterFreqs

ANI is an abbreviation of Auditory Nerve Image and is used for a (multi-channel) signal coming out of the auditory model. ANIFreq specifies the sample frequency of such an ANI, and ANIFilterFreqs specifies the center frequencies of the band pass filters that were used for calculating the ANI.

7.2.4 FileName and FilePath

These are used for functions that take input from / generate output to a file.

FileName is a Matlab string representing the name of the file, while FilePath is a Matlab string representing the complete directory path to the file.

For example:

```
FileName = 'music.wav'  
FilePath = 'E:\Tests\Sounds'
```

7.2.5 Empty and not specified input arguments

Empty input arguments

Sometimes, input arguments can be left empty, which means you can specify the empty matrix [] for them. If possible, a default value will be used in such a case.

Optional input arguments

Arguments at the end of the input arguments list do not always have to be specified, which means you do not have to type a value for them at all. This again is mostly used for allowing the caller to use default values.

7.2.6 PlotFlag

Always used as inPlotFlag, this argument specifies whether or not plots should be generated by the function.

The following convention is used: if it is zero, no plots are generated, if it is non-zero then plots are generated.

Chapter 8

Function reference

The following pages contain the entire list of functions that are written for general use (and in future versions also the ones that were written to setup a particular demo).

The information provided here is essentially the same as the one you would get from typing `help` followed by the function's name at the Matlab command prompt.

General functions

IPEMAMTone

Usage:

```
[outSignal] = IPEMAMTone(inCarrierFreq,inModulationFreq,inModulationDepth,...  
inDuration,indBLevel,inSampleFreq,inFadeInOut)
```

Description:

Generates an amplitude modulated tone according to:

$$s(t) = (1 + \text{inModulationDepth} \cdot \sin(2\pi \cdot \text{inModulationFreq} \cdot t)) \\ \cdot \sin(2\pi \cdot \text{inCarrierFreq} \cdot t)$$

Input arguments:

inCarrierFreq = carrier frequency (in Hz)
 inModulationFreq = modulation frequency (in Hz)
 inModulationDepth = modulation depth
 inDuration = wanted duration (in s)
 indBLevel = wanted dB level (in dB)
 if empty or not specified, -20 dB is used by default
 inSampleFreq = wanted sample frequency (in Hz)
 if empty or not specified, 22050 is used by default
 inFadeInOut = duration of linear fade in and out of the sound signal (in s)
 given as [FadeIn FadeOut] or as a scalar (FadeIn == FadeOut)
 if empty or not specified, 0.010 is used by default

Output:

outSignal = generated sound signal

Example:

```
s = IPEMAMTone(1000,70,0.65,0.200,-20,44100,0.010);
```

Authors:

Koen Tanghe - 20040323

IPEMAadaptLevel

Usage:

```
outSignal = IPEMAadaptLevel (inSignal,indB)
```

Description:

This function adapts the RMS power level of the (multi-channel) signal to the specified dB level. Channels are represented by rows.

Input arguments:

inSignal = the input signal (each row represents a channel)
indB = wanted level (0 is maximum level without clipping, -20 is reasonable)

Output:

outSignal = the adapted signal

Remarks:

The reference value of 0 dB is the level of a square wave with amplitude 1
Thus, a sine wave with amplitude 1 yields -3.01 dB.

Example:

```
Signal = IPEMAadaptLevel(Signal,-6);
```

Authors:

Marc Leman
Koen Tanghe - 20000626

IPEMAnimateSlices

Usage:

```
IPEMAnimateSlices(inSignal,inSampleFreq,XAxisLabel,YAxisLabel, ...
    inXData,inXRange,inYRange,inTimeRange, ...
    inTimeStep,inAnimationInterval,inScales,inGraphicsHandle)
```

Description:

Displays animated plots of a 2D matrix where each column represents a slice of data at a certain moment in time.
Could be used as an alternative to a surf plot...

Input arguments:

```
inSignal = 2D vector in which each column contains a slice of data
inSampleFreq = sample frequency of inSignal (in Hz)
inXAxisLabel = label to display on the X axis of the plot
    if empty or not specified '' is used by default
inYAxisLabel = label to display on the Y axis of the plot
    if empty or not specified '' is used by default
inXData = vector containing the values that correspond with each row of the
    2D vector
    if empty or not specified, 1:size(inSignal,1) is used by default
inXRange = X range for which values need to be displayed in each plot
    specified as a vector [LowXValue HighXValue]
    if empty or not specified, [min(inXData) max(XData)] is used
    by default
inYRange = Y range for which values will be visible in the plot, specified
    as a vector [MinYValueMaxYValue]
    if empty or not specified, [min(min(inSignal)) max(max(inSignal))]
    is used by default
inTimeRange = time range that should be animated, specified as a vector
    [StartTime EndTime] (in s)
    if empty or not specified, the entire signal is animated
    by default (so: [0 (size(inSignal,2)-1)/inSampleFreq])
inTimeStep = time step to use for stepping through the animated data (in s)
    if empty or not specified, 1/inSampleFreq is used by default
inAnimationInterval = time to wait between two successive plots in
    the animation (in s)
    if -1, real-time is used
    if empty or not specified, 0 is used by default
inScales = 2 element cell vector specifying the type of scale to use for the
    X and Y axis specified as {XType YType}, where XType and YType can
```

```
    be either 'linear' or 'log'  
    if empty or not specified, {'linear' 'linear'} is used by default  
inGraphicsHandle = handle of a subplot or a figure in which the animation  
    should be shown  
    if empty or not specified, a new subplot in a new figure  
    will be used by default
```

Remarks:

Keys that can be pressed in the figure while the plot is animated:
'+' = forward
'-' = backward
' ' = pause
'x' = exit

Example:

```
[s,fs] = IPEMReadSoundFile('schum1.wav');  
[S,T,F] = IPEMCalcSpectrogram(s,fs,0.040,0.010);  
IPEMAnimateSlices(abs(S),1/T(2),'Frequency (in Hz)', 'Amplitude',F);
```

Authors:

Koen Tanghe - 20021021

IPEMBatchExtractSoundFragment

Usage:

```
IPEMBatchExtractSoundFragment(inFragment,inInputDirectory,inFilePattern, ...
    inOutputDirectory,inNewNameFormat, ...
    inCreateDir,inFadeInOutTime, ...
    inPreferNegativeTimes)
```

Description:

Extracts a fragment from all sound files in a specific directory matching the specified wildcard pattern, and saves each fragment to the specified output directory appending a suffix.

Input arguments:

inFragment = Fragment that should be extracted.
 Specified as [StartTime EndTime] where both elements are in s.
 If StartTime < 0, it is referenced from the end of file.
 If EndTime <= 0, it is referenced from the end of the file.

inInputDirectory = Directory where the sound files are located.

inFilePattern = Wildcard pattern for the files that should be processed.
 If empty or not specified, '*.wav' is used by default.

inOutputDirectory = Directory where the extracted sound fragments should be written. If empty or not specified, inInputDirectory is used by default.

inNewNameFormat = Format string that should be used for the file names of the new sound fragments. This must be a format string where the first type specifier is a %s (will be replaced by the original base file name), and the second and third are a %f (will be replaced by the start resp. end time in seconds).
 If empty or not specified, '%s_%gs_%gs' is used by default.

inCreateDir = If 1, the output directory is created if it doesn't exist yet.
 If empty or not specified, 1 is used by default.

inFadeInOutTime = Time over which the beginning and end should be faded in repectively out (in s).
 If empty or not specified, 0.010 s is used by default.

inPreferNegativeTimes = If 0, negative start or end times will be shown as the actually used times in the output file names.
 Otherwise, negative times will be kept.
 If empty or not specified, 1 is used by default.

Example:

```
IPEMBatchExtractSoundFragment([60 90], 'D:\Temp\', '*.wav', 'D:\Temp2\');
```

Authors:

Koen Tanghe - 20040702

IPEMBellShape

Usage:

```
outY = IPEMBellShape (inX,inCenter,inWidth,inPeak)
```

Description:

This function generates a bell shaped curve.

Input arguments:

```
inX = input data points  
inCenter = center of the curve  
inWidth = width of the curve:  
          the value at (inCenter + inWidth) is 10% of inPeak  
inPeak = maximum value (at the center of the curve)
```

Output:

```
outY = curve values for the input
```

Example:

```
Y = IPEMBellShape(0:1000,500,100,1);
```

Authors:

Koen Tanghe - 20000208

IPEMBlockDC

Usage:

```
[outSignal,outSignalFreq] = ...
IPEMBlockDC(inSignal,inSignalFreq,inCutoffFreq,inPlotFlag)
```

Description:

Blocks DC (very low frequencies).
 Based on the formula:
 $y(n) = g * (x(n) - x(n-1)) + R * y(n-1)$
 where
 $R = 1 - (2\pi F_c / F_s)$
 $g = (1 + R)/2$ (for gain correction)
 F_c = cutoff frequency in Hz (-3 dB point)
 F_s = sample frequency in Hz

Input arguments:

```
inSignal = one dimensional input signal
inSignalFreq = sample frequency of inSignal (in Hz)
inCutoffFreq = -3 dB cutoff frequency for HP filter (in Hz)
inPlotFlag = if non-zero, plots are generated
            if not specified or empty, 0 is used by default
```

Output:

```
outSignal = Processed signal.
outSignalFreq = Sample frequency of output signal (same as input) (in Hz).
```

Example:

```
[s2,fs2] = IPEMBlockDC(s,fs,50,1);
```

Authors:

Koen Tanghe - 20050120

IPEMCalcANI

Usage:

```
[outANI,outANIFreq,outANIFilterFreqs] = ...
    IPEMCalcANI (inSignal,inSampleFreq,inAuditoryModelPath, ...
        inPlotFlag,inDownsamplingFactor, ...
        inNumOfChannels,inFirstCBU,inCBUStep)
```

Description:

This function calculates the auditory nerve image for the given signal.

Input arguments:

```
inSignal = the sound signal to be processed
inSampleFreq = the sample frequency of the input signal (in Hz)
inAuditoryModelPath = path to the working directory for the auditory model
    if empty or not specified, IPEMRootDir('code')\Temp
    is used by default
inPlotFlag = if non-zero, plots the ANI
    if empty or not specified, 0 is used by default
inDownsamplingFactor = the integer factor by which the outcome of the
    auditory model is downsampled
    (use 1 for no downsampling)
    if empty or not specified, 4 is used by default
inNumOfChannels = number of channels to use
    if empty or not specified, 40 is used by default
inFirstCBU = frequency of first channel (in critical band units)
    if empty or not specified, 2.0 is used by default
inCBUStep = frequency difference between channels (in cbu)
    if empty or not specified, 0.5 is used by default
```

Output:

```
outANI = a matrix of size [N M] representing the auditory nerve image,
    where N is the number of channels (currently 40) and
    M is the number of samples
outANIFreq = sample freq of ANI (in Hz)
outANIFilterFreqs = center frequencies used by the auditory model (in Hz)
```

Remarks:

The outcome of the auditory model normally has a sample frequency of 11025 Hz, but for most calculations this can be downsampled to a lower value (11025/4 is the default).

Example:

```
[ANI,ANIFreq,ANIFilterFreqs] = IPEMCalcANI(Signal,SampleFreq,[],1);
```

Authors:

Koen Tanghe - 20010129

IPEMCalcANIFromFile

Usage:

```
[outANI,outANIFreq,outANIFilterFreqs] = ...
    IPEMCalcANIFromFile (inFileName,inFilePath,inAuditoryModelPath, ...
        inPlotFlag,inDownsamplingFactor, ...
        inNumOfChannels,inFirstCBU,inCBUStep)
```

Description:

This function calculates the auditory nerve image for the given sound file.

Input arguments:

```
inFileName = the sound file to be processed
inFilePath = the path to the sound file
            if empty or not specified, IPEMRootDir('input')\Sounds
            is used by default
inAuditoryModelPath = path to the working directory for the auditory model
            if empty or not specified, IPEMRootDir('code')\Temp
            is used by default)
inPlotFlag = if non-zero or empty, plots the ANI
            if empty or not specified, 0 is used by default
inDownsamplingFactor = the integer factor by which the outcome of the
            auditory model is downsampled
            (use 1 for no downsampling)
            if empty or not specified, 4 is used by default
inNumOfChannels = number of channels to use
            if empty or not specified, 40 is used by default
inFirstCBU = frequency of first channel (in critical band units)
            if empty or not specified, 2.0 is used by default
inCBUStep = frequency difference between channels (in cbu)
            if empty or not specified, 0.5 is used by default
```

Output:

```
outANI = a matrix of size [N M] representing the auditory nerve images,
        where N is the number of channels (currently 40) and
        M is the number of samples
outANIFreq = sample freq of ANI
outANIFilterFreqs = center frequencies used by the auditory model
```

Remarks:

The outcome of the auditory model normally has a sample frequency of 11025 Hz, but for most calculations this can be downsampled to a lower value (11025/4 is the default).

Example:

```
[ANI,ANIFreq,ANIFilterFreqs] = IPEMCalcANIFromFile ('music.wav',[],[],1);
```

Authors:

Koen Tanghe - 20010129

IPEMCalcCentroid

Usage:

```
outCentroid = IPEMCalcCentroid (inWeights,inWeightFreq, ...
                                inFrameWidth,inFrameInterval,inDistances, ...
                                inPlotFlag)
```

Description:

Calculates the centroid for the given (time-varying) weights and their (constant) distances.

This can be used to calculate the (time-varying) centroid of the spectrum of a musical signal.

Input arguments:

inWeights = a matrix (of size [N M]) in which each column represents weight data at a given moment in time
 inWeightFreq = frequency at which the weight data is sampled (in Hz)
 inFrameWidth = width of 1 frame (in s)
 inFrameInterval = interval between successive frames (in s)
 inDistances = a vector of size [N 1], representing distances for each weight
 if empty or not specified, (1:N)' is used
 inPlotFlag = if non-zero, plots the centroid
 if empty or not specified, 0 is used by default

Output:

outCentroid = the (time-varying) centroid of the weights
 outCentroidFreq = sample frequency for centroid

Example:

```
[Centroid,CentroidFreq] = IPEMCalcCentroid(ANI,ANIFreq,0.05,0.01, ...
                                             ANIFilterFreqs);
```

Authors:

Koen Tanghe - 20010221

IPEMCalcCentroidWidth

Usage:

```
[outWidth,outWidthFreq] =
    IPEMCalcCentroidWidth(inWeights,inWeightFreq,
                          inFrameWidth,inFrameInterval,
                          inDistances,inCentroid,
                          inPlotFlag)
```

Description:

Calculates the weighted difference between the spectral components and the (already) calculated centroid of the given multi-channel signal.

Input arguments:

inWeights = a matrix (of size [N M]) in which each column represents weight data at a given moment in time
 inWeightFreq = frequency at which the weight data is sampled (in Hz)
 inFrameWidth = width of 1 frame (in s)
 inFrameInterval = interval between successive frames (in s)
 inDistances = a vector of size [N 1], representing distances for each weight if empty or not specified, (1:N)' is used
 inCentroid = this provides the centroid for the data that was calculated with IPEMCalcCentroid, using the same parameters if empty or not specified, the centroid itself is calculated in this function
 inPlotFlag = if non-zero, plots three curves:
 the centroid itself (central line), plus both the centroid + the centroid width and the centroid - the centroid width
 if empty or not specified, 0 is used by default

Output:

outWidth = the calculated width
 outWidthFreq = sample frequency of the width

Example:

```
[Width,WidthFreq] = IPEMCalcCentroidWidth(ANI,ANIFreq,0.05,0.01,....
                                         ANIFilterFreqs);
```

Authors:

Koen Tanghe - 20010221

IPEMCalcFFT

Usage:

```
[outAmpl,outPhase,outFreqs] = IPEMCalcFFT (inSignal, inSampleFreq, ...
                                             inFFTSize, inPlotFlag, ...
                                             inUseLogScale, inShowPhase)
```

Description:

This function calculates (and shows) the amplitude and phase of a real signal.

Input arguments:

inSignal = the signal to be analyzed
 inSampleFreq = the frequency at which the signal was sampled (in Hz)
 inFFTSize = size of the FFT
 if empty or not specified, the power of two nearest to the full length of the signal is used by default
 inPlotFlag = if non-zero, shows FFT plot
 if empty or not specified, 1 is used by default
 inUseLogScale = if non-zero, a logarithmic scale is used for the frequency
 if empty or not specified, 1 is used by default
 inShowPhase = if non-zero, the phase is shown as well
 if empty or not specified, 0 is used by default

Output:

outAmpl = amplitude of fft components
 outPhase = phase of fft components
 outFreqs = frequencies used for fft calculation

The indices in these output matrices (of length L) correspond to the following frequencies:

index 1 corresponds to the DC component
 index i corresponds to frequency $(i-1)*inSampleFreq/inFFTSize$
 index L corresponds to the Nyquist frequency (if inFFTSize was even)

Remarks:

If inFFTSize is smaller than the length of inSignal, inSignal is truncated (a warning is issued when this happens).

If inFFTSize is bigger than the length of inSignal, inSignal is padded with zeroes.

Example:

```
[Ampl,Phase,Freqs] = IPEMCalcFFT (Signal,SampleFreq,1024,1,1,0);
```

Authors:

Koen Tanghe - 20010627

IPEMCalcFlux

Usage:

```
outFlux = IPEMCalcFlux (inSignal,inSampleFreq,inPlotFlag)
```

Description:

Calculates the flux of the signal, where the flux is defined as the norm of the difference (vector) between the current and the previous values of the (multichannel) signal.

Input arguments:

```
inSignal = input signal (spectral data, for example)  
inSampleFreq = sample frequency of inSignal (in Hz)  
inPlotFlag = if non-zero, plots are generated  
              if not specified, 1 is used by default
```

Output:

```
outFlux = calculated flux
```

Example:

```
Flux = IPEMCalcFlux (Signal,SampleFreq);
```

Authors:

Koen Tanghe - 20010906

IPEMCalcMeanAndVariance

Usage:

```
[outMean,outVariance,outFreq] =
IPEMCalcMeanAndVariance(inSignal,inSampleFreq,
                         inFrameWidth,inFrameInterval,
                         inPlotFlag,inPlotTitle)
```

Description:

Calculates 'running' mean and variance of multi-channel signal:
for each channel, the mean and variance is calculated within successive frames.

Input arguments:

```
inSignal = signal to be analyzed
inSampleFreq = sample frequency of the input signal (in Hz)
inFrameWidth = width of 1 frame (in s)
inFrameInterval = interval between successive frames (in s)
inPlotFlag = if non-zero, plots are generated
              if not specified or empty, 1 is used by default
inPlotTitle = title for the plot
              if not specified or empty, no title is shown
```

Output:

```
outMean = 'running' mean of inSignal
outVariance = 'running' variance of inSignal
outFreq = sample frequency for both outMean and outVariance
```

Example:

```
[Mean, Variance, Freq] = IPEMCalcMeanAndVariance(RMS,RMSFreq,0.05,0.01,1, ...
'RMS');
```

Authors:

Koen Tanghe - 20001012

IPEMCalcNoteFrequency

Usage:

```
outFrequency =
    IPEMCalcNoteFrequency(inNoteNr,inOctaveNr,
                           inRefNoteNr,inRefOctaveNr,
                           inRefFreq,inNotesPerOctave,
                           inOctaveRatio)
```

Description:

Calculates the frequency of a note.

Supports non-standard tone scales that can be calculated like this:

$$\text{Frequency} = \text{RefFreq} * \text{OctaveRatio}^{\text{Exponent}}$$

where:

$$\text{Exponent} = (\text{OctaveNr}-\text{RefOctaveNr}) + (\text{NoteNr}-\text{RefNoteNr})/\text{NotesPerOctave}$$

Input arguments:

inNoteNr = the rank number of the note (for example: 1 for C, 2 for C#, ...)
 For the standard octave division (12 notes, ref. A4), this can
 also be a note string like 'A#' (in which case inOctaveNr must
 still be specified), or 'A#4' (in which case inOctaveNr can be
 omitted or left empty). Both sharps (#) and flats (b) can be used.

inOctaveNr = the octave number of the note
 if empty or not specified, inNoteNr is assumed to contain the
 octave specification

inRefNoteNr = the rank number of the reference note
 if empty or not specified, 10 is used by default

inRefOctaveNr = the octave number of the reference note
 if empty or not specified, 4 is used by default

inRefFreq = the frequency of the reference note (in Hz)
 if empty or not specified, 440 Hz is used

inNotesPerOctave = the number of notes in one octave
 if empty or not specified, 12 is used by default

inOctaveRatio = the frequency ratio between two octaves
 if empty or not specified, 2 is used

Output:

`outFrequency = the frequency for the note (in Hz)`

Example:

```
Frequency = IPEMCalcNoteFrequency('A#5');
```

Authors:

Koen Tanghe - 20000509

IPEMCalcOnsets

Usage:

```
[outOnsetSignal,outOnsetFreq] =  
    IPEMCalcOnsets(inSignal,inSampleFreq,inPlotFlag)
```

Description:

This function calculates the onsets for the given signal using the auditory model and an integrate-and-fire neural net layer.

Input arguments:

```
inSignal = signal to be processed  
inSampleFreq = sample frequency of input signal (in Hz)  
inPlotFlag = if non-zero, a plot is generated at the end  
            if empty or not specified, 0 is used by default
```

Output:

```
outOnsetSignal = signal having a non-zero value for an onset, and zero  
                otherwise (the higher the non-zero value, the more our  
                system is convinced that the onset is really an onset)  
outOnsetFreq = sample frequency of outOnsetSignal (in Hz)
```

Example:

```
[OnsetSignal,OnsetFreq] = IPEMCalcOnsets(s,fs);
```

Authors:

Koen Tanghe - 20030327

IPEMCalcOnsetsFromANI

Usage:

```
[outOnsetSignal,outOnsetFreq] =  
    IPEMCalcOnsetsFromANI(inANI,inANIFreq,inPlotFlag)
```

Description:

This function calculates the onsets for a given signal represented by its auditory nerve image, using an integrate-and-fire neural net layer.

Input arguments:

```
inANI = auditory nerve image to be processed  
inANIFreq = sample frequency of auditory nerve image (in Hz)  
inPlotFlag = if non-zero, a plot is generated at the end  
            if empty or not specified, 1 is used by default
```

Output:

```
outOnsetSignal = signal having a non-zero value for an onset, and zero  
                otherwise (the higher the non-zero value, the more our  
                system is convinced that the onset is really an onset)
```

Example:

```
[OnsetSignal,OnsetFreq] = IPEMCalcOnsetsFromANI(ANI,ANIFreq);
```

Authors:

Koen Tanghe - 20010122

IPEMCalcPeakLevel

Usage:

```
[outPeakSignal,outPeakFreq] = ...
    IPEMCalcPeakLevel(inSignal,inSampleFreq,inFrameWidth,
                      inFrameInterval,inUseAbs,inPlotFlag)
```

Description:

This function calculates the running maximum of the given signal: a maximum is generated every `inFrameInterval` seconds over a period of `inFrameWidth` seconds.

Input arguments:

`inSignal` = the input signal (if this is a matrix, peak values are calculated for each channel (ie. row) in the signal)
`inSampleFreq` = the sample frequency of the input signal (in Hz)
`inFrameWidth` = the period over which the maximum is calculated (in s)
`inFrameInterval` = the period between two successive frames (in s)
`inUseAbs` = if non-zero, `abs(inSignal)` is used instead of the original input
 if not specified or empty, 1 is used by default
`inPlotFlag` = if non-zero, plots are generated
 if not specified or empty, 0 is used by default

Output:

`outPeakSignal` = the running maximum value of the signal
`outPeakFreq` = the sample frequency of the maximum signal

Example:

```
[s,fs] = IPEMReadSoundFile;
[Max,MaxFreq] = IPEMCalcPeakLevel(s,fs,0.020,0.010,1,1);
```

Authors:

Koen Tanghe - 20040323

IPEMCalcRMS

Usage:

```
[outRMSSignal,outRMSFreq] = IPEMCalcRMS (inSignal,inSampleFreq,inFrameWidth,  
inFrameInterval,inPlotFlag)
```

Description:

This function calculates the running RMS value of the given signal: an RMS value is generated every `inFrameInterval` seconds over a period of `inFrameWidth` seconds.

Input arguments:

`inSignal` = the input signal (if this is a matrix, RMS values are calculated for each channel (ie. row) in the signal)
`inSampleFreq` = the sample frequency of the input signal (in Hz)
`inFrameWidth` = the period over which the RMS is calculated (in s)
`inFrameInterval` = the period between two successive frames (in s)
`inPlotFlag` = if non-zero, plots are generated
 if not specified or empty, 0 is used by default

Output:

`outRMSSignal` = the running RMS value of the signal
`outRMSFreq` = the sample frequency of the RMS signal

Example:

```
[RMSSignal,RMSFreq] = IPEMCalcRMS(ANI,ANIFreq,0.050,0.010);
```

Authors:

Koen Tanghe - 20010221

IPEMCalcRoughnessOfToneComplex

Usage:

```
[outSignal,outSignalFreq,outRoughness,outRoughnessFreq] = ...
    IPEMCalcRoughnessOfToneComplex(inBaseFreq,inOctaveRatio,
                                    inNumOfHarmonics,inDuration,inPlotFlag)
```

Description:

Calculates the roughness for a superposition of two complex tones having inNumOfHarmonics harmonics:

- a constant one with fundamental frequency inBaseFreq
- one with linearly increasing fundamental frequency (from inBaseFreq to inBaseFreq*inOctaveRatio)

Any octave ratio can be specified.

Input arguments:

inBaseFreq = fundamental frequency for the tone complex (in Hz)
 inOctaveRatio = frequency ratio for the octave
 inNumOfHarmonics = total number of harmonics for each tone complex
 inDuration = duration of the signal (in s)
 inPlotFlag = if non-zero, plots are generated
 if empty or not specified, 1 is used by default

Output:

outSignal = signal that was analyzed
 outSignalFreq = sample frequency for outSignal (in Hz)
 outRoughness = roughness calculated for the signal
 outRoughnessFreq = sample frequency for outRoughness (in Hz)

Example:

```
[s,fs,r,rfreq] = IPEMCalcRoughnessOfToneComplex(440,2,5,5);
```

Authors:

Koen Tanghe - 20050120

IPEMCalcRoughnessOverSubparts

Usage:

```
[outRoughness,outDiffCents] =
    IPEMCalcRoughnessOverSubparts(inOriginalFileName,inOriginalFilePath,
                                    inCentsToRaise,inCentsPerStep,
                                    inMixedFileName,inMixedFilePath)
```

Description:

Calculates the roughness over the subparts of a sound file that was made by CoolEdit2000 using a script file generated by IPEMGeneratePitchShiftScript. An array is returned containing a value for the roughness for each subpart of the sound file.

Input arguments:

inOriginalFileName = name of the original sound file that was processed
 inOriginalFilePath = path to the location of the original sound file
 if empty or not specified, IPEMRootDir('input')\Sounds'
 is used by default
 inCentsToRaise = maximum number of cents the pitch was raised
 if empty or not specified, 1200 is used by default
 inCentsPerStep = number of cents the pitch was raised per step
 if empty or not specified, 100 is used by default
 inMixedFileName = name of the file containing the mix of the original
 and the pitch shifted sounds
 if empty or not specified, the same name as the original
 file is used, but now ending on '_mixed.wav' instead of
 just '.wav'
 inMixedFilePath = path to the mixed file
 if empty or not specified, the same location as the
 original sound file is used by default

Output:

outRoughness = the roughness values
 outDiffCents = the corresponding difference in pitch (in cents)

Example:

```
[Roughness,DiffCents] = IPEMCalcRoughnessOverSubparts('bottle.wav');
```

Authors:

Koen Tanghe - 20011008

IPEMCalcSpectrogram

Usage:

```
[outSpectrogram,outTimes,outFrequencies] = ...
    IPEMCalcSpectrogram(inSignal,inSampleFreq,inFrameSize,inFrameInterval,
                         inWindowType,inFFTSize,inPlotFlag,inPlotThreshold)
```

Description:

Utility function for calculating/plotting the spectrogram of a signal.

Input arguments:

```
inSignal = signal to analyze
inSampleFreq = sample frequency of the incoming signal (in Hz)
inFrameSize = size of one analysis frame (in s)
    if empty or not specified, 0.040 is used by default
inFrameInterval = interval between successive frames (in s)
    if empty or not specified, 0.010 is used by default
inWindowType = if this is a string: type of window to be used
    supported types are: 'bartlett','blackman','boxcar','hamming',
    'hann','hanning','triang'
    if this is a 1-column matrix: any user defined window of size
    [Round(inFrameSize*inSampleFreq) 1]
    if empty or not specified, 'hanning' is used by default
inFFTSize = size of FFT to be used
    if -1 is specified, a value corresponding to inFrameSize is used
    if empty or not specified, the power of two >= inFrameSize
    in samples is used
inPlotFlag = if non-zero, a plot is generated (showing dB levels)
    if empty or not specified, 1 is used by default
inPlotThreshold = values below this level are ignored in the plot (in dB)
    if empty or not specified, -Inf is used by default
```

Output:

```
outSpectrogram = spectrogram data (amplitude values, so not in dB)
outTimes = instants at which an analysis was calculated (in s)
outFrequencies = frequencies used in decomposition (in Hz)
```

Example:

```
[S,T,F] = IPEMCalcSpectrogram(s,fs,0.040,0.010);
```

Authors:

Koen Tanghe - 20001130

IPEMCalcZeroCrossingRate

Usage:

```
[outZCR,outZCRFreq] =
    IPEMCalcZeroCrossingRate(inSignal,inSignalFreq,
                            inFrameWidth,inFrameInterval,
                            inZeroTolerance,inPlotFlag)
```

Description:

Calculates the number of zero-crossings per second for successive frames.

Input arguments:

inSignal = signal to be analyzed
 inSignalFreq = sample frequency of inSignal
 inFrameWidth = width of 1 frame (in s)
 inFrameInterval = interval between successive frames (in s)
 inZeroTolerance = defines a region around zero where zero-crossings are not
 counted (useful for noisy parts of the signal)
 if empty or not specified, 0 is used by default
 inPlotFlag = if non-zero, plots the ZCR
 if not specified or empty, 1 is used by default

Output:

outZCR = number of zero-crossings per second for each frame
 outZCRFreq = sample frequency of outZCR

Example:

```
[ZCR,ZCRFreq] = IPEMCalcZeroCrossingRate(s,fs,0.05,0.01,0.001);
```

Authors:

Koen Tanghe - 20001012

IPEMCheckVersion

Usage:

```
function [outCompatibility,outCurrentVersion] = ...
    IPEMCheckVersion(inComponent,inReferenceVersion)
```

Description:

Checks compatibility using version numbers

Input arguments:

inComponent = string identifying the component to check
inReferenceVersion = version number string to compare with

Output:

outCompatibility = compatibility result:
 -1 if the current version is lower than the reference
 0 if the current version is the same as the reference
 +1 if the current version is higher than the reference
 empty if the component was not present at all
outCurrentVersion = current version number string of requested component
 empty if not present

Remarks:

For comparison, str2num is used on the version number strings.

Example:

```
[Comp,CurrVer] = IPEMCheckVersion('signal','5.0');
```

Authors:

Koen Tanghe - 20011204

IPEMClip

Usage:

```
outSignal = IPEMClip(inSignal,inLowLimit,inHighLimit,
                      inClipLowTo,inClipHighTo)
```

Description:

This function clips the incoming (multi-channel) signal at the given limits. Specify empty values for either one of the limits if you don't want clipping at that side of the signal range.

Signal channels are represented by rows.

Input arguments:

inSignal = the (multi-channel) signal to be clipped
 inLowLimit = specifies the low level clipping value: values lower than this are replaced by either the limit itself or inClipLowTo
 if empty, no clipping occurs
 inHighLimit = specifies the high level clipping value: values higher than this are replaced by either the limit itself or inClipHighTo
 if empty or not specified, no clipping occurs
 inClipLowTo = if non-empty, this is a replacement value for too low values
 if empty or not specified, inLowLimit is used
 inClipHighTo = if non-empty, this is a replacement value for too high values
 if empty or not specified, inHighLimit is used

Output:

outSignal = the clipped signal

Example:

```
Signal = IPEMClip(Signal,0.05,1,0,[]);
```

Authors:

Koen Tanghe - 20000419

IPEMCombFilter

Usage:

```
outSignal = IPEMCombFilter(inSignal,inSampleFreq,inDelay,
                           inAttenuation,inSync)
```

Description:

Filters the incoming signal with a comb filter according to the following scheme:

$$\text{OUT} = \text{IN} * (1 - \text{inAttenuation} * z^{-m})$$

where $m = \text{round}(\text{inDelay} * \text{inSampleFreq})$

In the frequency domain, this produces a comb-like frequency response, in which either the peaks (if $\text{inSync} == 1$) or the valleys (if $\text{inSync} == 0$) of the comb are at frequencies $n.F_0$ ($F_0 = 1/\text{inDelay}$).

Input arguments:

- inSignal = input signal (each row represents a channel)
- inSampleFreq = input signal sample frequency (in Hz)
- inDelay = delay of feedforward branch (in s)
- inAttenuation = attenuation ratio (0 to 1)
- inSync = if 1, the comb's peaks are at frequencies $n/delay$
if 0, the comb's valleys are at frequencies $n/delay$

Output:

outSignal = the filtered signal

Example:

```
scomb = IPEMCombFilter(s,fs,1/500,0.98,1);
```

Authors:

Koen Tanghe - 20000208

IPEMContextualityIndex

Usage:

```
[outChords,outToneCenters,...  
outContextuality1,outContextuality2,outContextuality3] = ...  
IPEMContextualityIndex(inPeriodicityPitch,inSampleFreq,inPeriods,...  
inSnapShot,inHalfDecayChords,...  
inHalfDecayToneCenters,inEnlargement,inPlotFlag)
```

Description:

This function calculates the contextuality index. Two methods are used:

- inspection: compares fixed chord images with running chord images
and running tone center images
- comparison: compares running chord images with running tone center images

Input arguments:

```
inPeriodicityPitch = periodicity pitch image  
inSampleFreq = sample frequency of the input signal (in Hz)  
inPeriods = periods of periodicity analysis (in s)  
inSnapShot = time where the snapshot should be taken (in s)  
              if negative, the time is taken at abs(inSnapShot) from the end  
              of the sample  
              if empty or not specified, the time of the last sample is used  
              by default  
inHalfDecayChords = half decay time for leaky integration into chord image  
              if empty or not specified, 0.1 is used by default  
inHalfDecayToneCenters = half decay time for leaky integration into tone  
              center image  
              if empty or not specified, 1.5 is used by default  
inEnlargement = time by which the input signal is enlarged (in s)  
              if -1, 2*inHalfDecayToneCenters is used  
              if empty or not specified, 0 is used by default  
inPlotFlag = if non-zero, plots are generated  
              if empty or not specified, 0 is used by default
```

Output:

```
outChords = local integration of inPeriodicityPitch into chord image  
outToneCenters = global integration of inPeriodicityPitch into tone center  
              image  
outContextuality1 = correspondence between chord taken at snapshot position  
              and running chord image
```

```
outContextuality2 = correspondence between chord taken at snapshot position  
and running tone center image  
outContextuality3 = correspondence between running chord image and running  
tone center image
```

Remarks:

Sample frequency of output signals is the same as inSampleFreq.

Example:

```
[Chords,ToneCenters,Contextuality1,Contextuality2,Contextuality3] = ...  
IPEMContextualityIndex(PP,PPFreq,PPPeriods);
```

Authors:

Marc Leman - 19991221 (originally made)
Koen Tanghe - 20010129 (minor code changes + documentation)

IPEMConvertToAMNoise

Usage:

```
outAMNoise = ...
IPEMConvertToAMNoise(inSignal,inSampleFreq,inWantedSampleFreq, ...
                     inPlaySound,inNoiseBand)
```

Description:

Converts a signal to amplitude modulated noise (the given signal is used as modulator) and plays back the sound (optionally).
Can be used for a quick listening to envelope-like signals.

Input arguments:

```
inSignal = signal to convert to AM noise
inSampleFreq = sample frequency of incoming signal (in Hz)
inWantedSampleFreq = wanted sample frequency for the resulting AM noise
                     (in Hz)
                     if empty or not specified, 22050 is used by default
inPlaySound = if empty, non-zero or not specified, the AM noise will be
              played back
              otherwise, no sound is produced
inNoiseBand = two element row vector specifying the wanted the noise band
              (in Hz)
              if empty or not specified, broad band noise is used by default
```

Output:

```
outAMNoise = the resulting amplitude modulated noise
```

Example:

```
AMNoise = IPEMConvertToAMNoise(RMS,RMSFreq,22050,1);
```

Authors:

Koen Tanghe - 20010221

IPEMConvertToClickSound

Usage:

```
outClickSignal = IPEMConvertToClickSound(inClickTimes,inWantedSampleFreq,...  
inTimeRange,inAmplitudes)
```

Description:

Converts a sequence of times into a click sound signal.

Input arguments:

inClickTimes = Times where a click should be generated (in s).
inWantedSampleFreq = Sample frequency of the output signal (in Hz).
inTimeRange = Time range for which clicks should be generated (only times from inClickTimes that fall into this range will be used). Specified as [min max] (in s). If empty or not specified, the full range will be used.
inAmplitudes = Amplitudes corresponding to the click times. Should be in [0...1] range. If this is a single scalar, this value will be used for all click times. If empty or not specified, 1 is used by default.

Output:

outClickSignal = The click sound signal.

Remarks:

The clicks are very simple 0 to 1 transitions (Dirac pulses). The end time of the time range is exclusive.

Example:

```
s = IPEMConvertToClickSound(0:0.5:5,44100,[1 4],0.8);
```

Authors:

Koen Tanghe - 20040423

IPEMConvertToMIDINoteNr

Usage:

```
outMIDINoteNr = IPEMConvertToMIDINoteNr(inNoteString)
```

Description:

Converts a cell array of note specification strings to an array of MIDI note numbers.

Input arguments:

```
inNoteString = cell array of note specification string
```

Output:

```
outMIDINoteNr = MIDI note number array
```

Remarks:

Both sharps ('#') or flats ('b') can be used.

You can specify a sign in front of the octave part, if you want to...

Things like 'Cb' and 'E#' will be interpreted as, respectively, 'B' and 'F'.

If you want to compare note strings, you'll have to make sure that you always stick to the same notation method (sharps OR flats).

If you don't, you'll have to use the note NUMBERS to test for equality, and not the note STRINGS !

Example:

```
IPEMConvertToMIDINoteNr('A4')
```

Authors:

Koen Tanghe - 20000315

IPEMConvertToString

Usage:

```
outNoteString = IPEMConvertToString (inMIDINoteNr,inUseFlats)
```

Description:

Converts an array of MIDI note numbers to a cell array of note strings.

Input arguments:

inMIDINoteNr = array containing MIDI note numbers
inUseFlats = if non-zero, flats will be used instead of sharps
 if empty or not specified, 0 is used by default (sharps)

Output:

```
outNoteString = cell array with note specification strings
```

Remarks:

Where inMIDINoteNr contains an illegal MIDI note number (NaN), the note string corresponding with that value will be left empty...

Example:

```
IPEMConvertToString([69 70 ; 58 57])
```

Authors:

Koen Tanghe - 20000315

IPEMCountZeroCrossings

Usage:

```
outZeroCrossings = IPEMCountZeroCrossings(inSignal,inZeroTolerance)
```

Description:

Counts the number of zero crossings for each row of inSignal.

Input arguments:

inSignal = signal to be analyzed for zero-crossings
inZeroTolerance = defines a region around zero where zero-crossings are not
counted (useful for noisy parts of the signal)
if empty or not specified, 0 is used by default

Output:

```
outZeroCrossings = number of zero crossings for each row
```

Example:

```
ZeroCrossing = IPEMCountZeroCrossings(MusicExcerpt,0.01);
```

Authors:

Koen Tanghe - 20000509

IPEMCreateMask

Usage:

```
outMask = IPEMCreateMask(inTime,inStartTimes,inAmplitudes,inDecayPeriod,...  
inFractionAtDecayPeriod)
```

Description:

This function creates an exponentially decaying mask from the given start times and amplitudes.

Input arguments:

inTime = the time span for which a mask is needed (in samples !)
 inStartTimes = the start times of the masking events (in samples !)
 inAmplitudes = the amplitudes of the masking events
 inDecayPeriod = the time (in samples !) it takes before a single mask reaches
 inFractionAtDecayPeriod of its initial amplitude
 inFractionAtDecayPeriod = fraction of original amplitude at inDecayPeriod
 if empty or not specified, 0.5 is used by default
 (i.e. half decay time)

Output:

outMask = the requested mask

Remarks:

The relation between a decay time (DT) + fraction at decay time (DF) specification and a half decay time (HDT) specification is as follows:

$$\text{HDT} = \text{DT} * \log(0.5) / \log(\text{DF})$$

Example:

```
Mask = IPEMCreateMask(1:1000,[20 250 400 860 900],[1 0.5 1 0.5 1],100);  
or  
fs = 22050; t = 0:1/fs:5-1/fs;  
Mask = IPEMCreateMask(t,[1 2 3],[1 0.5 1],0.2);
```

Authors:

Koen Tanghe - 20040106

IPEMDoOnsets

Usage:

```
[Ts,Tsmp] = IPEMDoOnsets(inFileName,inFilePath,inPlotFlag);
```

Description:

Convenience function for detecting onsets in a sound signal.
The onset times are also written to a text file.

Input arguments:

inFileName = name of the sound file to process (with extension!)
 inFilePath = path to the directory where the sound file is located
 if empty or not specified, IPEMRootDir('input')\Sounds is used
 by default
 inPlotFlag = if non-zero, plots are generated (no plots if zero)
 if empty or not specified, 0 is used by default

Output:

Ts = onset times in seconds
 Tsmp = onset times in samples

Additionally, a text file containing info about the detected onsets is stored in the same directory as the sound file. It has the same name as the sound file without extension, but followed by '_onsets.txt'.

Example:

```
[Ts,Tsmp] = IPEMDoOnsets('music.wav');
```

Authors:

Koen Tanghe - 20011107

IPEMEnsureDirectory

Usage:

```
outResult = IPEMEnsureDirectory (inDirectoryPath,inCreate)
```

Description:

Checks whether the given directory exists (and optionally creates it, if not)

Input arguments:

```
inDirectoryPath = full directory path to be checked (created)  
inCreate = if non-zero, the directory will be created if it doesn't exist  
          if empty or not specified, 1 is used by default
```

Output:

```
outResult = 1 if the directory already existed before calling this function
```

Remarks:

Use this function with care: it's good policy to only create a directory

1. if this is really necessary (let user specify directory if possible)
2. within the default output directory (that way, the user's disk doesn't get scattered with new directories)

Example:

```
IPEMEnsureDirectory(fullfile(IPEMRootDir('output'),'Test','Example1'),1);
```

The above line makes sure that a directory Test\Example1 exists within the default output directory.

Authors:

Koen Tanghe - 20040504

IPEMEnvelopeFollower

Usage:

```
[outEnvelope,outEnvelopeFreq] = ...
IPEMEnvelopeFollower(inSignal,inSignalFreq,inAttackTime,inReleaseTime, ...
inPlotFlag);
```

Description:

Does envelope following on a signal.

Input arguments:

inSignal = one-dimensional input signal
 inSignalFreq = sample frequency for inSignal (in Hz)
 inAttackTime = time it takes to reach 0.5 for a 0 to 1 step signal (in s)
 inReleaseTime = time it takes to reach 0.5 for a 1 to 0 step signal (in s)
 if not specified or empty, inAttackTime is used by default
 inPlotFlag = if not specified or empty, 0 is used by default

Output:

outEnvelope = envelope signal
 outEnvelopeFreq = sample frequency for outEnvelope (same as inSignalFreq)

Remarks:

Uses one-pole low-pass filters with different coefficients for attack and release on the absolute value of the signal. The attack and release times are specified as "half-value times".

Example:

```
fs = 44100;
s = zeros(1,fs*5);
s(fs:fs*3) = 1;
IPEMEnvelopeFollower(s,fs,0.010,0.020,1);
```

Authors:

Koen Tanghe - 20030403

IPEMExportFigures

Usage:

```
IPEMExportFigures(inInputPath,inOutputPath,inFormat,inUseIPEMPLayout)
```

Description:

Exports figures to specified graphics format.

Input arguments:

```
inInputPath = path to scan for .fig files
              if empty or not specified, the currently opened figures are
              exported, with the names being just 'Figure_No_1', etc...
              (in this case, inOutputPath must be given)
inOutputPath = path to be used for saving the exported figures
              if empty or not specified, inInputPath is used by default
inFormat = graphics format to be used
              this can be one of the following:
              'lowpng' = low resolution png ('-dpng -r75')
              'highpng' = high resolution png ('-dpng -r300')
              'screenpng' = screen resolution png ('-dpng -r0')
              'eps' = black & white eps ('-deps -r600')
              'epsc' = colored eps ('-depsc -r600')
              or an entire specification conforming to Matlab's 'print' options
              (to be specified as a cell array of strings, see 'help print')
              if empty or not specified, 'lowpng' is used by default
inUseIPEMPLayout = if non-zero, the default IPEM layout settings are used
                     if empty or not specified, each figure is used "as is"
```

Output:

Files having the same name as the figures, but with an extension depending on inFormat.

Remarks:

This thing is NOT COMPILABLE because in Matlab R11.1 the print command does not work when called from a compiled M-file (does work in R12 however).

Example:

```
IPEMExportFigures('E:\Koen\Docs\Research\TeacupPaper\Originals');
```

Authors:

Koen Tanghe - 20010528

IPEMExtractSegments

Usage:

```
outSignalSegments = IPEMExtractSegments(inSignal,inSampleFreq,inTimeSegments)
```

Description:

Extracts segments from the signal according to the given time segments.

Input arguments:

```
inSignal = input signal (1 row vector)
inSampleFreq = sample frequency of inSignal (in Hz)
inTimeSegments = time segments for cutting up the input signal (in s)
                  (each row contains the start and end of a time segment)
                  if this is a positive scalar, the signal is divided into
                  inTimeSegments parts of equal size + an additional segment
                  containing the remaining part of the signal (which will
                  contain less than inTimeSegments samples)
                  if this is a negative scalar, the signal is divided into
                  -inTimeSegments parts where some segments will contain
                  1 sample more than other segments
```

Output:

```
outSignalSegments = 1 column cell (!) vector of signal segments
                    (each row contains a segment)
```

Example:

```
Parts = IPEMExtractSegments(s,fs,OnsetSegments);
FirstSegment = Parts{1}; % etc...
```

Authors:

Koen Tanghe - 20000510

IPEMFadeLinear

Usage:

```
outSignal = IPEMFadeLinear (inSignal,inSampleFreq,inFadeTime)
```

Description:

Does a linear fade in and out of the signal over inFadeTime seconds.
Can be used for smoothing beginnings and endings of a synthesized signal.

Input arguments:

```
inSignal = the input signal  
inSampleFreq = sample frequency of the signal (in Hz)  
inFadeTime = the time to fade in and out (in s)  
    if this is a scalar, both fade in and out time are the same  
    if this is a two element row vector, inFadeTime(1) is the fade  
    in time and inFadeTime(2) is the fade out time  
(fade times should be smaller than the signal duration)
```

Output:

```
outSignal = the faded signal
```

Remarks:

If the sum of fade in and fade out time is \geq the sound length,
then the fades will accumulate at the overlapping section, thus introducing
a quadratic fade at that section.

Example:

```
Signal = IPEMFadeLinear(Signal,SampleFreq,0.01);
```

Authors:

Koen Tanghe - 20000913

IPEMFindAllPeaks

Usage:

```
thePeakIndices = IPEMFindAllPeaks (inSignal,inFlatPreference,inPlotFlag)
```

Description:

This function finds the indices of all peaks in the given signal vector. A peak is taken to occur whenever a rising part in the signal is followed by a falling part.

Input arguments:

```
inSignal = the signal that is scanned for peaks (a row vector)
inFlatPreference = preference for choosing exact position in case of
                  flat peaks:
                  if 'left', the leftmost point of a flat peak is chosen
                  if 'center', the center of a flat peak is chosen
                  if 'right', the rightmost point of a flat peak is chosen
                  if empty or not specified, 'center' is used by default
inPlotFlag = if non-zero, plots are generated
              if empty or not specified, 0 is used by default
```

Output:

```
outPeakIndices = a vector with the indices of the peaks (could be empty!)
```

Example:

```
PeakIndices = IPEMFindAllPeaks(Signal,'center');
```

Authors:

Koen Tanghe - 20000509

IPEMFindNearestMinima

Usage:

```
[outLeftIndex, outRightIndex] = IPEMFindNearestMinima(inSignal,inPeakIndex)
```

Description:

Finds the nearest minima to the left and the right of the specified peak.
Peaks can be found with IPEMFindAllPeaks.

Input arguments:

inSignal = the signal to analyze
inPeakIndex = index in inSignal of a peak

Output:

outLeftIndex = the index of the nearest minimum to the left of the peak
if no real minimum was found, 1 is returned
outRightIndex = the index of the nearest minimum to the right of the peak
if no real minimum was found, length(inSignal) is returned

Example:

```
[LeftIndex,RightIndex] = IPEMFindNearestMinima(Signal,PeakIndex);
```

Authors:

Koen Tanghe - 20000208

IPEMFindNoteFromFrequency

Usage:

```
[outNoteNr,outOctaveNr,outExactFrequency,outNoteString] =
IPEMFindNoteFromFrequency(inFrequency,inRefNoteNr,inRefOctaveNr,
                            inRefFreq,inNotesPerOctave,inOctaveRatio)
```

Description:

Finds nearest note corresponding to the given frequency.

Supports non-standard tone scales that can be calculated like this:

$$\text{Frequency} = \text{RefFreq} * \text{OctaveRatio}^{\text{Exponent}}$$

where:

$$\text{Exponent} = (\text{OctaveNr}-\text{RefOctaveNr}) + (\text{NoteNr}-\text{RefNoteNr})/\text{NotesPerOctave}$$

Input arguments:

inFrequency = frequency for which a corresponding note must be found
 inRefNoteNr = the rank number of the reference note
 if empty or not specified, 10 is used by default
 inRefOctaveNr = the octave number of the reference note
 if empty or not specified, 4 is used by default
 inRefFreq = the frequency of the reference note (in Hz)
 if empty or not specified, 440 Hz is used
 inNotesPerOctave = the number of notes in one octave
 if empty or not specified, 12 is used by default
 inOctaveRatio = the frequency ratio between two octaves
 if empty or not specified, 2 is used

Output:

outNoteNr = rank number of the note (1 to inNotesPerOctave)
 outOctaveNr = octave number of the note (integer)
 outExactFrequency = exact frequency of the found note (in Hz)
 outNoteString = string representing the found note
 (only for standard tone scale, otherwise empty)

Remarks:

Getting the note name for standard tone scale:

```
MIDINoteNr = (outNoteNr-1) + 12*(outOctaveNr+1);  
NoteString = IPEMConvertToNoteString(MIDINoteNr);
```

Example:

```
[NoteNr,OctaveNr,Freq,Name] = ...  
IPEMFindNoteFromFrequency([441.01 153 ; 222 94]);
```

Authors:

Koen Tanghe - 20000629

IPEMGenerateBandPassedNoise

Usage:

```
outSignal = ...
    IPEMGenerateBandPassedNoise(inPassBands,inDuration,inSampleFreq, ...
                                indBLevel,inFFTWidth)
```

Description:

Generates a band passed noise signal using inverse FFT.

Input arguments:

```
inPassBands = 2 column vector containing specification of the pass bands:
              each row specifies a low and high frequency (in Hz)
inDuration = duration of the sound (in s)
inSampleFreq = wanted sample frequency (in Hz)
              if empty or not specified, 22050 is used by default
indBLevel = dB level for the signal (in dB)
              if empty or not specified, -6 is used by default
inFFTWidth = FFT width to use for the inverse FFT
              if -1, the next power of 2 >= number of samples is used
              if empty or not specified, the number of samples is used
              by default
```

Output:

outSignal = generated band passed noise signal

Remarks:

When using an inFFTWidth that is smaller than the number of requested samples, the resulting sound will consist of the appropriate number of repetitions of the inverse fft of size inFFTWidth, thus introducing an artificial period of the size of 1 inverse fft. For non dense noise, this can become noticeable.

Example:

```
s = IPEMGenerateBandPassedNoise([1000 1200],1);
```

Authors:

Koen Tanghe - 20000926

IPEMGenerateFrameBasedSegments

Usage:

```
[outSegments,outActualWidth,outActualInterval] = ...
    IPEMGenerateFrameBasedSegments(inSignalOrDuration,inSampleFreq, ...
        inFrameWidth,inFrameInterval,... ...
        inIncludeIncompleteFrames)
```

Description:

Generates (overlapping) time segments using equally sized frames.

Input arguments:

`inSignalOrDuration` = if this is a vector, it is the signal for which to calculate the frame-based time segments
 if this is a scalar, it is the duration of a signal (in s) (allows to generate segments even if you don't have a signal)
`inSampleFreq` = sample frequency of `inSignal` (in Hz)
`inFrameWidth` = wanted width of 1 frame (in s)
`inFrameInterval` = wanted interval between frames (in s)
`inIncludeIncompleteFrames` = if non-zero, incomplete frames at the end are included as well (different lengths!)
 if empty or not specified, 0 is used by default

Output:

`outSegments` = 2-column array of segments:
 each row contains the start and end moment of a segment (in s)
`outActualWidth` = width used (to get an integer number of samples)
`outActualInterval` = interval used (to get an integer number of samples)

Remarks:

Because of the discrete nature of sampled signals, the frame width and interval requested for by the caller might not be representable in the resolution specified by the sample frequency.
 In that case, the nearest width and interval that are integer multiples of the sample period will be used. This is what the `outActualWidth` and `outActualInterval` arguments are for.

Example:

```
[Segments,Width,Interval] = IPEMGenerateFrameBasedSegments(1,100,0.1,0.025)
```

Width will be 0.1 s (as requested)

Interval will become 0.03 s, since 0.025 s does not exist in a 1/100 s resolution

Authors:

Koen Tanghe - 20000418

IPEMGeneratePitchShiftScript

Usage:

```
IPEMGeneratePitchShiftScript(inFileName,inFilePath,inCentsToRaise,...  
inCentsPerStep,inScriptFilePath,inOnlyShift)
```

Description:

Generates a CoolEdit2000 script for creating a sound file containing appended pitch-shifted versions of the original sound file.
Generated pitch-shifted sounds can be mixed with original sound.

Way to go:

1. run this function on your sound file
2. start CoolEdit2000 and open the sound file
3. go to 'Options\Scripts & Batch Processing' and open the script file
4. deselect 'Pause at Dialogs', select the script 'shift' and press 'Run'
5. wait...
6. save your sound file (either using suffix '_shifted' or '_mixed')

Input arguments:

inFileName = name of the sound file to generate a script for
 inFilePath = path to the location of the sound file
 if empty or not specified, IPEMRootDir('input')\Sounds is used
 by default
 inCentsToRaise = maximum number of cents for the pitch to raise
 if empty or not specified, 1200 is used by default
 inCentsPerStep = number of cents the pitch will raise per step
 if empty or not specified, 100 is used by default
 inScriptFilePath = path to the location where the script file will be saved
 if empty or not specified, the same location as the sound
 file is used by default
 inOnlyShift = if 1, the script will only generate the shifted versions
 of the sound signal (from low to high)
 otherwise, it will immediately mix these shifted versions with
 the appropriate number of copies of the original sound signal
 if empty or not specified, 0 is used by default
 (mixing enabled)

Output:

A CoolEdit2000 script with the same name as the sound file, but with the extension .scp

Remarks:

CoolEdit2000 scripts contain selection ranges and sample rates (which are different for different sound files), so you can only use the script for the sound file you created it for (or for sound files with exactly the same number of samples and sampling rate)...

You can download a trial version of CoolEdit2000 from:

<http://www.syntrillium.com>

Make sure you select function groups 1 (Save...) and 3 (Stretching...) in order to be able to run the script.

Example:

```
IPEMGeneratePitchShiftScript('bottle.wav');
```

Authors:

Koen Tanghe - 20000418

IPEM GetAllCombinations

Usage:

```
outCombinations = IPEM GetAllCombinations(inValueArrays)
```

Description:

Finds all combinations of values by picking a value from a specified set of values for each of the multiple dimensions.

Input arguments:

inValueArrays = Cell array containing for each dimension an array of the values from which one should be chosen for a combination.

Output:

outCombinations = Array containing all possible combinations in rows.

Example:

```
Comb = IPEM GetAllCombinations({[1 2 3],[10 20],[111 222 333]})
```

Authors:

Koen Tanghe - 20040510

IPEMGetCrestFactor

Usage:

```
outCrestFactor = IPEMGetCrestFactor(inSignal)
```

Description:

Gets the crest factor of the signal, which is defined as:
$$CF = \max(\text{abs}(\text{signal})) / \text{rms}(\text{signal})$$

Input arguments:

inSignal = input signal

Output:

outCrestFactor = the crest factor over the entire signal

Example:

```
[s,fs] = IPEMReadSoundFile;  
CF = IPEMGetCrestFactor(s);
```

Authors:

Koen Tanghe - 20030415

IPEMGetFileList

Usage:

```
[outFullFiles,outFileNames,outFilePaths] = ...
IPEMGetFileList(inDirectory,inFilePattern,inRecurse)
```

Description:

Gets a list of files located in a given directory and whose names adhere to a specific pattern (option for recursion).

Input arguments:

inDirectory = Path to the directory to search in.
inFilePattern = Pattern to which files should adhere in order to be listed.
If empty or not specified, '*' is used by default.
inRecurse = If non-zero, also subdirectories are scanned.
If empty or not specified, 0 is used by default.

Output:

outFullFiles = Cell array containing the full file specifications.
outFileNames = Cell array containing only the file names.
outFilePaths = Cell array containing only the paths to the files.

Example:

```
[F, FN, FP] = IPEMGetFileList('D:\Koen\', '*.txt', 1);
```

Authors:

Koen Tanghe - 20040507

IPEMGetKurtosis

Usage:

```
outKurtosis = IPEMGetKurtosis(inData)
```

Description:

Gets the kurtosis of the data, which is defined here as:

$$\text{Kurtosis}(X) = \frac{\sum((X-\mu)^4)}{(N\sigma^4)} - 3$$

where

μ = average of X

σ = standard deviation (normalized using N) of X

N = number of data points in X

Input arguments:

```
inData = data vector
```

Output:

```
outKurtosis = the kurtosis of the data
```

Example:

```
X = [1 2 3 1 5 6 4 3];  
Kurtosis = IPEMGetKurtosis(X);
```

Authors:

Koen Tanghe - 20030415

IPEMGetLevel

Usage:

```
outRMSLevel = IPEMGetLevel(inSignal,inUseDecibels)
```

Description:

Calculates the average RMS power level of a signal in dB (or not).

Input arguments:

```
inSignal = signal to be analyzed  
inUseDecibels = if non-zero, dB units are used instead of plain RMS  
                if empty or not specified, 1 is used by default
```

Output:

```
outRMSLevel = average RMS level (in dB if inUseDecibels is non-zero)
```

Remarks:

The reference value of 0 dB is the level of a square wave with amplitude 1. Thus, a sine wave with amplitude 1 yields -3.01 dB.

Example:

```
RMSLevel = IPEMGetLevel(Signal);
```

Authors:

Koen Tanghe - 20040323

IPEMGetRolloff

Usage:

```
outRolloffIndex = IPEMGetRolloff(inData,inRolloffFraction)
```

Description:

Gets the index R in inData for which:

$$\text{sum}(\text{inData},1,R) = \text{inRolloffFraction} * \text{sum}(\text{inData},1,\text{length}(\text{inData}))$$

Input arguments:

inData = Data to be analyzed (each row is analyzed separately).

inRolloffFraction = Column vector containing a rolloff fraction for each row
in inData. If a scalar is given, the same value is used
for all rows.

Output:

outRolloffIndex = column vector containing the rolloff indices for each row
in inData

Example:

```
v = rand(1,20);  
ri = IPEMGetRolloff(v,0.85);  
v(ri)
```

Authors:

Koen Tanghe - 20030523

IPEMGetRoughnessFFTReference

Usage:

```
outReferenceValue = IPEMGetRoughnessFFTReference(indBLevel,inLength)
```

Description:

Returns the reference value for IPEMRoughnessFFT at the specified dB level.

Input arguments:

indBLevel = dB level for which the reference value is requested
inLength = length (in s) used for calculating roughness with FFT method
if empty or not specified, 1 is used by default

Output:

outReferenceValue = the requested reference value

Remarks:

The reference is taken to be an AM sine tone at the specified dB level with the following properties:

- carrier frequency = 1000 Hz
- modulation frequency = 70 Hz
- modulation index = 1
- duration = 1 s

The reference value is calculated over the reference sound as a whole (not as the mean of the calculated frames over the sound)

Example:

```
Ref = IPEMGetRoughnessFFTReference(-20,0.5)
```

Authors:

Koen Tanghe - 20010816

IPEMGetSkew

Usage:

```
outSkew = IPEMGetSkew(inData)
```

Description:

Gets the skew of the data, which is defined here as:

$$\text{Skew}(X) = \frac{\sum((X-\mu)^3)}{(N\sigma^3)}$$

where

μ = average of X

σ = standard deviation (normalized using N) of X

N = number of data points in X

Input arguments:

```
inData = data vector
```

Output:

```
outSkew = the skew of the data
```

Example:

```
X = [1 2 3 1 5 6 4 3];  
Skew = IPEMGetSkew(X);
```

Authors:

Koen Tanghe - 20030415

IPEMHandleInputArguments

Usage:

```
[...] = IPEMHandleInputArguments(inPassedArgs,inNumFirstOptionalArg,
                                inDefaults)
```

Description:

This function is a simple but general input handling routine. Just pass what you got from your caller, say what variables you expect and which of them are optional, and specify the defaults (which will be used to replace arguments that were left empty). After calling this function, you'll have the valid arguments you need to proceed with.

Input arguments:

inPassedArguments = cell array with the arguments as passed to the caller
 inNumFirstOptionalArg = number of the first optional argument
 if empty, length(inDefaults)+1 is used by default
 (this means: no optional arguments)
 inDefaults = cell array with default values for the arguments that were
 left empty

Output:

The specified variables are set in the caller's name space or an appropriate error message is returned if something was wrong with the arguments.

Example:

This is a function using IPEMHandleInputArguments:

```
function Test (varargin)
% Usage:
%   Test(inArg1,inArg2,inArg3,inArg4)
% -----
%
% Handle input arguments
[inArg1,inArg2,inArg3,inArg4] = ...
    IPEMHandleInputArguments (varargin, 3, {'def1' 'def2' 0.1 0.2});
%
% Rest of code for Test
```

.....

If you now call Test the following way:

```
Test('abc',[],5);
```

the internal variables of Test after the call to IPEMHandleInputArguments will be:

```
inArg1 = 'abc'  
inArg2 = 'def2'  
inArg3 = 5  
inArg4 = 0.2
```

Authors:

Koen Tanghe - 19991109

IPEMHarmonicTone

Usage:

```
outSignal = IPEMHarmonicTone (inFundamentalFreq,inDuration,inSampleFreq, ...
                             inPhaseFlag,indBLevel,inNumOfHarmonics)
```

Description:

This function generates a tone with a number of harmonics.
The amplitudes of the harmonics vary as $1/N$ (N = number of harmonic).

Input arguments:

inFundamentalFreq = the main frequency in Herz
inDuration = the duration (in s)
if empty or not specified, 1 is used by default
inSampleFreq = the desired sample frequency for the output signal (in Hz)
if empty or not specified, 22050 is used by default
inPhaseFlag = for choosing whether random phase has to be used or not
(1 to use random phase, 0 otherwise)
if empty or not specified, 1 is used by default
indBLevel = dB level of generated tone (in dB)
if empty or not specified, no level adjustment is performed
inNumOfHarmonics = number of harmonics (including fundamental frequency)
if empty or not specified, 10 is used by default

Output:

outSignal = the signal for the tone

Example:

```
Signal = IPEMHarmonicTone(440,1,22050,1,-20,5);
```

Authors:

Koen Tanghe - 20010116

IPEMHarmonicToneComplex

Usage:

```
outSignal = IPEMHarmonicToneComplex(inToneVector,inDuration,inSampleFreq, ...
                                     inPhaseFlag,indBLevel,inNumOfHarmonics)
```

Description:

This function generates a tone complex built up of harmonic tones.

Input arguments:

inToneVector = a 12 elements vector representing the amplitude for each tone
 C, C#, D, ... in the tone complex (0 = no tone, 1 = full
 amplitude)
 inDuration = the duration (in s)
 if empty or not specified, 1 is used by default
 inSampleFreq = the desired sample frequency for the output signal (in Hz)
 if empty or not specified, 22050 is used by default
 inPhaseFlag = for choosing whether random phase has to be used or not
 (1 to use random phase, 0 otherwise)
 if empty or not specified, 1 is used by default
 indBLevel = dB level of generated tone complex (in dB)
 if empty or not specified, no level adjustment is performed
 inNumOfHarmonics = number of harmonics for each tone (including the
 fundamental frequency)
 if empty or not specified, 10 is used by default

Output:

outSignal = the signal for the tone complex

Remarks:

This is a rather simple routine because the frequencies are fixed:

```

notes = C4,      C#4,      D4,      D#4,      E4,      F4,      F#4,      G4,      G#4,
freqs = 261.6,  277.2,  293.7,  311.1,  329.6,  349.2,  370.0,  392.0,  415.3,  

       ...  

notes = A,      A#4,      B
freqs = 440.0,  466.2,  493.9
  
```

Example:

```
Signal = IPEMHarmonicToneComplex([1 0 0 0 1 0 0 1 0 0 0 0],1,22050,1,-20,5);
```

Authors:

Marc Leman - 19990528

Koen Tanghe - 20010116

IPEMLeakyIntegration

Usage:

```
outLeakyIntegration = ...
    IPEMLeakyIntegration(inSignal,inSampleFreq,inHalfDecayTime, ...
                           inEnlargement,inPlotFlag)
```

Description:

Calculates leaky integration with specified half decay time.

Input arguments:

```
inSignal = (multi-dimensional) input signal, each row representing a channel
          leaky integration is performed for each channel
inSampleFreq = sample frequency of inSignal (in Hz)
inHalfDecayTime = time (in s) at which an impulse would be reduced to half
                  its value
                  if empty or not specified, 0.1 is used by default
inEnlargement = time (in s) to enlarge the input signal
                  if -1, 2*inHalfDecayTime is used
                  if empty or not specified, 0 is used by default
inPlotFlag = if non-zero, plots are generated
                  if empty or not specified, 0 is used by default
```

Output:

outLeakyIntegration = the leaky integration of the input signal

Remarks:

Sample frequency of outLeakyIntegration is the same as inSampleFreq.

Example:

```
LeakyIntegration = ...
    IPEMLeakyIntegration(PeriodicityPitchMatrix,PeriodicityPitchFreq,0.1,0.1);
```

Authors:

Marc Leman - 19991224
Koen Tanghe - 20010129

IPEMLoadANI

Usage:

```
[outANI,outANIFreq,outANIFilterFreqs] = IPEMLoadANI(inName,inPath,...  
inANIName,inANIFreqName,inANIFilterFreqsName)
```

Description:

Loads an auditory nerve image and its corresponding sample frequency and filter frequencies from a .mat file on disk.

Input arguments:

```
inName = the name of the .mat file containing the nerve image  
        if empty or not specified, 'ANI.mat' is used by default  
inPath = path to the .mat file  
        if empty or not specified, IPEMRootDir('code')\Temp is used  
        by default  
inANIName = name of the auditory nerve image array  
        if empty or not specified, 'ANI' is used by default  
inANIFreqName = name of the sample frequency of the auditory nerve image  
        if empty or not specified, 'ANIFreq' is used by default  
inANIFilterFreqsName = name of the filter frequencies array  
        if empty or not specified, 'ANIFilterFreqs' is used  
        by default
```

Output:

```
outANI = the auditory nerve image  
outANIFreq = the sample frequency of the ANI  
outANIFilterFreqs = center frequencies used for calculating the ANI
```

Example:

```
[ANISchum1,ANIFreqSchum1,ANIFilterFreqsSchum1] = ...  
IPEMLoadANI('ANIs','c:\','Schum1','Schum1Freq','Schum1FilterFreqs');
```

Authors:

Koen Tanghe - 20000221

IPEMMECAAnalysis

Usage:

```
[outValues,outValuesFreq,outPeriods] = ...
    IPEMMECAAnalysis(inSignal,inSampleFreq,inMinPeriod,inMaxPeriod, ...
        inStepSize,inHalfDecayTime,inPlotFlag)
```

Description:

Performs a periodicity analysis of a (multi-channel) signal using the Minimal Energy Change algorithm. This part calculates the difference values. Use IPEMMECFindBestPeriods to find the best matching periods.

Input arguments:

```
inSignal = (multi-channel) signal to be analyzed
inSampleFreq = sample frequency of inSignal (in Hz)
inMinPeriod = minimum period to look for (in s)
    if empty, 1/inSampleFreq is used by default
inMaxPeriod = maximum period to look for (in s)
inStepSize = interval between successive evaluations (in s)
    if empty or not specified, 1/inSampleFreq is used by default
    (this corresponds to a step size of 1 sample)
inHalfDecayTime = half decay time for leaky integration (in s)
    if empty or not specified, no integration is performed
    (this corresponds to inHalfDecayTime = 0)
inPlotFlag = if non-zero, plot will be generated
    if empty or not specified, 0 is used by default
```

Output:

```
outValues = cell array containing the calculated values for each channel
    and for each period at each evaluated moment
    (outValues{i}{j,k} represents the value calculated for the j-th
    period at the time corresponding with index k for channel i)
outValuesFreq = sample frequency of outValues (in Hz)
outPeriods = column vector containing the periods for which values were
    calculated (in s)
```

Example:

```
[Values,ValuesFreq,Periods] = IPEMMECAAnalysis(RMS,RMSFreq,0.5,5,[],1.5);
```

Authors:

Koen Tanghe - 20010225

IPEMMECEExtractPatterns

Usage:

```
[outPatterns,outPatternLengths,outPatternFreq] = ...
    IPEMMECEExtractPatterns(inAnalyzedSignal,inSampleFreq,inPeriods, ...
        inBestPeriodIndices,inAnalysisFreq, ...
        inRescalePatterns,inPlotFlag);
```

Description:

Extracts the best pattern from the original signal using the results of an IPEMMECAanalysis run.

Input arguments:

inAnalyzedSignal = (multi-channel) signal that was used for IPEMMECAanalysis
 inSampleFreq = sample frequency of inAnalyzedSignal (in Hz)
 inPeriods = column vector with the periods that were evaluated (in s)
 inBestPeriodIndices = 2D matrix containing the indices in inPeriods of the best periods (each row represents a channel)
 (if inBestPeriodIndices is single-channel while inAnalyzedSignal is multi-channel, inBestPeriodIndices is used for each channel)
 inAnalysisFreq = sample frequency of the IPEMMECAanalysis output (in Hz)
 inRescalePatterns = if non-zero, the patterns are rescaled between 0 and 1
 if empty or not specified, 0 is used by default
 inPlotFlag = if non-zero, plots are generated
 if empty or not specified, 0 is used by default

Output:

outPatterns = cell array containing the extracted patterns for each channel at each moment in time
 (outPatterns{i}(j,k) represents the j-th sample of the pattern extracted at the time corresponding with index k for channel i)
 outPatternLengths = 2D array containing the length of the patterns in outPatterns (in samples) (each row represents a channel)
 outPatternFreq = sample frequency of outPatterns (same as inAnalysisFreq)
 (each pattern itself is of course sampled at inSampleFreq)

Remarks:

The input arguments inPeriods, inBestPeriodIndices and inAnalysisFreq are direct results of IPEMMECAanalysis and IPEMMECFindBestPeriods.

Example:

```
[P,PLengths,PFreq] = ...
IPEMMECEExtractPatterns(RMS,RMSFreq,Periods,BestPeriodIndices, ...
AnalysisFreq,0,1);
```

Authors:

Koen Tanghe - 20010925

IPEMMECFindBestPeriods

Usage:

```
[outBestPeriodIndices,outSampleFreq] = ...
    IPEMMECAalysis(inValues,inValuesFreq,inPeriods,inPlotFlag, ...
        inMergeType,inDecisionType)
```

Description:

Finds the best matching periods corresponding to the given values obtained from IPEMMECAalysis.

Input arguments:

inValues = cell array containing difference values from IPEMMECAalysis
 inValuesFreq = sample frequency for inValues
 inPeriods = periods corresponding to the rows of the elements in inValues
 (in s)
 inPlotFlag = if non-zero, plot will be generated
 if empty or not specified, 0 is used by default
 inMergeType = type of merging in case of multi-channel input
 supported types are:
 'sum' = decision is made on sum of given values
 'separate' = decision is made for each channel separately
 if empty or not specified, 'sum' is used by default
 inDecisionType = type of decision used to find the best matching periods
 supported types are:
 'min' = period corresponding to smallest difference value
 if empty or not specified, 'min' is used by default

Output:

outBestPeriodIndices = indices corresponding to the best matching periods
 (so Periods(outBestPeriodIndices(i,j),1) is the best
 matching period for channel i at moment j, where
 Periods is the row matrix from IPEMMECAalysis)
 outSampleFreq = sample frequency of outBestPeriodIndices

Remarks:

For a single channel inValues, the merge type has no influence of course.

Example:

```
[Best,BestFreq] = IPEMMECAAnalysis(Values,ValuesFreq,Periods,1);
```

Authors:

Koen Tanghe - 20010225

IPEMMECReSynthUI

Usage:

```
IPEMMECReSynthUI(inAction,inData)
```

Description:

User Interface callback function for interactively handling the resynthesis of MEC analysis results.

Input arguments:

```
inAction = string representing the type of action to take
          (should be 'Start' when launched)
inData = struct containing the data needed by the interface
         contains the following fields:
           Sound = original sound
           SampleFreq = sample frequency of Sound
           Periods = periods that were analyzed
           BestPeriodIndices = indices for detected best periods
           AnalysisFreq = sample frequency of BestPeriodIndices
           P = extracted patterns (see remark below)
           PLenghts = lengths of the patterns in P
           PFreq = sample frequency of P (same as AnalysisFreq)
           OriginalFreq = sample frequency of contents of P
           NoiseBands = noise bands to use for resynthesis
           ANIFilterFreqs = filter frequencies
                           (empty in case of non-ANI data)
```

Remarks:

1. For more info on the data stored in inData, see input and output arguments of the functions IPEMMECAAnalysis, IPEMMECExtractPatterns, IPEMMECSynthesis.
2. For the 'P' field, you have to specify P between curly braces:

```
DataStruct = struct('Sound',s,'SampleFreq',fs,...,'P',{P},...);
```

This is because of Matlab's way of handling cell arrays in struct definitions.

Example:

```
UIData = struct('Sound',s,'SampleFreq',fs,'Periods',Periods,...
               'BestPeriodIndices',BestPeriodIndices,'AnalysisFreq',Freq,...
```

```
'P',{P}, 'PLengths', PLengths, 'PFreq', PFreq, 'OriginalFreq', ...
RMSFreq, 'NoiseBands', NoiseBands, 'ANIFilterFreqs', []);
IPEMMECReSynthUI('Start', UIData);
```

Authors:

Koen Tanghe - 20010925

IPEMMECSaveResults

Usage:

```
IPEMMECSaveResults(inFileName,inFilePath,inSound,inSampleFreq, ...
    inPeriods,inBestPeriodIndices,inAnalysisFreq, ...
    inPatterns,inPatternLengths,inPatternFreq,
    inOriginalSignalFreq,inNoiseBands,inANIFilterFreqs, ...
    inAppend,inSuffix)
```

Description:

Utility function that saves the results of IPEMMECExtractPatterns, so that resynthesis can still be done at a later date, after reloading this data.

Input arguments:

```
inFileName = name of the .mat file to save the results to
inFilePath = full path to the location where the file must be saved
            if empty, IPEMRootDir('output') is used by default
inSound = original sound
inSampleFreq = sample frequency of inSound (in Hz)
inPeriods = (see IPEMMECExtractPatterns)
inBestPeriodIndices = (idem)
inAnalysisFreq = (idem)
inPatterns = (see IPEMMECSynthesis)
inPatternLengths = (idem)
inPatternFreq = (idem)
inOriginalSignalFreq = (idem)
inNoiseBands = (idem)
inANIFilterFreqs = ANI filter frequencies (empty if non-ANI data)
inAppend = if non-zero, the data is appended to the data already in the file
            if zero, empty or not specified, all old data is cleared
inSuffix = suffix that should be added to the general names
            (this facilitates saving multiple result sets)
            if empty or not specified, no suffix is used by default
```

Remarks:

1. For generality, the given input arguments are saved under the following names:
 'Sound', 'SampleFreq', 'Periods', 'BestPeriodIndices', 'AnalysisFreq', 'P',
 'PLengths', 'PFreq', 'OrginalFreq', 'NoiseBands' and 'ANIFilterFreqs'.
 A suffix can be added to these names if inSuffix is given.
2. You can later reload the data and perform another resynthesis using

IPEMMECSynthesis like shown in the example below

Example:

```
IPEMMECSaveResults('MEC_ligeti.mat','e:\temp',s,fs,Periods,...  
    BestPeriodIndices,AnalysisFreq,P,PLengths,PFreq,...  
    RMSFreq,NoiseBands,ANIFilterFreqs);  
...  
DataStruct = load('e:\temp\MEC_ligeti.mat');  
IPEMMECReSynthUI('Start',DataStruct);
```

Authors:

Koen Tanghe - 20010228

IPEMMECSynthesis

Usage:

```
[outResynthesizedSound,outDetectedPeriods,outIndividualSounds] = ...
    IPEMMECSynthesis(inPatterns,inPatternLengths,inPatternFreq, ...
        inOriginalSignalFreq,inSelectionTime,inDuration, ...
        inOutputFreq,inNoiseBands,indBLevel,inEnhanceContrast, ...
        inReSynthWidth,inPlotFlag);
```

Description:

Generates an AM modulated noise signal constructed from a repetition of the pattern found at the specified time moment.

Input arguments:

inPatterns = the patterns extracted from the original signal
inPatternLengths = the lengths of the patterns in **inPatterns** (in samples)
inPatternFreq = sample frequency for **inPatterns** (in Hz)
inOriginalSignalFreq = sample frequency of the contents of the patterns
 (in Hz)
inSelectionTime = moment in time specifying the pattern to be used (in s)
inDuration = wanted duration of the synthesized sound (in s)
 if empty or not specified, 3 times the maximum pattern length
 is used by default
inOutputFreq = wanted sample frequency (in Hz)
 if empty or not specified, 22050 is used by default
inNoiseBands = 2 column matrix with wanted noise bands, one row per channel
 (in Hz)
 if empty or not specified, broad band noise is used by default
indBLevel = wanted dB level for the resynthesized sound (in dB)
 if empty or not specified, -20 dB is used by default
inEnhanceContrast = if non-zero, the pattern used for modulation of a noise
 band is first reshaped according to:

$$(P/\max(\text{abs}(P))).^3 * \max(\text{abs}(P))$$

 if empty or not specified, 0 is used by default
inReSynthWidth = if non-zero, specifies the width of one piece of
 resynthesized noise in samples
 (see IPEMGenerateBandPassedNoise for more details)
 if empty or not specified, the entire wanted duration
 is used by default
inPlotFlag = if non-zero, plot are generated
 if empty or not specified, 0 is used by default

Output:

```
outResynthesizedSound = the resynthesized sound  
outDetectedPeriods = duration of the periodic patterns that were used in  
each channel (in s)  
outIndividualSounds = resynthesized sounds per channel
```

Example:

```
[ss,periods,sschan] = IPEMMECSynthesis(P,PL,PFreq,RMSFreq,3,10,22050);
```

Authors:

Marc Leman - 20000907 (original prototype version)
Koen Tanghe - 20010926 (code generalization + documentation)

IPEMNormalize

Usage:

```
outSignal = IPEMNormalize (inSignal, inNormalizeSeparately)
```

Description:

This function normalizes the given (multi-channel) signal, so that all values are in the range [-1,1]. Channels are represented by rows.

Input arguments:

inSignal = the signal to be normalized (each row represents a channel)
inNormalizeSeparately = if non-zero, each channel is normalized separately
 if zero, empty or not specified, all channels are multiplied with the same normalisation factor,
 so that channel level ratio's are conserved

Output:

```
outSignal = the normalized signal
```

Example:

```
Signal = IPEMNormalize(Signal);
```

Authors:

Koen Tanghe - 20000919

IPEMOnsetPattern

Usage:

```
[outOnsetPattern,outOnsetPatternFreq] = ...
    IPEMOnsetPattern(inSignal,inSampleFreq)
```

Description:

Integrate-and-fire neural net for onset-detection.

Based on an article by Leslie S. Smith (1996)

Neuron dynamics are given by:

$$\frac{dA}{dt} = I(t) - \text{diss}*A \quad \text{with: } A = \text{neurons' accumulated value}$$

I = input
diss = dissipation

Input arguments:

```
inSignal = a matrix of size [N M], where N = the number of channels
           and M = the number of samples
inSampleFreq = a scalar representing the frequency at which the signal was
               sampled (in Hz)
```

Output:

```
outOnsetPattern = a matrix of size [N M] where an element is 1 if an onset
                  was detected
outOnsetPatternFreq = sample frequency of result (same as inSampleFreq)
```

Example:

```
[OnsetPattern,OnsetPatternFreq] = IPEMOnsetPattern(Signal,SampleFreq);
```

Authors:

Koen Tanghe - 20000510

IPEMOnsetPatternFilter

Usage:

```
[outOnsetSignal,outOnsetSignalFreq] = ...
    IPEMOnsetPatternFilter (inOnsetPattern, inSampleFreq)
```

Description:

This function takes an (multi-channel) OnsetPattern generated by IPEMOnsetPattern, and outputs a (one-dimensional) signal in which a non-zero value means an onset occurred (the value is a measurement of the likeliness for the onset).

Currently, the following (simple) strategy is followed:
an onset is detected at a certain moment, if:

1. a certain fraction of channels have an onset within a specific period of time
2. the moment falls at least a minimum period of time behind the last detected onset

Input arguments:

```
inOnsetPattern = a matrix of size [N M],  

    where N = the number of channels and  

        M = the number of samples  

inSampleFreq = sample frequency of the input signal (in Hz)
```

Output:

```
outOnsetSignal = a 1D signal having a non-zero value at detected onset-times  

    (the clearer the onset, the higher the value)  

outOnsetSignalFreq = sample frequency of the result (same as inSampleFreq)
```

Example:

```
[OnsetSignal,OnsetSignalFreq] = ...
    IPEMOnsetPatternFilter(OnsetPattern,SampleFreq);
```

Authors:

Koen Tanghe - 20031209

IPEMOnsetPeakDetection

Usage:

```
[outOnsetResults,outOnsetResultsFreq] = ...  
    IPEMOnsetPeakDetection (inSignal,inSampleFreq,inPlotFlag)
```

Description:

Returns a pattern having a non-zero value on moments of possible onset peaks

Input arguments:

```
inSignal = a matrix of size [N,M] where N is the number of channels and  
          M is the length in samples  
inSampleFreq = sample frequency of the incoming signal (in Hz)  
inPlotFlag = if non-zero, plots are generated  
             if empty or not specified, 0 is used by default
```

Output:

```
outOnsetResults = a matrix of size [N,M] with a value between 0 and 1  
                  (non-zero where an onset occurs, and proportional to the  
                   importance of the onset)  
outOnsetResultsFreq = sample frequency of the result (same as inSampleFreq)
```

Example:

```
[OnsetResults,OnsetResultsFreq] = IPEMOnsetPeakDetection(Signal,SampleFreq);
```

Authors:

Koen Tanghe - 20010122

IPEMOnsetPeakDetection1Channel

Usage:

```
[outPeakIndices, outImportances] = ...
    IPEMOnsetPeakDetection1Channel(inSignal, inSampleFreq)
```

Description:

Returns the indices of the "important" peaks in the input signal, together with an indication of the "importance" of the peak.

Input arguments:

inSignal = the signal to be scanned for peaks
inSampleFreq = the sample frequency of the signal (in Hz)

Output:

outPeakIndices = indices within the input signal where the peaks are found
outImportance = value between 0 and 1 showing the "importance" of a peak

Example:

```
[PeakIndices, Importances] = ...
    IPEMOnsetPeakDetection1Channel(Signal, SampleFreq);
```

Authors:

Koen Tanghe - 20000510

IPEMPeriodicityPitch

Usage:

```
[outSignal,outSampleFreq,outPeriods,outFANI] = IPEMPeriodicityPitch ...
    (inMatrix,inSampleFreq,inLowFrequency, ...
     inFrameWidth,inFrameStepSize,inPlotFlag)
```

Description:

This function calculates the periodicity pitch of a matrix containing neural firing patterns in different auditory channels. It is based on the idea that periodicity pitch is calculated as the summed autocorrelation over bandpass filtered fluctuations in auditory channels. Due to the fact that the output of the auditory model gives the envelopes of the neural firing probabilities (< 1250 Hz) it suffices to use a low-pass filter in order to obtain the pitch e.g. between 80 - 1250 Hz. Only the lowest frequency can be changed by inLowFrequency.

Input arguments:

inMatrix = the input signal is a matrix of size [N M] where
 N is the number of auditory channels (<= 40) and
 M is the number of samples
 the input signal is the output of the auditory model (ANI).
 inSampleFreq = sample frequency of the input signal (in Hz)
 inLowFrequency = cutoff frequency (in Hz) of a first order lowpass filter
 applied before calculating the autocorrelation
 if empty or not specified, 80 is used by default
 inFrameWidth = width of the frame used for the accumulation of the
 autocorrelation (in s)
 if empty or not specified, 0.064 is used by default
 inFrameStepSize = stepsize or timeinterval between two inFrameWidths (in s)
 if empty or not specified, 0.010 is used by default
 inPlotFlag = if non-zero, plots are generated
 if empty or not specified, 0 is used by default

Output:

outSignal = periodicity pitch: a matrix of size
 inFrameWidth * length(inMatrix) / outSampleFreq
 outSampleFreq = sampling rate, equal to inSampleFreq/inFrameStepSize (in Hz)
 outPeriods = analyzed periods (in s)
 outFANI = bandpass filtered auditory nerve images (at the original sample
 frequency)

Remarks:

As for any frame-based function, the first value of the output signal is the value calculated for the first complete frame in the input signal.

This means the following:

if you have an input signal of length 1 s at a sample frequency of 1000 Hz, a frame width of 0.050 s, and a frame step size of 0.010 s, then there will be $\text{ceil}(((1 - 0.050)*1000 + 1)/(0.010*1000)) = 96$ values in the output signal, where the first value corresponds to the first complete frame (the interval 0 to 0.050 s)

Example:

```
[PP,PPFreq,PPPeriods,PPFANI] = IPEMPeriodicityPitch(ANI,ANIFreq);
```

Authors:

Marc Leman - 20011204 (originally made)

Koen Tanghe - 20050119

IPEMPlaySoundWithCursor

Usage:

```
IPEMPlaySoundWithCursor(inSignal,inSignalFreq,inAxisHandles,inTimeRange,  
inTimeOffset,inSpeedFactor)
```

Description:

Plays the given sound signal while showing a moving cursor in the given axes.

Input arguments:

inSignal = Sound signal.
inSignalFreq = Sample frequency of inSignal (in Hz).
inAxisHandles = Handles of the axes for which a pointer should be shown.
If empty or not specified, all axes of the current figure
will be used by default (if no current figure, -1 is used).
If -1, the sound file is drawn in a new figure.
inTimeRange = Range of playback in the form [starttime endtime] (in s).
If empty or not specified, the entire signal will be used.
inTimeOffset = Time offset of first sample (in s). This is useful when the
given signal is part of a larger signal.
If empty or not specified, 0 is used by default.
inSpeedFactor = Factor by which the sound should be sped up (>1) or slowed
down (<1). If empty or not specified, 1 is used by default.

Remarks:

The plots associated with the given handles are supposed to have an x-axis
that has a time scale in seconds.

Example:

```
IPEMCalcRMS(s,fs,0.040,0.020,1);  
IPEMPlaySoundWithCursor(s,fs);
```

Authors:

Koen Tanghe - 20040330

IPEMPlotMultiChannel

Usage:

```
IPEMPlotMultiChannel(inData,inSampleFreq,inTitle,inXLabel,inYLabel, ...
    inFontSize,inChannelLabels,inChannelLabelStep, ...
    inMinY,inMaxY,inType,inTimeOffset,inYScaleFactor)
```

Description:

Plots a multi-channel signal.

Input arguments:

```
inData = the multi-channel signal (each row represents a channel)
inSampleFreq = sample frequency of the data (in Hz)
    if empty or not specified, 1 is used by default
inTitle = title for the plot
    if empty or not specified, '' is used by default
inXLabel = name for the X-values
    if empty or not specified, '' is used by default
inYLabel = name for the Y-values
    if empty or not specified, '' is used by default
inFontSize = size of the font to be used for title (size of axes labels will
    be set to inFontSize-2)
    if empty or not specified, no special fontsize is set by default
inChannelLabels = labels for the rows of the multi-channel data
    if empty or not specified, 1:Rows(inData) is used
    by default
inChannelLabelStep = step size for displaying the labels of the channels:
    labels 1:inChannelLabelStep:NumOfRows will be used
    as ticks
    if -1, the initial automatic tick division is kept
    if empty or not specified, 1 is used by default
inMinY = minimum Y-value to be used for scaling the line plot of a channel
    if empty or not specified min(min(inData)) is used by default
inMaxY = maximum Y-value to be used for scaling the line plot of a channel
    if empty or not specified max(max(inData)) is used by default
inType = plot type to use:
    0 : uses line plots for each channel, starting from tick label
    1 : uses imagesc
    2 : uses line plots for each channel, centered around tick label
    if empty or not specified, 0 is used by default
inTimeOffset = time offset to use for the first sample (in s): the first
    sample will be assumed to be taken at time inTimeOffset
```

```
    if empty or not specified, 0 is used by default
inYScaleFactor = factor to be used on the Y values (can be used to scale the
                  range of Y values to better show the signal dynamics)
    if empty or not specified, 1 is used by default
```

Output:

A graph in the current subplot of the current figure.

Example:

```
IPEMPlotMultiChannel(ANI,ANIFreq,'Auditory nerve image','Time (in s)',...
    'Auditory channels (center freqs. in Hz)',14,ANIFilterFreqs,3);
```

Authors:

Koen Tanghe - 20040322

IPEMReadSoundFile

Usage:

```
[outSignal,outSampleFreq,outSize,outFile] = ...
    IPEMReadSoundFile(inFileName,inFilePath,inSection,inPlotFlag)
```

Description:

Reads (part of) a sound file into memory.

Input arguments:

```
inFileName = name of the sound file (with extension)
            if empty or not specified, the standard "open file" dialog box
            can be used to select a sound file (if the user then presses
            'Cancel', all outputs will be empty)
inFilePath = path to the sound file (if it does not exist, the current
            directory is used)
            use the empty string '' if inFileName already contains a full
            path specification (no extra path will be prepended)
            if empty or not specified, IPEMRootDir('input')\Sounds is used
            by default (if it exists)
inSection = section of the sound file to read, given as [start end] (in s)
            if 'size', no samples are read: outSignal will be empty, and
            only outSampleFreq and outSize will have relevant values
            if empty or not specified, the entire sound file is read
inPlotFlag = if non-zero, a plot is generated
            if empty or not specified, 0 is used by default
```

Output:

```
outSignal = the sound signal (each row represents a channel)
outSampleFreq = the frequency at which the sound was sampled (in Hz)
outSize = the size of the entire sound file returned as
            [NumSamples NumChannels] (optional)
            where NumSamples = number of samples in each channel
            NumChannels = number of channels
outFile = entire path specification of the sound file
```

Remarks:

The 'end' in the section specification is exclusive:
 if you want to read the first second of a sound file sampled at 1000 Hz,

you should specify [0 1] for inSection (and not [0 0.999]). This will read samples 1 to 1000, which is a full second (the sample at time 1 s is not read!).

Example:

```
[s,fs] = IPEMReadSoundFile('schum1.wav');
```

Authors:

Koen Tanghe - 20040409

IPEMReplaceSubStringInFileNames

Usage:

```
IPEMReplaceSubStringInFileNames(inDir,inOriginalString,inNewString,...  
                                inRecursive,inAffectExtensions)
```

Description:

Replaces a substring in the filenames of the files contained in a specific directory (and its subdirectories if needed) by a new substring.

Input arguments:

```
inDir = full directory path to the root directory for which the changes  
       should be applied  
inOriginalString = substring to be replaced in the filenames  
                   if empty or not specified, '_' is used by default  
inNewString = substring by which inOriginalString will be replaced  
                   if empty or not specified, '' is used by default  
inRecursive = if non-zero, subdirectories are also affected  
                   if empty or not specified, 0 is used by default  
inAffectExtensions = if non-zero, extensions of filenames are also affected  
                   if empty or not specified, 0 is used by default
```

Example:

```
% Remove all underscores in the filenames (not the extensions) of all files  
% under the given directory and its subdirectories  
IPEMReplaceSubStringInFileNames('D:\Koen\Temp\Test','_','','',1,0);
```

Authors:

Koen Tanghe - 20010522

IPEMRescale

Usage:

```
outSignal = IPEMRescale (inSignal,inLowLevel,inHighLevel,inScaleGlobally)
```

Description:

This function rescales each channel of the incoming signal between the given lower and upper bounds. Signal channels are represented by rows.

Input arguments:

inSignal = the signal that is to be rescaled (each row represents a channel)
inLowLevel = the new value for the lowest value in the incoming signal
 if empty or not specified, 0 is used by default
inHighLevel = the new value for the highest value in the incoming signal
 if empty or not specified, 1 is used by default
inScaleGlobally = if non-zero, all channels are scaled in exactly the
 same way
 if zero, each channel is scaled separately
 if empty or not specified, 0 is used by default

Output:

```
outSignal = the rescaled signal
```

Remarks:

1. For constant signals, inLowLevel is used as constant output level
2. You can specify a value for inLowLevel that is higher than the inHighLevel, in which case the signal will be inverted and scaled between the two given levels.

Example:

```
Signal = IPEMRescale(Signal,-1,1);
```

Authors:

Koen Tanghe - 20000224

IPEMReshape

Usage:

```
outSignal = IPEMReshape(inData,inShaper,inMethod,inSymmetry,inPlotFlag)
```

Description:

Reshapes the incoming data by substituting each original value by a new value using the given shaper matrix.

Input arguments:

```
inData = 1D data vector to be reshaped
inShaper = shaper matrix: two column matrix of which the first column
            represents the original value and the second column represents
            its substituted value
inMethod = string with the method to be used for interpolating substituted
            values (see description of methods in Matlab's interp1 function)
            if empty or not specified, 'linear' is used by default
inSymmetry = if 1, inShaper is assumed to be one half of a curve symmetric
            around the point zero:  $Y(-X) = -Y(X)$ 
            if 2, inShaper is assumed to be one half of a curve symmetric
            around the Y-axis:  $Y(-X) = Y(X)$ 
            if empty or not specified, no symmetry is assumed by default
inPlotFlag = if non-zero, a plot is generated
            if empty or not specified, 0 is used by default
```

Output:

```
outSignal = reshaped data
```

Remarks:

Only values between the `min(inShaper(:,1))` and `max(inShaper(:,1))` can be substituted, so always make sure to include these bounds and their substitutes in `inShaper`.

In case you make use of `inSymmetry`, only `max(inShaper(:,1))` will play a role in defining the region where values can be substituted.

Example:

```
X = 0:0.1:1;
Y = X.^3;
s2 = IPEMReshape(s,[X' Y'], 'linear',1,1);
```

Authors:

Koen Tanghe - 20010110

IPEMRippleFilter

Usage:

```
outSignal = IPEMRippleFilter(inSignal,inSampleFreq,inDelay, ...
                             inAttenuation,inSync)
```

Description:

Filters the incoming signal with a ripple filter according to the following scheme:

$$\text{OUT} = \text{IN} / (1 + \text{inAttenuation} * z^{-m})$$

where $m = \text{round}(\text{inDelay} * \text{inSampleFreq})$

In the frequency domain, this produces a rippled frequency response, in which either the peaks (if $\text{inSync} == 1$) or the valleys (if $\text{inSync} == 0$) of the ripples are at frequencies $n.F_0$ ($F_0 = 1/\text{inDelay}$).

Input arguments:

- `inSignal` = input signal (each row represents a channel)
- `inSampleFreq` = input signal sample frequency (in Hz)
- `inDelay` = delay of feedback branch (in s)
- `inAttenuation` = attenuation ratio (0 to 1)
- `inSync` = if 1, the ripple's peaks are at frequencies $n/delay$
if 0, the ripple's valleys are at frequencies $n/delay$

Output:

`outSignal` = the filtered signal

Example:

```
Signal = IPEMRippleFilter(Signal,SampleFreq,1/100,0.95,1);
```

Authors:

Koen Tanghe - 20000209

IPEMRootDir

Usage:

```
outRootDirPath = IPEMRootDir(inDirType)
```

Description:

This function returns a default root directory for the requested type. These directories are used by some of the IPEM functions.

Input arguments:

inType = string specifying the type of info for which you want the standard directory (the supported types are defined below)

Output:

outRootDirPath = the requested directory path

Remarks:

*** Description of supported types ***

'input'

The directory that will be used as the root of the subdirectories containing data to be processed by the IPEM Toolbox code.
(e.g. sound files could be stored in a subdirectory called 'Sounds', located in this 'input' directory)

'output'

Same as 'input', but for output of the IPEM Toolbox code.
In some special cases, the directory of the processed file is used (but this will always be explicitly mentioned in the comments of such a function)

'code'

This is the directory where the IPEM Matlab code is located and can NOT be set by setpref.
(e.g. 'E:\Code\Matlab\IPEMToolbox')

*** Supports preferences ***

You can specify the default 'input' and 'output' root directories by typing:

```
setpref('IPEMToolbox','RootDir_Input','directory1')
```

and/or

```
setpref('IPEMToolbox','RootDir_Output','directory2')
```

where directory1 and directory2 should be replaced by the complete paths to existing directories on your system.

If you don't set these directory preferences, the 'Temp' subdirectory in your 'IPEMToolbox' directory is used by default.

Authors:

Koen Tanghe - 20010117

IPEMRotateMatrix

Usage:

```
outMatrix = IPEMRotateMatrix (inMatrix,inRows,inColumns)
```

Description:

Rotates a matrix along its rows and/or columns.

Input arguments:

```
inMatrix = matrix to process  
inRows = number of rows by which to rotate  
        positive values mean rotation towards increasing row numbers  
inColumns = number of columns by which to rotate  
           positive values mean rotation towards increasing column numbers
```

Output:

```
outMatrix = the rotated matrix
```

Example:

```
RotMat = IPEMRotateMatrix([1 2 3; 4 5 6; 7 8 9],-1,2);
```

```
--> RotMat == [5 6 4; 8 9 7; 2 3 1]
```

Authors:

Koen Tanghe - 20000209

IPEMRoughnessFFT

Usage:

```
[outRoughness,outSampleFreq,outFFTMATRIX1,outFFTMATRIX2] =
IPEMRoughnessFFT(inANI,inANIFreq,inANIFilterFreqs, ...
inFrameWidth,inFrameStepSize,inPlotFlag)
```

Description:

Based on Leman (2000) (paper on calculation and visualization of Synchronization & Roughness)

This function calculates three outputs:

1. The total energy
2. The energy in all channels
3. The energy spectrum of the neuronal synchronization

Input arguments:

inANI = auditory nerve image to process, where each row represents a channel

inANIFreq = sample frequency of the input signal (in samples per second)

inANIFilterFreqs = filterbank frequencies used by the auditory model

inFrameWidth = the width of the window for analysing the signal (in s)
if empty or not specified, 0.2 s is used by default

inFrameStepSize = the stepsize or time interval between two
inFrameWidthInSamples (in s)

if empty or not specified, 0.02 s is used by default

inPlotFlag = if non-zero plots are generated
if empty or not specified, 0 is used by default

Output:

outRoughness = roughness over signal

outSampleFreq = sampling rate of outRoughness (in Hz)

outFFTMATRIX1 = visualisation of energy over channels

outFFTMATRIX2 = visualisation of energy spectrum for synchronization
(synchronisation index SI)

Remarks:

For now, the roughness values are dependend on the used frame width.

So, to make usefull comparisons, only results obtained using the same frame width should be compared (this should be fixed in the future...)

Example:

```
[Roughness,RoughnessFreq] = ...
IPEMRoughnessFFT(ANI,ANIFreq,ANIFilterFREQs,0.20,0.02,1);
```

Authors:

Marc Leman - 20000910 (made)

Koen Tanghe - 20010816 (bug fix + documentation corrections + code cleanup)

IPEMRoughnessOfSoundPairs

Usage:

```
[outMeanRoughness, outSortedList, outRoughness, outRoughnessFreq, ...
outUsedSection] = ...
    IPEMRoughnessOfSoundPairs(inFiles, inPaths, inSection, ...
        inSaveSoundCombinations, inOutputPath, ...
        inUseReference, inPlotFlag)
```

Description:

Calculates roughness of 2-by-2 combinations of sounds.

Input arguments:

inFiles = set of sound files to be analyzed (cell array of strings)
 inPaths = paths for the sound files
 either a cell array of strings containing a path for each sound
 file, or a single character array containing one path for all
 sound files
 if empty or not specified, IPEMRootDir('input')\Sounds is used
 by default
 inSection = section of the sound to be used (given as [start end] in s,
 where Inf for the end parameter means the minimum of all sound
 file lengths)
 if empty or not specified, the (minimum of the) full length of
 the sounds is used by default
 inSaveSoundCombinations = if non-zero, the sound combinations are saved to
 inOutputPath
 if empty or not specified, 1 is used by default
 inOutputPath = path where the sound combinations should be stored
 if empty or not specified, IPEMRootDir('output') is used
 by default
 inUseReference = if non-zero, a reference value is used depending on the dB
 level (see IPEMGetRoughnessFFTReference)
 if empty or not specified, 0 is used by default
 inPlotFlag = if non-zero, plots outRoughness as a 2D map
 if empty or not specified, 1 is used by default

Output:

outMeanRoughness = mean of outRoughness over the complete section
 outSortedList = sound combinations sorted according to roughness:
 column 1 = roughness

columns 2 and 3 = number of the first resp. second sound
outRoughness = 3D matrix containing the calculated roughness over the
combinations of sounds, where element outRoughness(i,j,:)
are the roughness values calculated over the combination
of sound i with sound j
outRoughnessFreq = sample frequency of outRoughness (Hz)
outUsedSection = section of the sound that was eventually used
(same as inSection, except when inSection(2) = Inf)

Remarks:

1. The sound signals (say s1 and s2) are mixed as follows:
 $s1 = \text{IPEMAadaptLevel}(s1, -20);$
 $s2 = \text{IPEMAadaptLevel}(s2, -20);$
 $s = \text{IPEMAadaptLevel}(s1+s2, -20);$
2. All sound files should be of the same sample frequency...

Example:

```
theFiles = {'sound1.wav', 'sound2.wav', 'sound3.wav'};
thePaths = {'E:\Koen\Data', 'e:\Koen\Data', 'e:\Koen\Data\Sounds\New'};
[MeanRoughness, SortedList] = ...
    IPEMRoughnessOfSoundPairs(theFiles, thePaths, [0.2 Inf]);
```

Authors:

Koen Tanghe - 20011004

IPEMSaveANI

Usage:

```
IPEMSaveANI(inANI,inANIFreq,inANIFilterFreqs,inName,inPath,...  
inANIName,inANIFreqName,inANIFilterFreqsName,inAppend)
```

Description:

Saves an auditory nerve image and its corresponding sample frequency and filter frequencies to a .mat file on disk.

Input arguments:

```
inANI = auditory nerve image array  
inANIFreq = sample frequency of the auditory nerve image  
inANIFilterFreqs = filter frequencies array  
inName = the name of the .mat file for storing the nerve image  
         if empty or not specified, 'ANI.mat' is used by default  
inPath = path to the .mat file  
         if empty or not specified, IPEMRootDir('code')\Temp is used  
         by default  
inANIName = name for the auditory nerve image array in the mat file  
         if empty or not specified, 'ANI' is used by default  
inANIFreqName = name for the sample frequency  
         if empty or not specified, 'ANIFreq' is used by default  
inANIFilterFreqsName = name for the filter frequencies array  
         if empty or not specified, 'ANIFilterFreqs' is used  
         by default  
inAppend = if 1, the variables are appended to the mat file (if it already  
           exists) otherwise, a new mat file is created  
         if empty or not specified, 0 is used by default
```

Example:

```
IPEMSaveANI(ANI,ANIFreq,ANIFilterFreqs,'ANIs','c:\','Schum1','FreqSchum1',...  
'FilterFreqsSchum1');
```

Authors:

Koen Tanghe - 20000515

IPEMSaveVar

Usage:

```
IPEMSaveVar (inFile,inName,inValue)
```

Description:

Saves a variable with name inValue and value inValue to a mat file.
If the file did not already exist, a new file is created first.

Input arguments:

```
inFile = full name (with path) of the mat file (the .mat extension will be  
added automatically if it's not there already)  
inName = name for the variable  
inValue = value for the variable
```

Remarks:

This thing is NOT COMPILABLE because of call stack/workspace differences
between M-file evaluation and mex-evaluation !

Example:

```
IPEMSaveVar('E:\Koen\Data\music.mat','NerveImage',ANI);
```

Authors:

Koen Tanghe - 20000209

IPEMSetFigureLayout

Usage:

```
IPEMSetFigureLayout(inFigureHandle)
```

Description:

Sets the layout of the figure to the IPEM defaults.

Input arguments:

```
inFigureHandle = handle of the figure to be adjusted  
if 'all' is specified, all open figures will be adjusted  
if empty or not specified, the current figure will be  
adjusted
```

Example:

```
FigH = figure;  
% ... some plots ...  
IPEMSetFigureLayout(FigH);
```

Authors:

Koen Tanghe - 20010129

IPEMSetup

Usage:

`IPEMSetup`

Description:

This function sets up the Matlab environment for working with the IPEM Toolbox.

Example:

`IPEMSetup;`

Authors:

Koen Tanghe - 20011204

IPEMSetupPreferences

Usage:

```
IPEMSetupPreferences(inReset)
```

Description:

This function sets up preferences for the IPEM Toolbox.

Use `getpref('IPEMToolbox', 'preference')` to get the value of a preference.
Use `setpref('IPEMToolbox', 'preference', value)` to set the value of a preference.

Input arguments:

```
inReset = if non-zero, resets all preferences to their initial values  
          if zero, initializes non-existing preferences, but keeps existing  
          ones  
          if empty or not specified, 0 is used by default
```

Remarks:

List of currently supported preferences:

'RootDir_Input'	- standard input directory
	- initial value: IPEMToolbox\Temp
'RootDir_Output'	- standard output directory
	- initial value: IPEMToolbox\Temp

Example:

```
IPEMSetupPreferences(1);
```

Authors:

Koen Tanghe - 20010117

IPEMShepardTone

Usage:

```
outSignal = IPEMShepardTone(inMainFreq,inDuration,inSampleFreq,...  
                           inPhaseFlag,indBLevel)
```

Description:

This function generates a Shepard tone.

Input arguments:

```
inMainFreq = the main frequency (in Hz)  
inDuration = the duration (in s)  
              if empty or not specified, 1 is used by default  
inSampleFreq = the desired sample frequency for the output signal (in Hz)  
              if empty or not specified, 22050 is used by default  
inPhaseFlag = for choosing whether random phase is to be used or not  
              (1 to use random phase, 0 otherwise)  
              if empty or not specified, 1 is used by default  
indBLevel = dB level of generated tone (in dB)  
              if empty or not specified, no level adjustment is performed
```

Output:

```
outSignal = the signal for the tone
```

Example:

```
Signal = IPEMShepardTone(440,1,22050,1,-20);
```

Authors:

Koen Tanghe - 20010116

IPEMShepardToneComplex

Usage:

```
outSignal = IPEMShepardToneComplex (inToneVector,inDuration,inSampleFreq, ...
                                     inPhaseFlag,indBLevel)
```

Description:

This function generates a tone complex built up of Shepard tones.

Input arguments:

```
inToneVector = a 12 elements vector representing the amplitude for each tone
               C, C#, D, ... in the tone complex (0 = no tone, 1 = full
               amplitude)
inDuration = the duration (in s)
               if empty or not specified, 1 is used by default
inSampleFreq = the desired sample frequency for the output signal (in Hz)
               if empty or not specified, 22050 is used by default
inPhaseFlag = for choosing whether random phase has to be used or not
               (1 to use random phase, 0 otherwise)
               if empty or not specified, 1 is used by default
indBLevel = dB level of generated tone complex (in dB)
               if empty or not specified, no level adjustment is performed
```

Output:

outSignal = the signal for the tone complex

Example:

```
Signal = IPEMShepardToneComplex([1 0 0 0 1 0 0 1 0 0 0 0],1,22050,1,-20);
```

Authors:

Koen Tanghe - 20010116

IPEMShowSoundFile

Usage:

```
IPEMShowSoundFile(inFileName,inFilePath,inNewFigureFlag)
```

Description:

This function shows the waveform of the specified sound file.
Only .wav and .snd sound files are supported.

Input arguments:

```
inFileName = the name of the sound file  
inFilePath = the path to the sound file  
           if empty or not specified, IPEMRootDir('input')\Sounds is used  
           by default  
inFigureFlag = if non-zero or not specified, a new figure is created  
              if 0, the graph will be shown in the current (sub)plot
```

Example:

```
IPEMShowSoundFile('music.wav',[],0);
```

Authors:

Koen Tanghe - 20000209

IPEMSineComplex

Usage:

```
[outSignal,outSubSignals] = ...
    IPEMSineComplex(inSampleFreq,inDuration, ...
                    inFrequencies,inAmplitudes,inPhases,indBLevel)
```

Description:

Generates a signal consisting of pure sinusoids with constant amplitude and phase.

Input arguments:

```
inSampleFreq = wanted sample frequency (in Hz)
inDuration = wanted duration (in s)
inFrequencies = frequency components (in Hz)
inAmplitudes = amplitudes for the frequency components
                if empty or not specified, 1 is used for all components
inPhases = phases for the frequency components (in radians)
                if you specify 'random', random phase between 0 and 2*pi is used
                if empty or not specified, 0 is used for all components
indBLevel = normalization level (in dB) for the generated sound
                (the signals in outSubSignals are not normalized)
                if empty or not specified, no normalization is performed
```

Output:

```
outSignal = the generated signal
outSubSignals = multi-channel signal where each row contains the generated
                sine wave for the corresponding frequency
```

Example:

```
s = IPEMSineComplex(22050,1,[110 220 550 660],[1 0.5 0.8 0.4],'random',-6);
```

Authors:

Koen Tanghe - 20020506

IPEMSnipSoundFile

Usage:

```
[outSignalSegments,outSampleFreq] = ...
IPEMSnipSoundFile(inFileName,inFilePath,inTimeSegments, ...
inWriteSegments,inOutputPath)
```

Description:

Snips the sound file into segments.

Input arguments:

```
inFileName = name of the sound file
inFilePath = location of the sound file
            if empty or not specified, IPEMRootDir('input')\Sounds is
            used by default
inTimeSegments = time segments for cutting up the input signal (in s)
                (each row contains the start and end of a time segment)
                if this is a positive scalar, the signal is divided into
                inTimeSegments parts of equal size + an additional segment
                containing the remaining part of the signal (which will
                contain less than inTimeSegments samples)
                if this is a negative scalar, the signal is divided into
                -inTimeSegments parts where some segments will contain
                1 sample more than other segments
inWriteSegments = if non-zero, the segments are written to separate sound
                  files in inOutputPath (the name of the files is the
                  original name appended with the number of the segment)
                  if empty or not specified, 0 is used by default
inOutputPath = ouput path used when sound files are written
                if empty or not specified, the same directory as the sound
                file is used by default
```

Output:

```
outSignalSegments = the wanted signal segments
outSampleFreq = sample frequency of the original sound file (and segments)
                (in Hz)
```

Example:

```
IPEMSnipSoundFile('dnb_drumloop_1patt.wav',[],-16,'E:\Koen\Data\Temp');
```

Authors:

Koen Tanghe - 20000419

IPEMSnipSoundFileAtOnsets

Usage:

```
IPEMSnipSoundFileAtOnsets (inFileName,inFilePath,inOutputPath)
```

Description:

Calculates the onsets for the specified sound file and cuts the sound file into new sound segments at these onset times.

Input arguments:

```
inFileName = name of the sound file  
inFilePath = location of the sound file  
           if empty or not specified, IPEMRootDir('input')\Sounds is  
           used by default  
inOutputPath = path where the sound segments should be stored  
              if empty or not specified, the same directory as the sound  
              file is used by default
```

Example:

```
IPEMSnipSoundFileAtOnsets('dnb_drumloop_1patt.wav',[],...  
                           'E:\Koen\SnippedSamples');
```

Authors:

Koen Tanghe - 20000209

IPEMStripFileSpecification

Usage:

```
outParts = IPEMStripFileSpecification (inFileSpecification)
```

Description:

Strips a complete file specification into its subparts.
Can also be used to strip directory paths in their subparts.

Input arguments:

inFileSpecification = the complete file specification

Output:

outParts = a struct containing the following information:
outParts.FileName = string containing the name of the file
outParts.Name = name of the file without extension
outParts.Extension = extension of the file
outParts.Path = complete path towards the file
outParts.PathParts = cell array containing the subparts of the path specification

Remarks:

1. Only tested on Windows!
2. If the specification doesn't end on an extension, it is assumed to be a directory path.

Example:

```
theParts = IPEMStripFileSpecification('D:\Koen\Data\Sounds\schum1.wav');
---->
theParts.FileName = 'schum1.wav'
theParts.Name = 'schum1'
theParts.Extension = 'wav'
theParts.Path = 'D:\Koen\Data\Sounds'
theParts.PathParts = {'D:' 'Koen' 'Data' 'Sounds'}
```

Authors:

Koen Tanghe - 20000607

IPEMToneScaleComparison

Usage:

```
[outError,outBestRatioDiffs] = IPEMToneScaleComparison;
```

Description:

Compares tonescales with different subdivisions to the 'ideal case'

Output:

```
outError = error compared to ideal case  
outBestRatioDiffs = difference of the best ratio's with the ideal ones
```

Remarks:

Works for the 3 ideal ratio's 6/5, 5/4, 3/2

Authors:

Koen Tanghe - 19990816

IPEMToolboxVersion

Usage:

```
IPEMToolboxVersion
```

Description:

Returns the version of the IPEM Toolbox

Output:

```
outVersionString = string with full version information  
outMajor = major version number  
outMinor = minor version number  
outType = version type string (either 'beta' or empty)  
outBuildDate = build date of this version using the format 'YYYYMMDD',  
               where YYYY = year, MM = month and DD = day
```

Authors:

Koen Tanghe - 20050120

IPEMWriteSoundFile

Usage:

```
outResult = IPEMWriteSoundFile(inSignal,inSampleFreq,inFileName, ...
                               inFilePath,inSection)
```

Description:

Writes (part of) a sound signal to file.

Input arguments:

```
inSignal = sound signal (each row represents a channel)
inSampleFreq = sample frequency of sound signal (in Hz)
inFileName = name for the sound file (with extension)
             if empty or not specified, the standard "Write file" dialog box
             can be used to select a sound file to write to
inFilePath = path for the sound file
             if empty or not specified, IPEMRootDir('input')\Sounds is used
             by default
inSection = section of the sound signal to be written, given as [start end]
             (in s)
             if empty or not specified, the entire sound signal is written
```

Output:

```
outResult = usually 1
           0 if the user pressed 'Cancel' when selecting a sound file
```

Remarks:

The 'end' in the section specification is exclusive:
 if you want to write the first second of a sound signal sampled at 1000 Hz,
 you should specify [0 1] for inSection (and not [0 0.999]). This will read
 samples 1 to 1000, which is a full second (the sample at time 1 s is not
 written!).

Always writes 16-bit.

Example:

```
IPEMWriteSoundFile(s,fs,'test.wav');
```

Authors:

Koen Tanghe - 20040323

Demo functions

MEC pattern extraction application

IPEMDemoMECRhythmExtraction

Usage:

IPEMDemoMECRhythmExtraction

Description:

Starts the demonstration of the MEC algorithm used for the extraction of rhythmic patterns in sound.

Authors:

Koen Tanghe - 20020912

IPEMDemoStartMEC

Usage:

```
[outOriginalSound,outResynthesizedSound,outSampleFreq,outDetectedPeriods,...  
outIndividualSounds] = ...  
IPEMDemoStartMEC(inSoundFile,inSoundFilePath,inMinPeriod,inMaxPeriod,...  
inSelectionTime,inStep,inHalfDecayTime,...  
inResynthDuration,inRescalePatterns,...  
inPatternIntegrationTime,inEnhanceContrast,...  
inSaveMECResults,inUseMode);
```

Description:

Starts an entire MEC analysis and resynthesis run for the given sound file.

Input arguments:

```
inSoundFile = sound file to be processed  
inSoundFilePath = path to the sound file  
          if empty or not specified, IPEMRootDir('input')\Sounds  
          is used by default  
inMinPeriod = minimum period to look for (in s)  
          if empty or not specified, 0.5 is used by default  
inMaxPeriod = maximum period to look for (in s)  
          if empty or not specified, 5 is used by default  
inSelectionTime = selected moment in time to extract current pattern (in s)  
          if empty or not specified, 3 is used by default  
inStep = integer specifying number of samples between successive  
        calculations  
          if empty or not specified, 1 is used by default (every sample)  
inHalfDecayTime = half decay time for leaky integration of energy  
        differences (in s)  
          if empty or not specified, 1.5 is used by default  
inResynthDuration = duration (in s) of resynthesized sound  
          if empty or not specified, 10 is used by default  
inRescalePatterns = if non-zero, patterns are rescaled between 0 and 1  
          if empty or not specified, 0 is used by default  
inPatternIntegrationTime = half decay time for leaky integration of patterns  
        (in s)  
          if empty or not specified, 0 is used by default  
inEnhanceContrast = if non-zero, the "contrast" in the resynthesized sound  
        is enhanced (see IPEMMECSynthesis for more details on  
        how this is done)  
          if empty or not specified, 0 is used by default
```

```

inSaveMECResults = if non-zero, intermediate results of the analysis that
                    are necessary for later resynthesis, are saved to a .mat
                    file
                    if zero, empty or not specified, results are not saved
inAnalysisType = if 1, the RMS of the sound signal is used
                    if 2, the RMS of 10 ANI ch. (summed diff. values) is used
                    if 3, the RMS of 10 ANI ch. (separate diff. values) is used
                    if empty or not specified, 1 is used by default

```

Output:

```

outOriginalSound = original sound
outResynthesizedSound = resynthesized sound
outSampleFreq = sample frequency for all returned sounds (Hz)
outDetectedPeriods = detected periods in each channel (in s)
outIndividualSounds = individual resynthesized sounds per channel

```

Remarks:

If saving of intermediate results is requested, data is saved to a .mat file with the name 'MEC_xxx.mat', where xxx is the name of the sound file without extension. The location of this file will be IPEMRootDir('output').

Example:

```
[s,ss,fs,periods,sschan] = ...
IPEMDemoStartMEC('ligeti.wav',[],0.4,5,3,2,1.5,10,1,0.5,1,1,1);
```

Authors:

Koen Tanghe - 20020222